

# Introduction to Ray Core & Ray Ecosystem

---

**Jules S. Damji, @2twitme**

Lead Developer Advocate, Anyscale

**TBD**

Software Engineer Ray Core , Anyscale



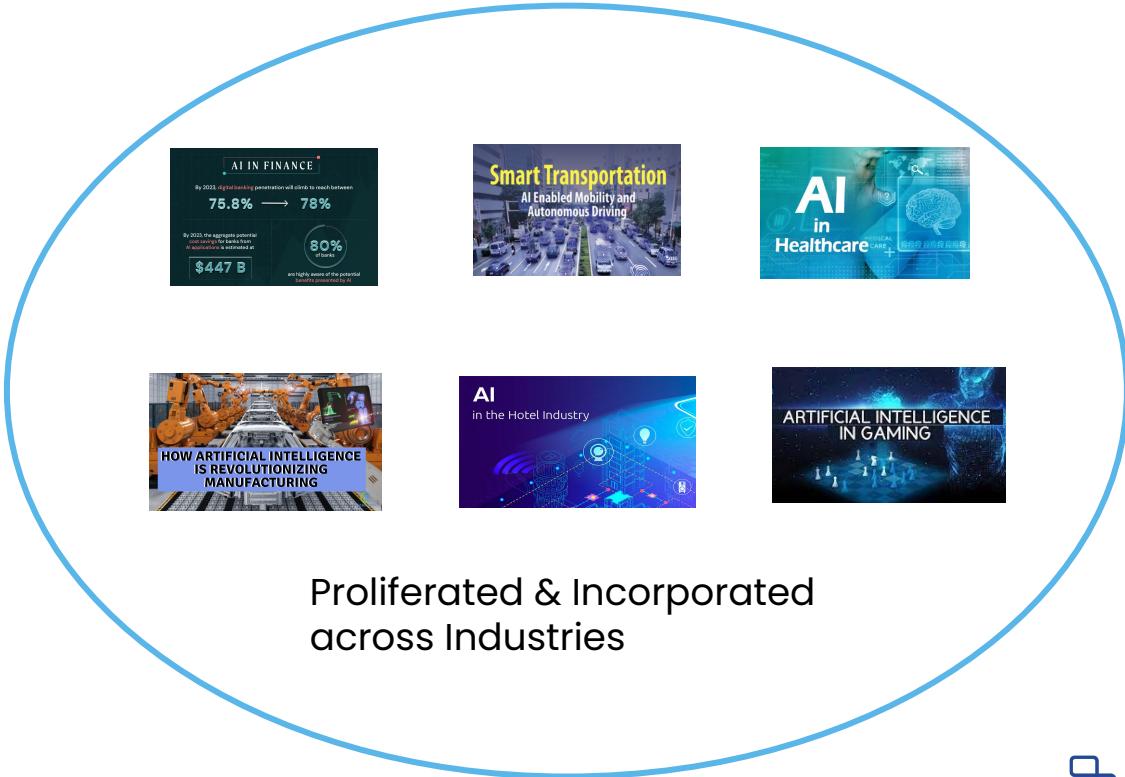
# Overview

- **What, Why Ray & Ray Ecosystem**
  - **Ray Architecture & Components**
  - **Ray Core APIs**
  - **Demo**
  - **Q & A**
-

# Why Ray?

- **Machine learning is pervasive in every domain**
- Distributed machine learning is becoming a necessity
- Distributed systems is notoriously hard

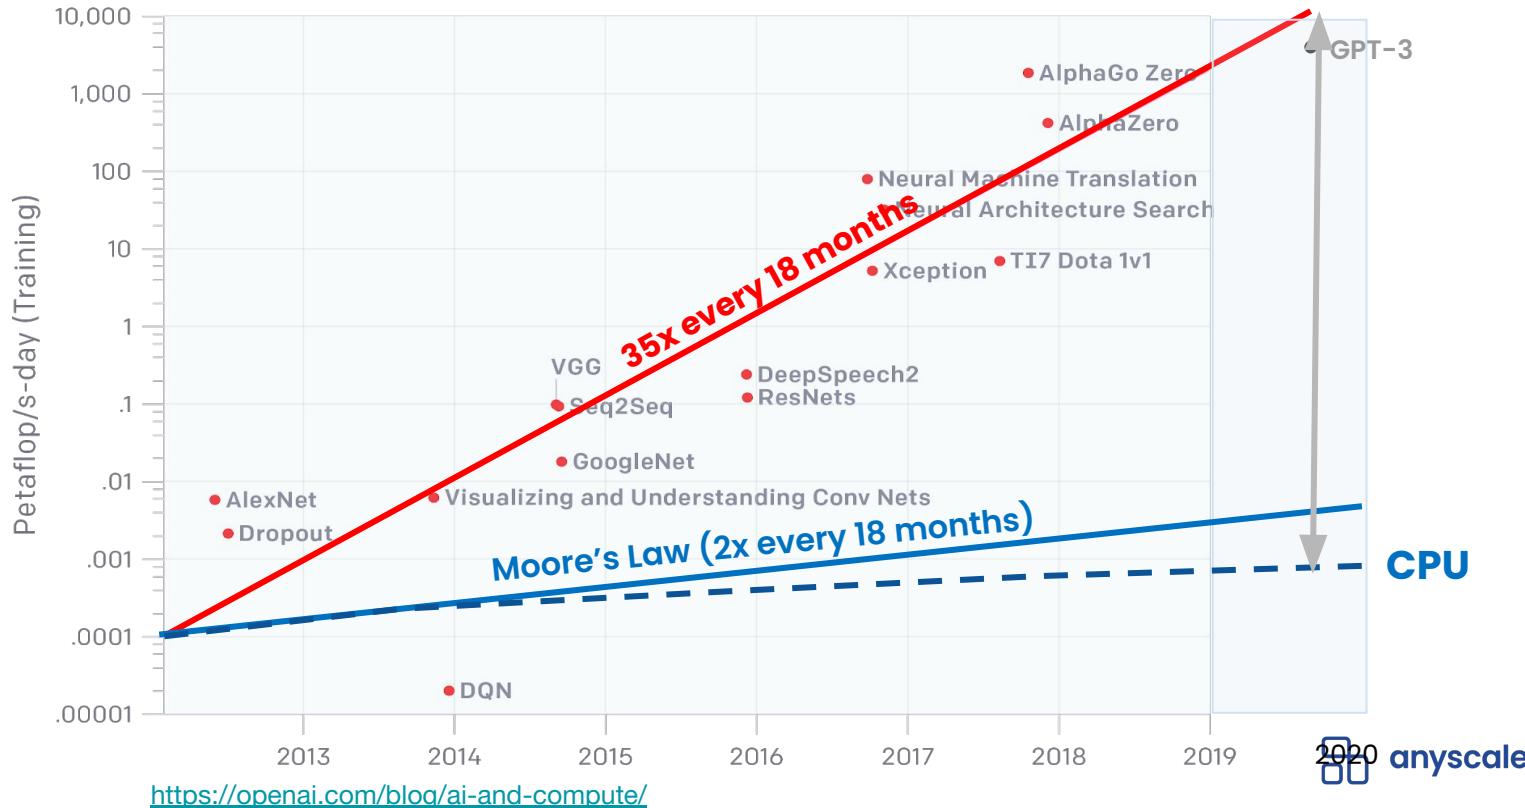
# Apps increasingly incorporate AI/ML



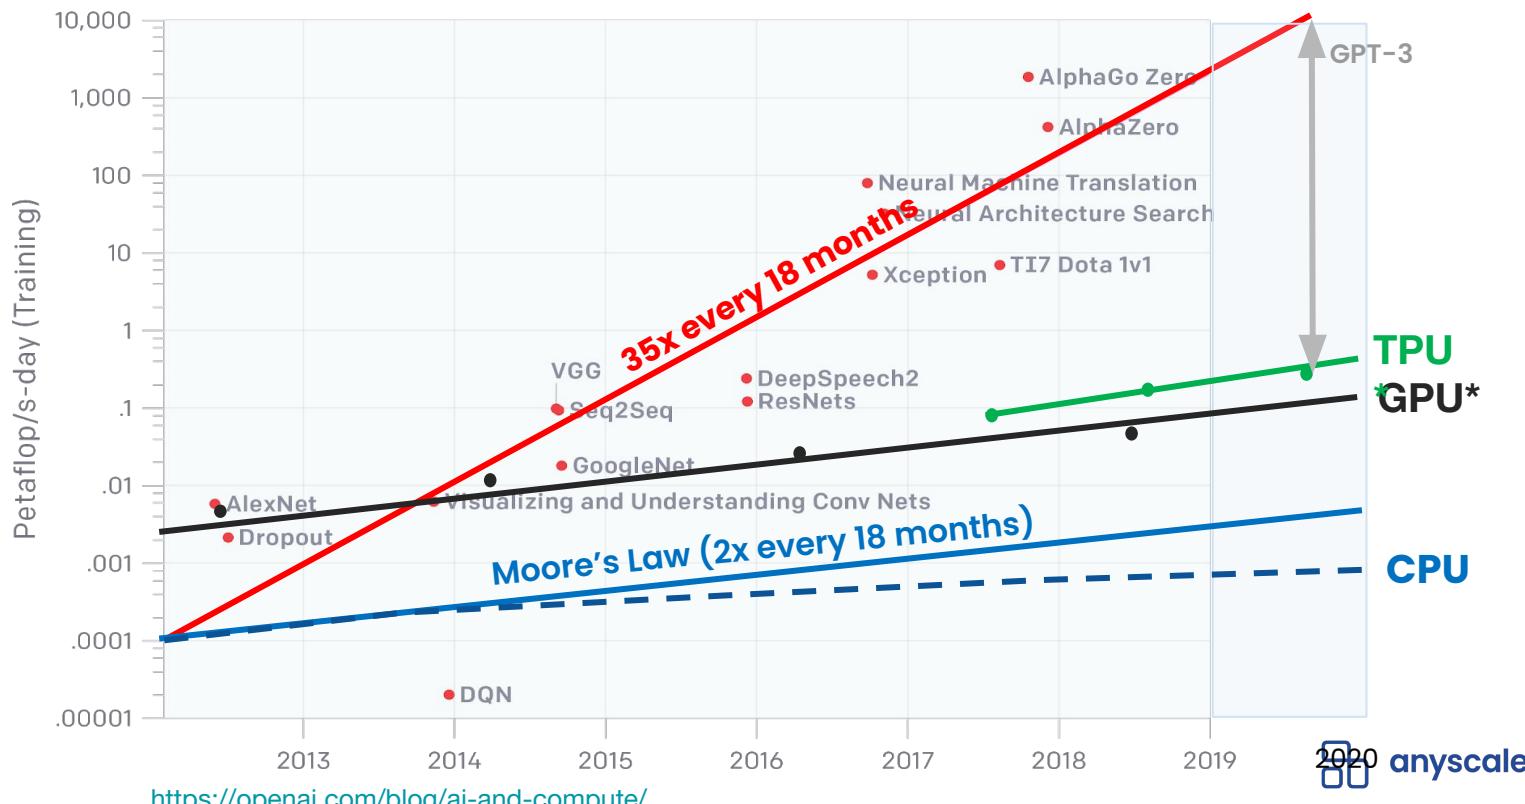
# Why Ray?

- Machine learning is pervasive in every domain
- **Distributed machine learning is becoming a necessity**
- Distributed systems is notoriously hard

# Compute demand growing faster than supply



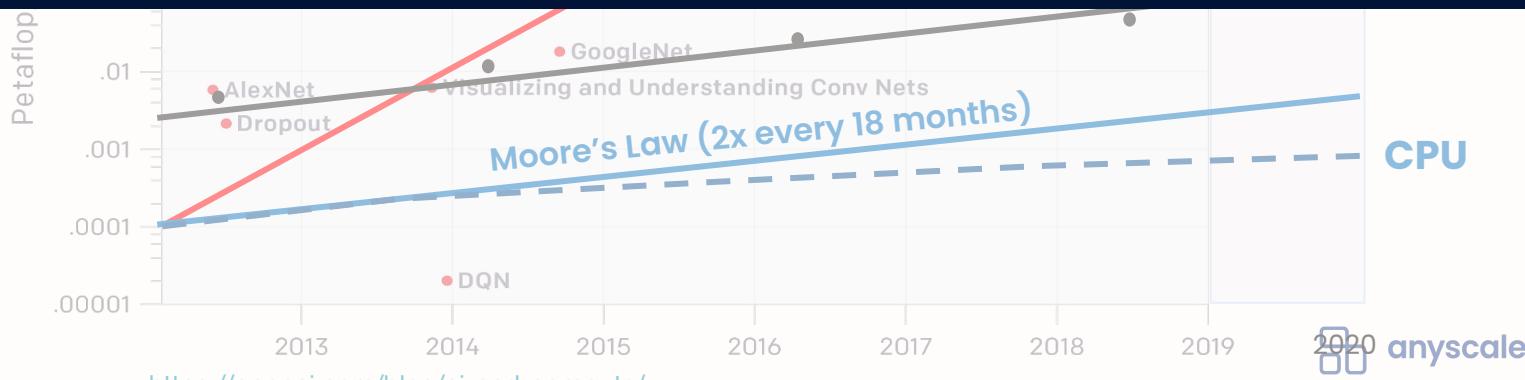
# Specialized hardware is also not enough



# Specialized hardware is also not enough



No way out but to distribute!



# Why Ray?

- Machine learning is pervasive in every domain
- Distributed machine learning is becoming a necessity
- **Distributed systems are notoriously hard**

# Existing solutions have many tradeoffs

Ease of development



Generality

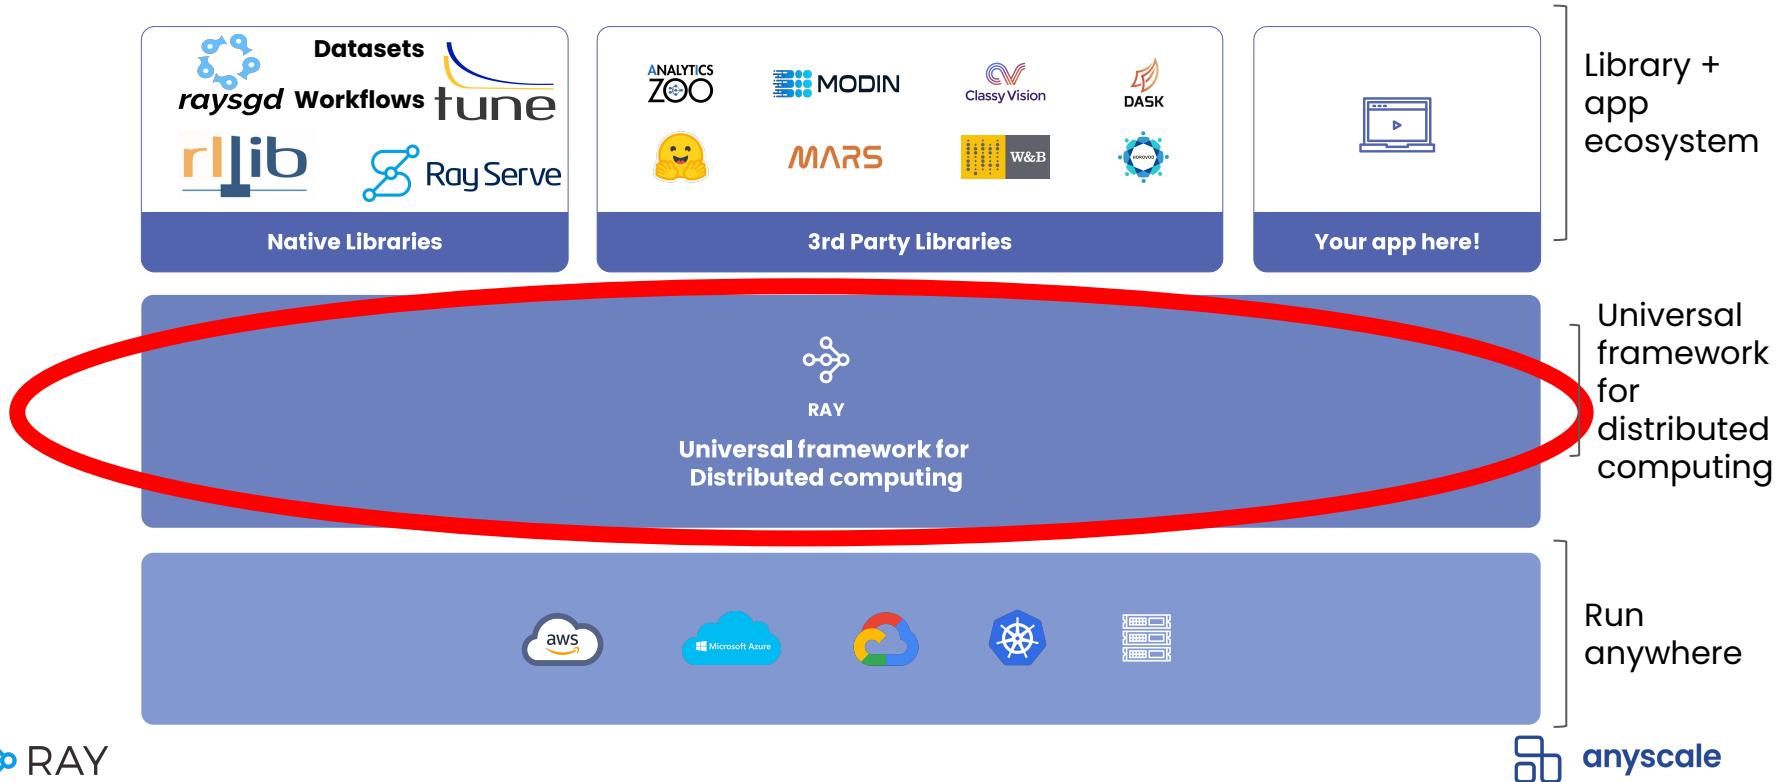
# Why Ray?

- Machine learning is pervasive in every domain
- Distributed machine learning is becoming a necessity
- Distributed systems are notoriously hard

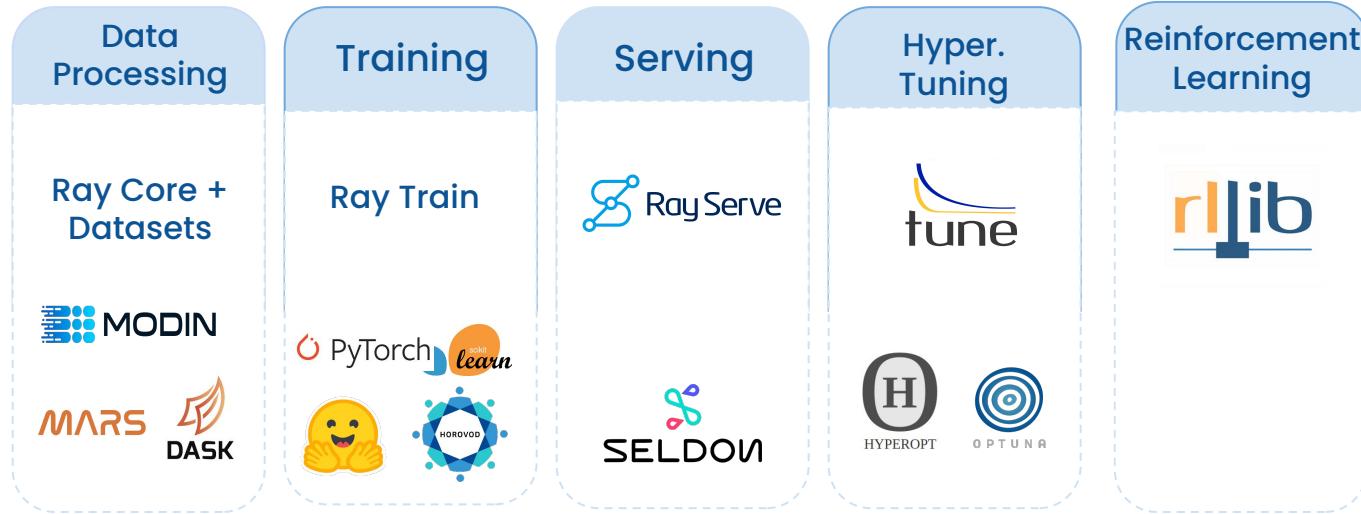
## Ray's vision:

Make distributed computing accessible to every developer

# The Ray Layered Cake and Ecosystem

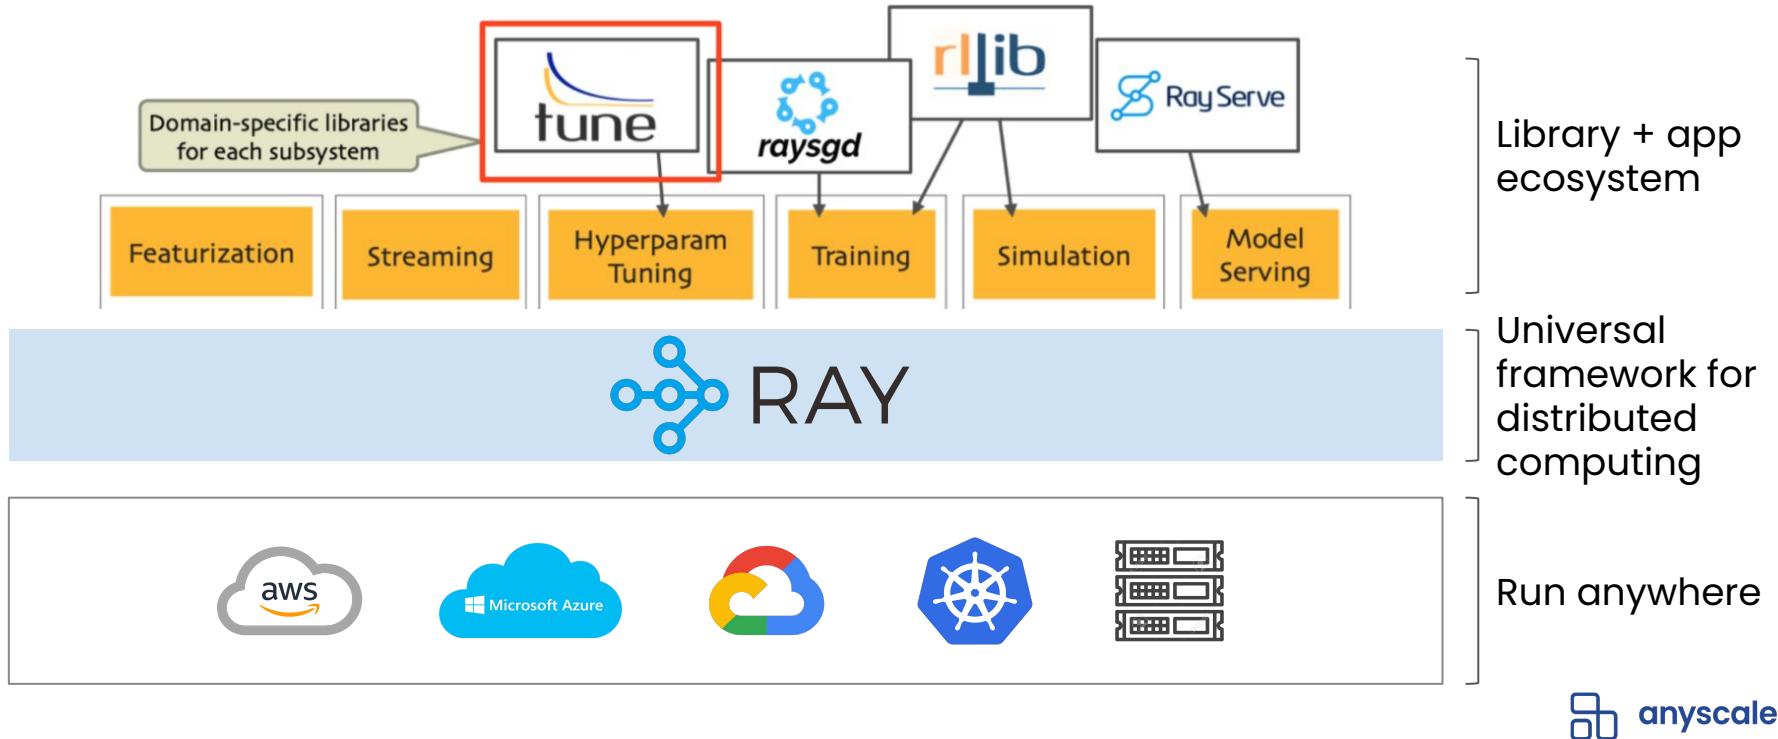


# Rich ecosystem for scaling ML workloads



Integrate Ray only based on your needs!

# Capitalize on Ray Ecosystem



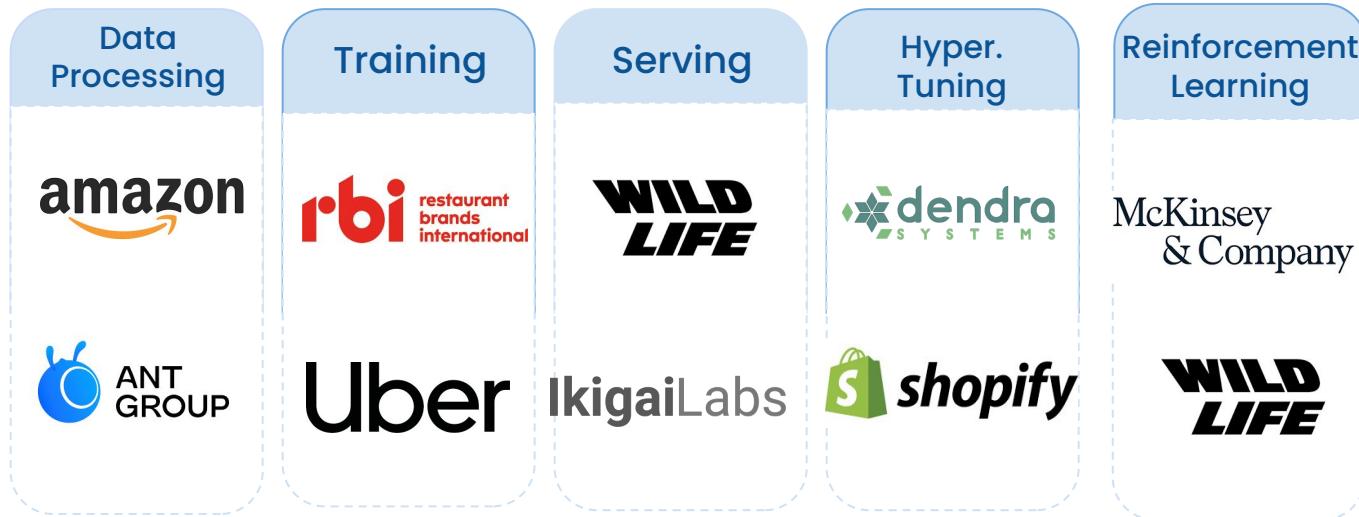
# Companies scaling ML with Ray



McKinsey  
& Company



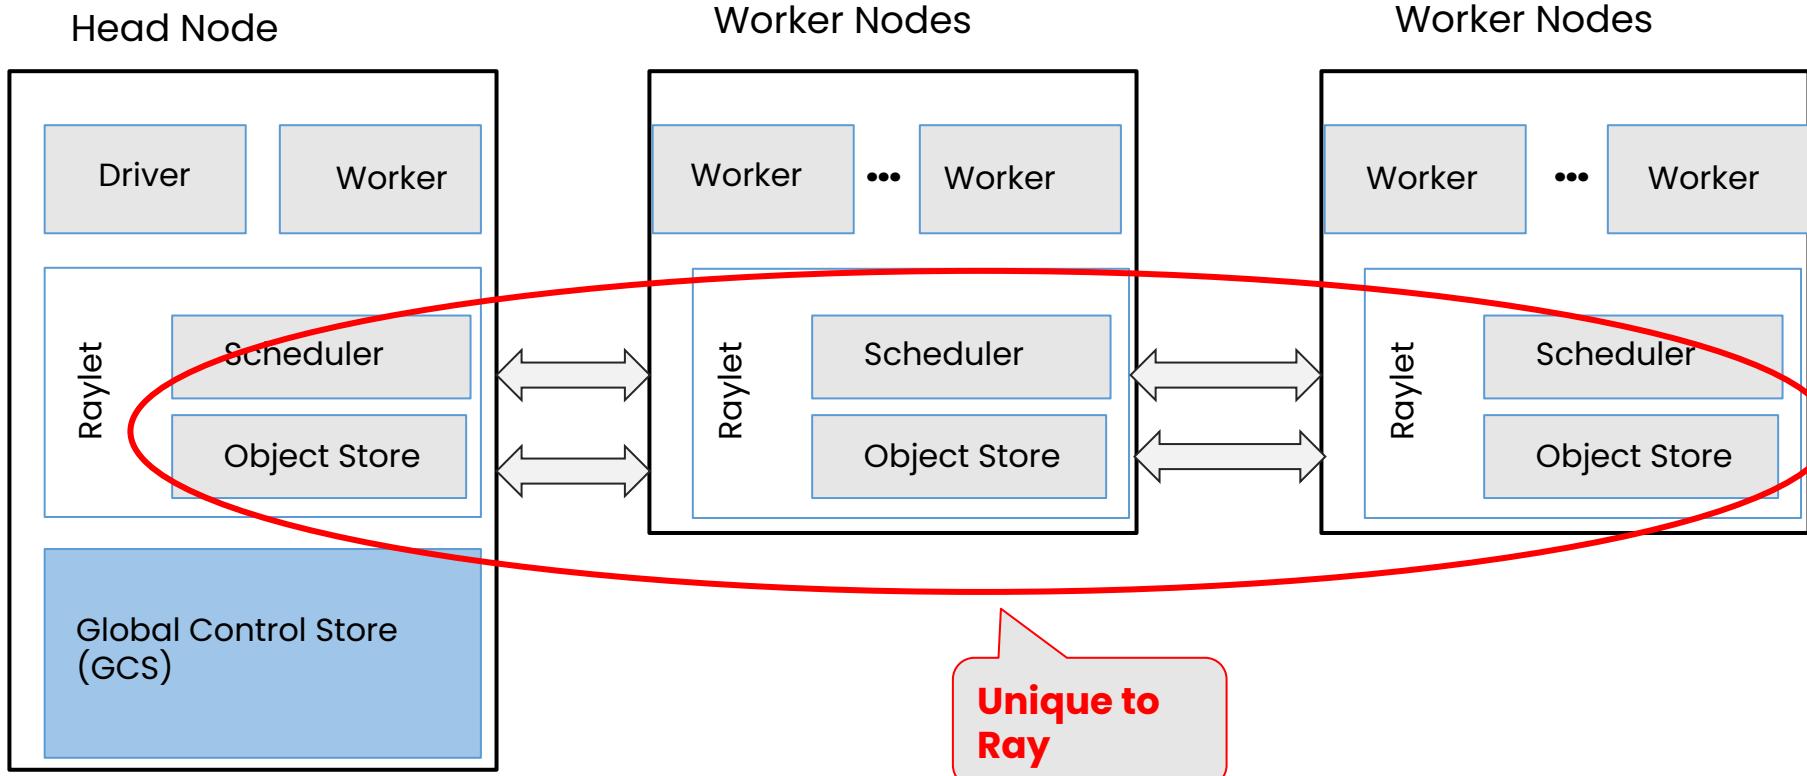
# Companies scaling ML with Ray

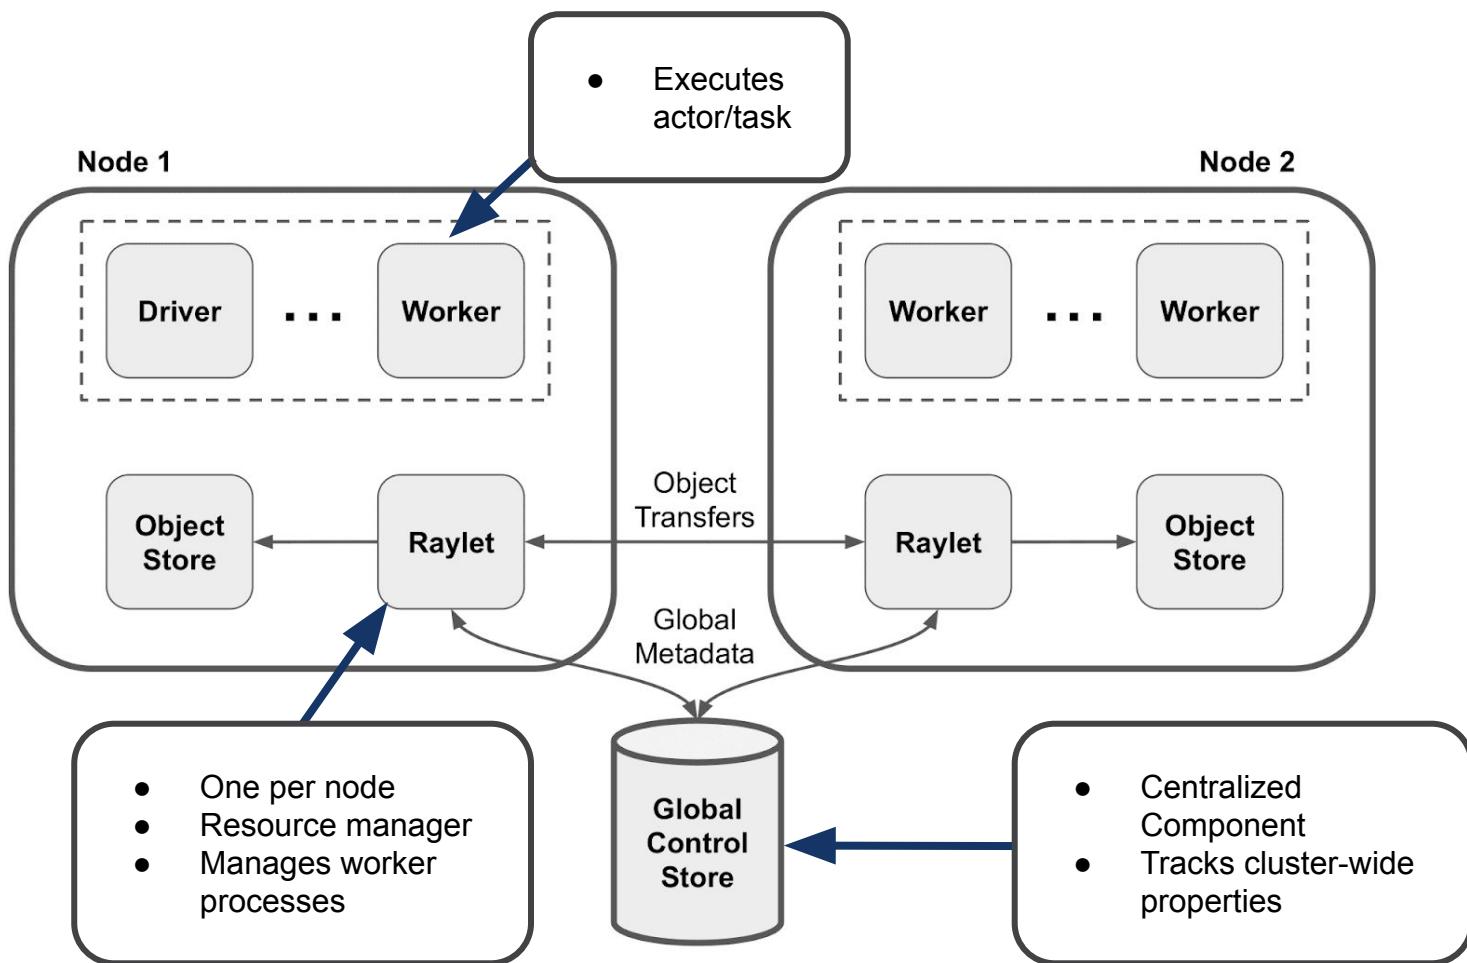


- <https://eng.uber.com/horovod-ray/>
- <https://www.anyscale.com/blog/wildlife-studios-serves-in-game-offers-3x-faster-at-1-10th-the-cost-with-ray>
- <https://www.ikigailabs.com/blog/how-ikigai-labs-serves-interactive-ai-workflows-at-scale-using-ray-serve>

# Ray Architecture & Components

# What does Ray Cluster Looks Like ...





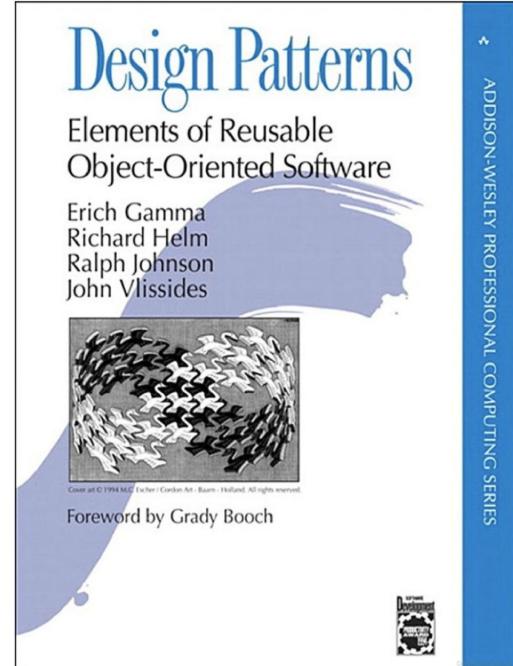
# Ray Distributed Design Patterns & APIs



# Ray Basic Design Patterns

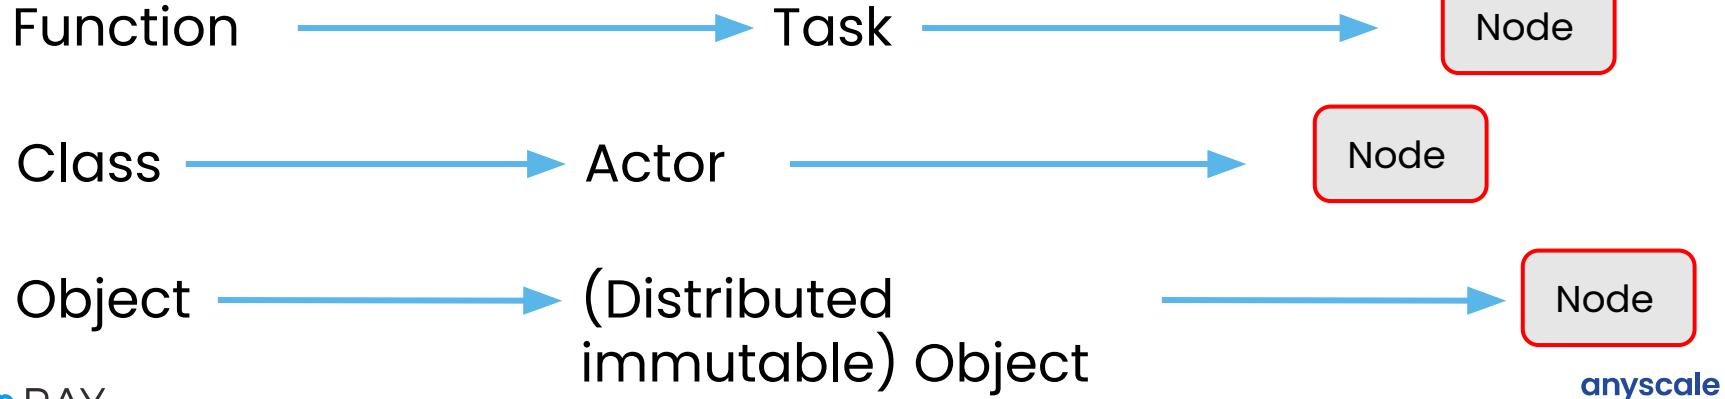
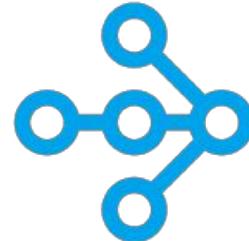
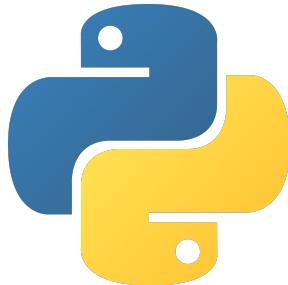
- **Ray Parallel Tasks**
  - Functions as stateless units of execution
  - Functions distributed across a clusters as tasks
- **Ray Objects or Futures**
  - Distributed (immutable) Object stored in cluster
  - Retrievable when available
  - Enable asynchronous execution of
- **Ray Actors**
  - Stateful service on a cluster
  - Message passing and maintains state

1. [Patterns for Parallel Programming](#)
2. [Ray Design Patterns](#)





# Python → Ray Basic Patterns



# Function

```
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

def add(a, b):
    return np.add(a, b)

a = read_array(file1)
b = read_array(file2)
sum = add(a, b)
```

# Class

```
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
    return self.value

c = Counter()
c.inc()
c.inc()
```

# Function → Task

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

a = read_array(file1)
b = read_array(file2)
sum = add(a, b)
```

# Class → Actor

```
@ray.remote
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
        return self.value

c = Counter()
c.inc()
c.inc()
```

# Function → Task

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```

# Class → Actor

```
@ray.remote(num_gpus=1)
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
        return self.value

c = Counter.remote()
id4 = c.inc.remote()
id5 = c.inc.remote()
```

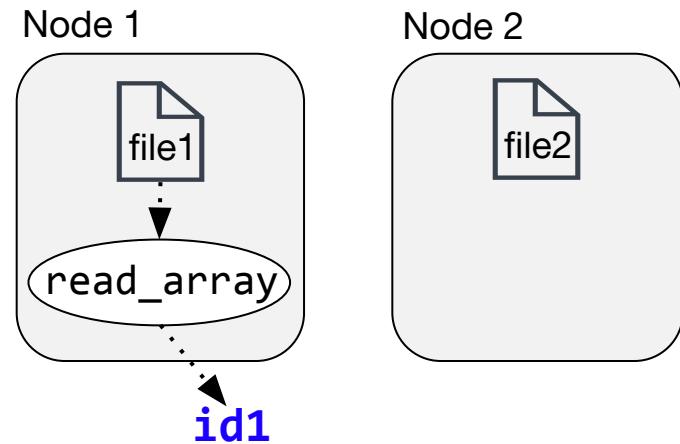
# Task API

Blue variables are ObjectRef IDs  
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```



Return **id1** (future) immediately,  
before **read\_array()** finishes

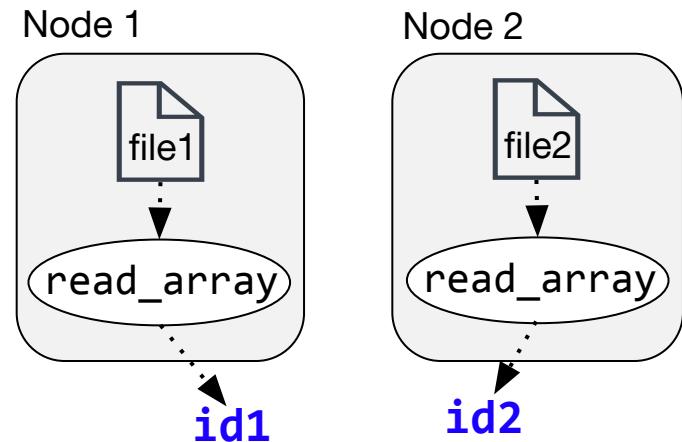
# Task API

Blue variables are Object IDs  
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```



Dynamic task graph:  
build at runtime

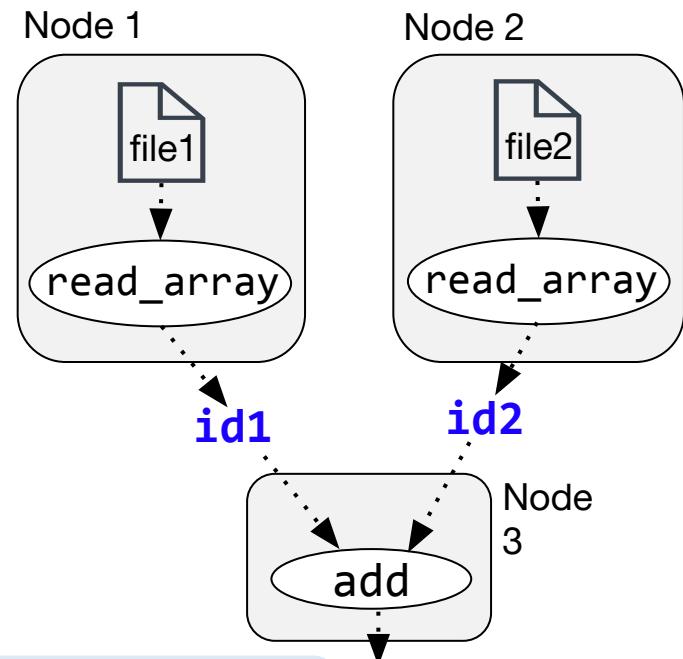
# Task API

Blue variables are Object IDs  
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```



Every task scheduled,  
but not finished yet

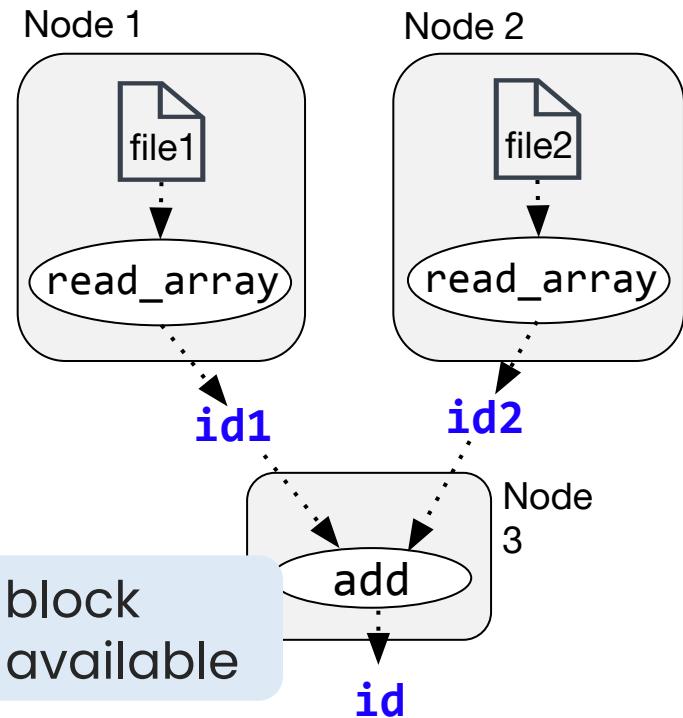
# Task API

Blue variables are Object IDs  
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```



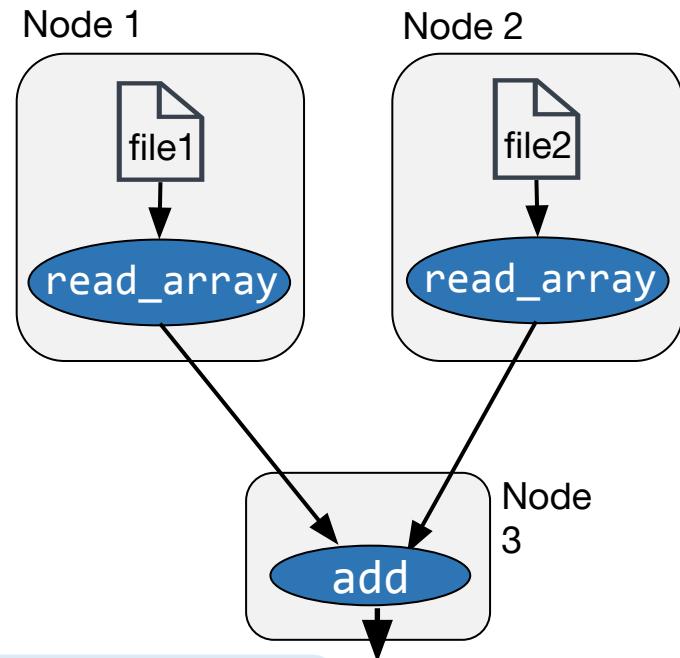
# Task API

Blue variables are Object IDs  
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

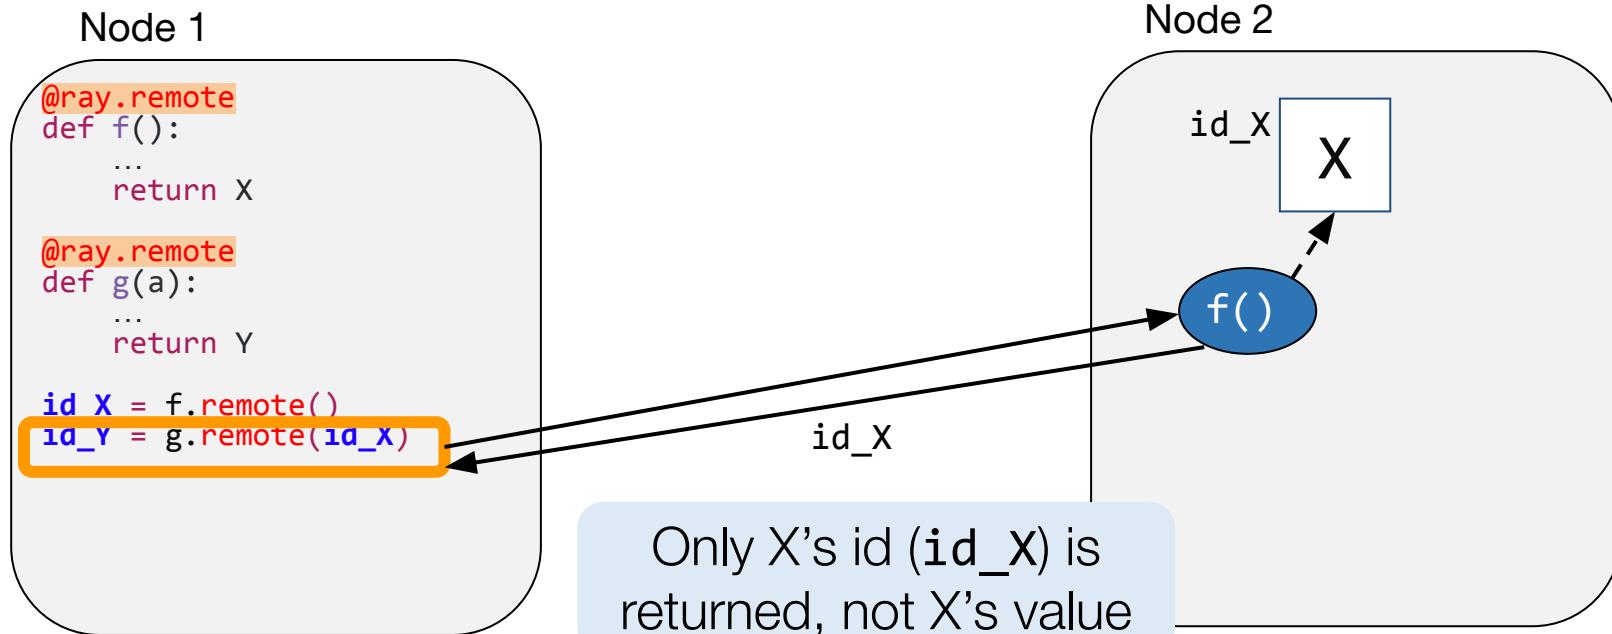
@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```

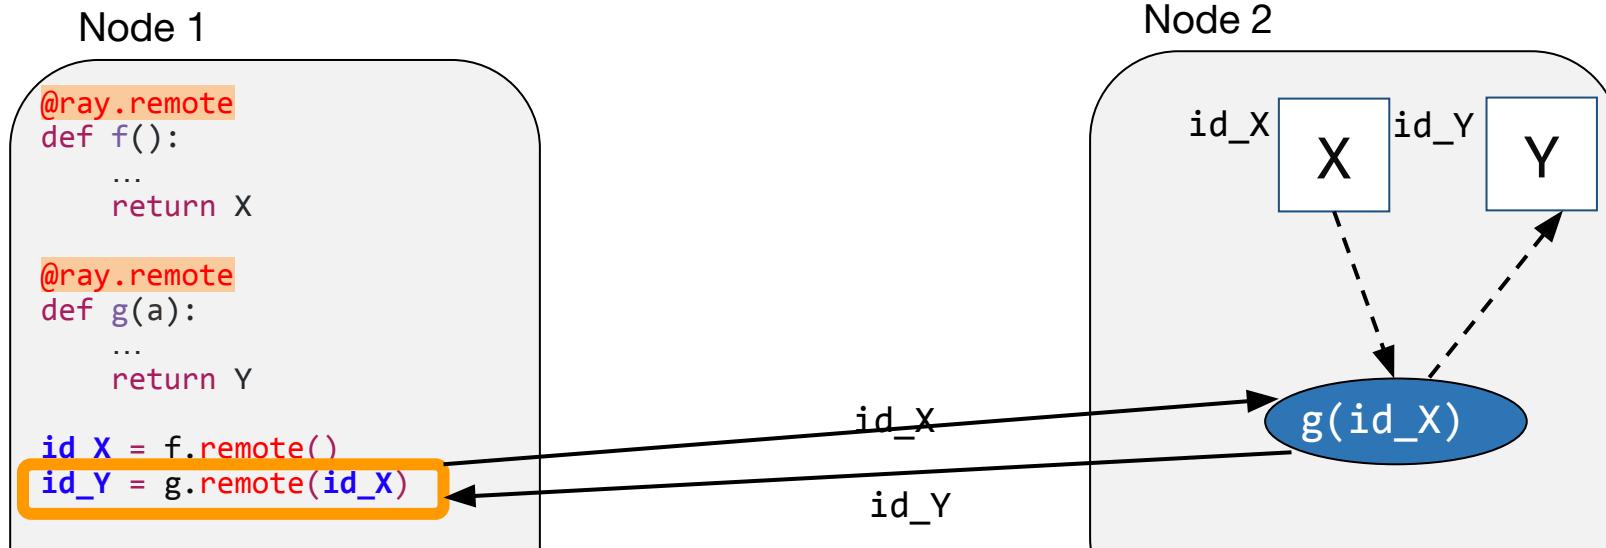


Task graph executed to sum  
compute sum

# Distributed Immutable object store



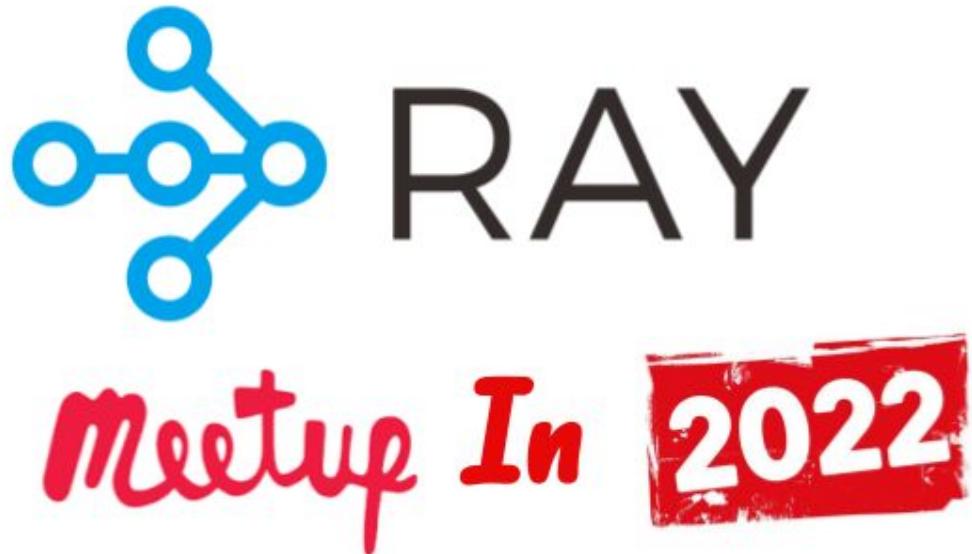
# Distributed object store



`g(id_X)` is scheduled on same node, so `X` is never transferred

# Takeaways

- **Distributed computing is a necessity & norm**
- **Scale your ML workloads with Ray Libraries**
- **Ray's vision: make distributed programming simple**
- **Don't have to be distributed systems expert. Just use `@ray.remote` :)**



<https://tinyurl.com/ray-meetup-jan2022>

Jan 19 th @ 6:00 PM PST

# Demo

---

<https://tinyurl.com/ray-core-demo>

## You might be interested in ...

- [How to scale ML workloads with Ray](#)
- [How to Effectively tune HPO with Ray Tune](#)
- [How to Scale Model Serving with Ray Serve](#)
- [Scalable RL training & serving in the cloud with Ray](#)
- Let us know what you are looking for ....

# Start learning Ray and contributing ...

## Getting Started:

### Documentation ([docs.ray.io](https://docs.ray.io))

*Quick start example, reference guides, etc*

### Join Ray Meetup

*Revive in Jan 2022. Publish recording to the members*

<https://www.meetup.com/Bay-Area-Ray-Meetup/>

### Forums ([discuss.ray.io](https://discuss.ray.io))

*Learn / share with broader Ray community, including core team*

### Ray Slack

*Connect with the Ray team and community*

### Social Media (@raydistrributed, @anyscalecompute)

*Follow us on Twitter and linkedIn*

### GitHub

*Check out sources, file an issue, become a contributor, give us a Star :)*

<https://github.com/ray-project/ray>

<https://github.com/dmatrix/ray-core-tutorial>

# Thank you!

---

Let's stay in touch:

jules@anyscale.com  
<https://www.linkedin.com/in/dmatrix/>



@2twitme

# Extras

---



# ray.wait() non-blocking usage

```
c = Counter.remote()
increment_refs = [c.inc.remote() for _ in range(5)]
while len(increment_refs) > 0:
    return_n = 2 if len(increment_refs) > 1 else 1
    ready_refs, remaining_refs = ray.wait(increment_refs,
                                           num_returns=return_n,
                                           timeout=10.0)
    if len(ready_refs) > 0:
        print(ray.get(ready_refs))
    # Update the remaining ones
    increment_refs = remaining_refs
[1, 2]
[3, 4]
[5]
```

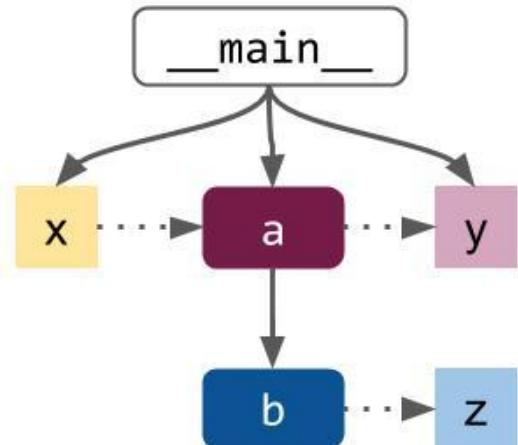




# Object Ownership

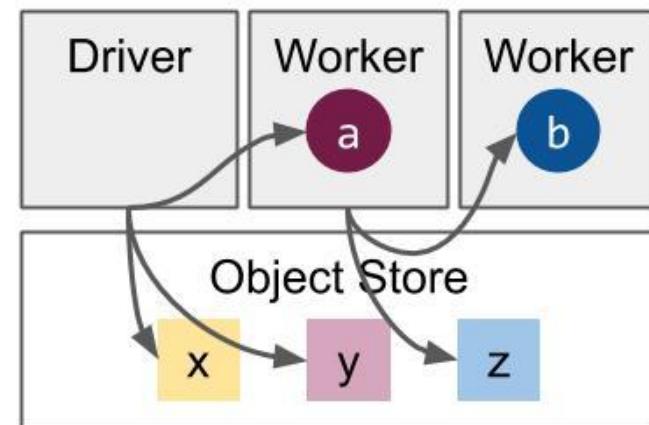
```
@ray.remote  
def b():  
    return  
  
@ray.remote  
def a(dep):  
    z = b.remote()  
  
x = ray.put(...)  
y = a.remote(x)
```

Program



Task graph

→ Ownership  
→ Dependency



Physical execution