

Distributed Python & ML with Ray: Hands-on with Ray Core APIs

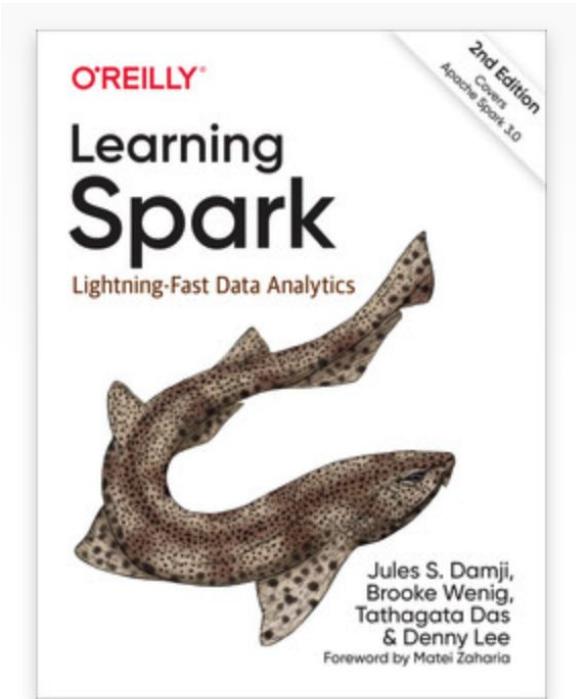
Jules S. Damji, @2twitme

Lead Developer Advocate, Ray Team @ Anyscale
PyCon US, Salt Lake City, UT, April 27, 2022



\$whoami

- Lead Developer Advocate @Anyscale
- Senior Developer Advocate @Databricks
- Led Developer Advocacy @Hortonworks
- Held software engineering positions:
 - Sun Microsystems
 - Netscape
 - @Home
 - LoudCloud/Opsware
 - Verisign
 - ProQuest/Ebrary

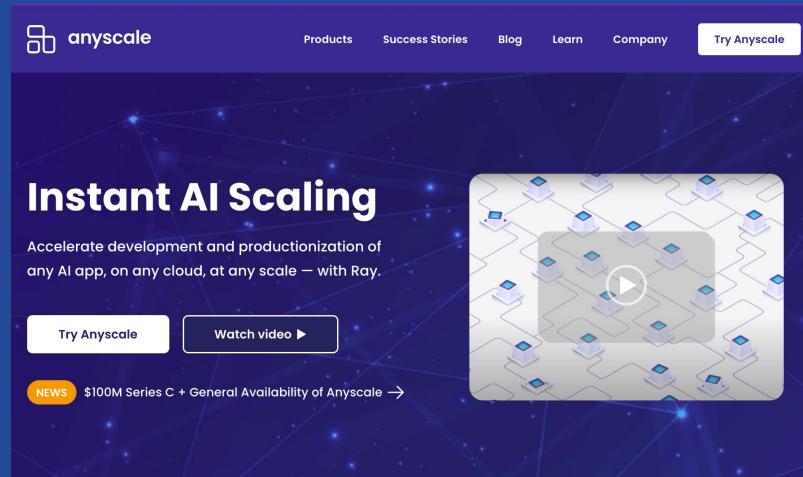


Anyscale

Who we are: Original creators of Ray

What we do: Provide managed service for Ray on public cloud to scale AI workloads

Why do it: Make distributed computing easy and simple for everyone



Overview

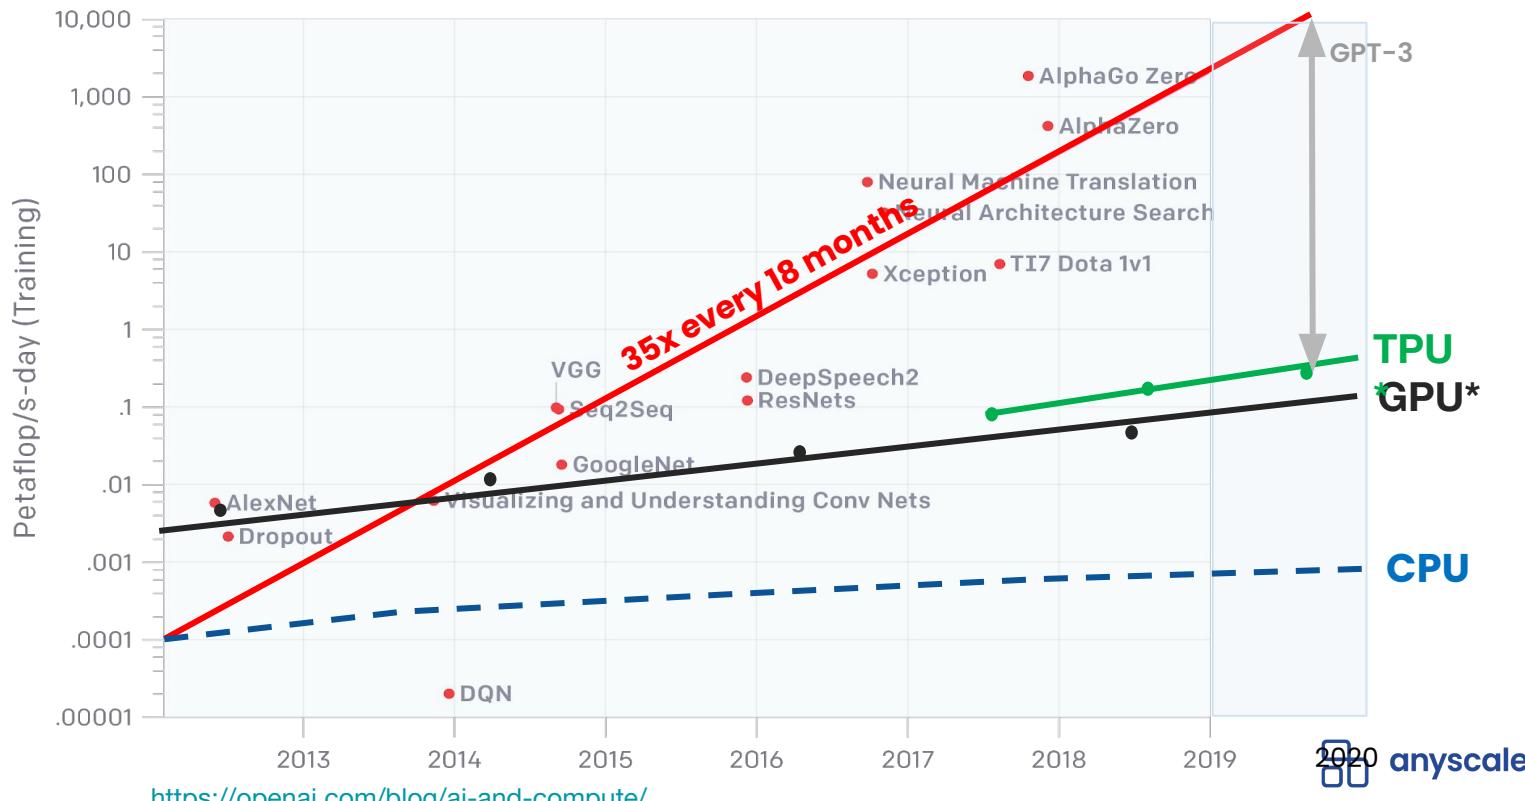
- **Why & What's Ray & Ray Ecosystem**
- **Ray Architecture & Components**
- **Ray Core Design Patterns & APIs**
- **Hands-on Tutorials**

<https://tinyurl.com/pycon-ray-tutorial>

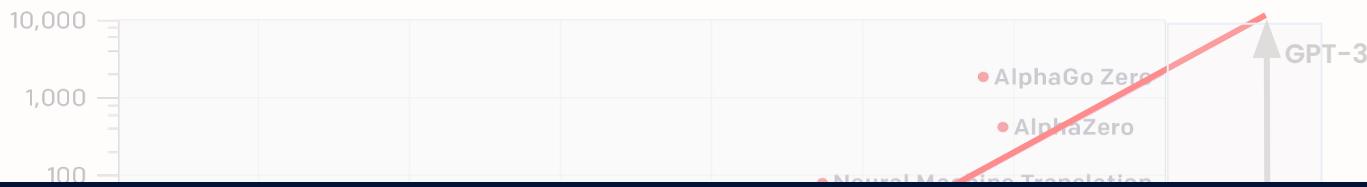
Why Ray ?

-  Machine learning is pervasive in every domain
-  Distributed machine learning is becoming a necessity
-  Python is the default language for DS & ML

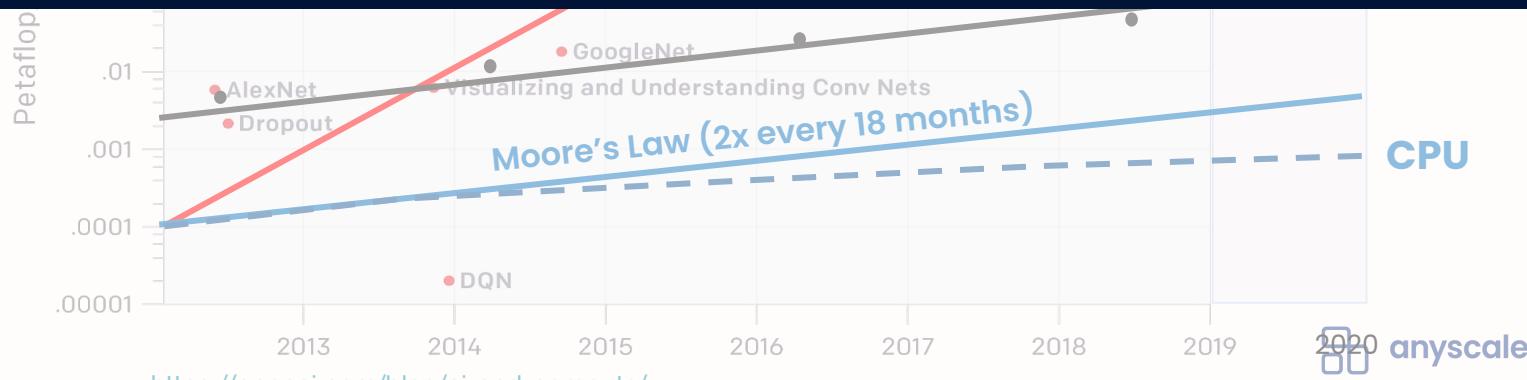
Compute Demand - supply problem



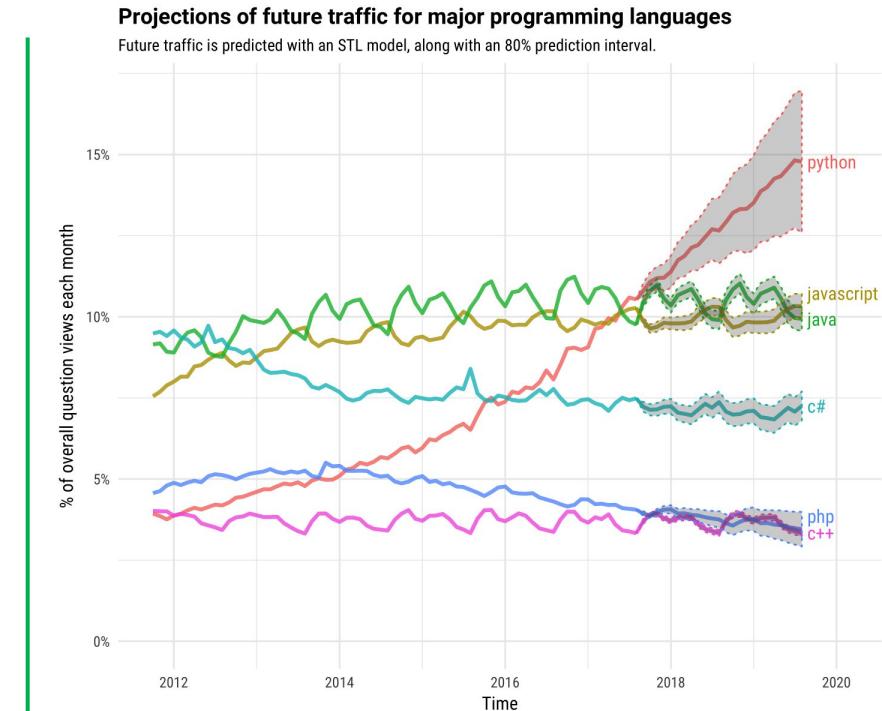
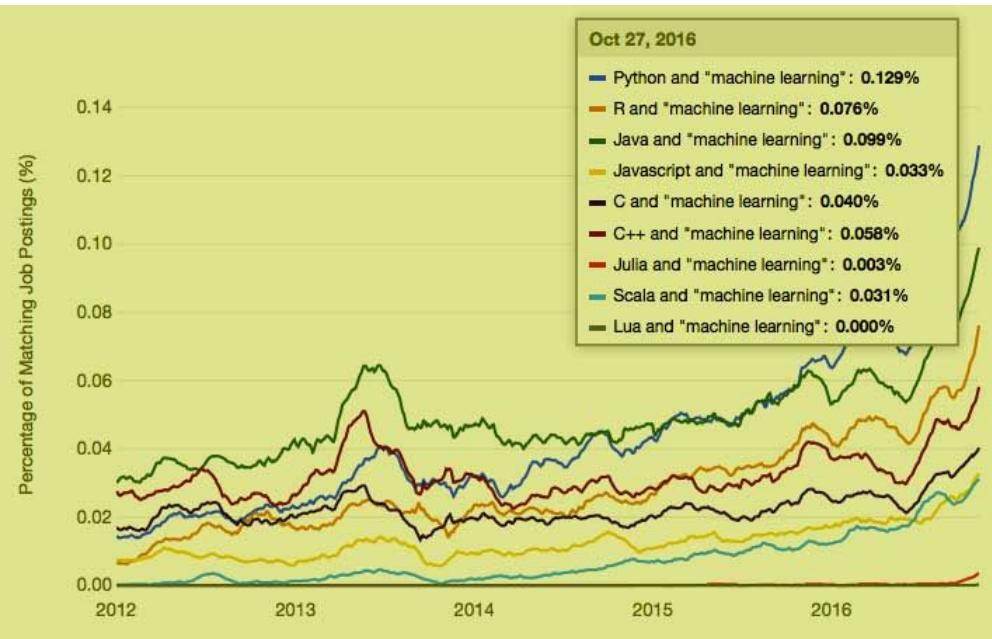
Specialized hardware is also not enough



No way out but to distribute!

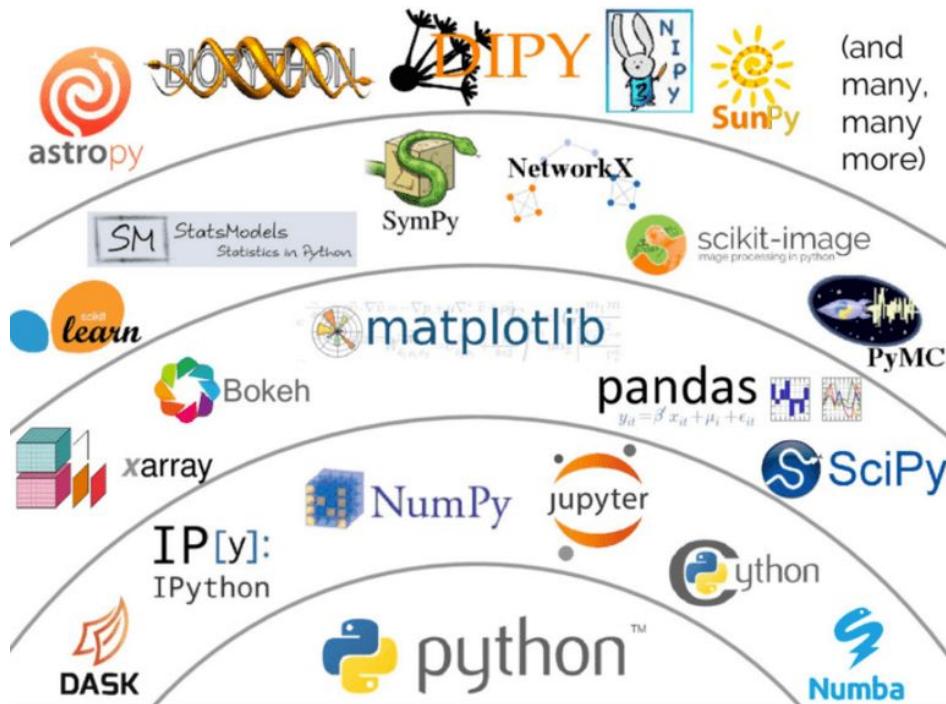


Python's growth as programming language

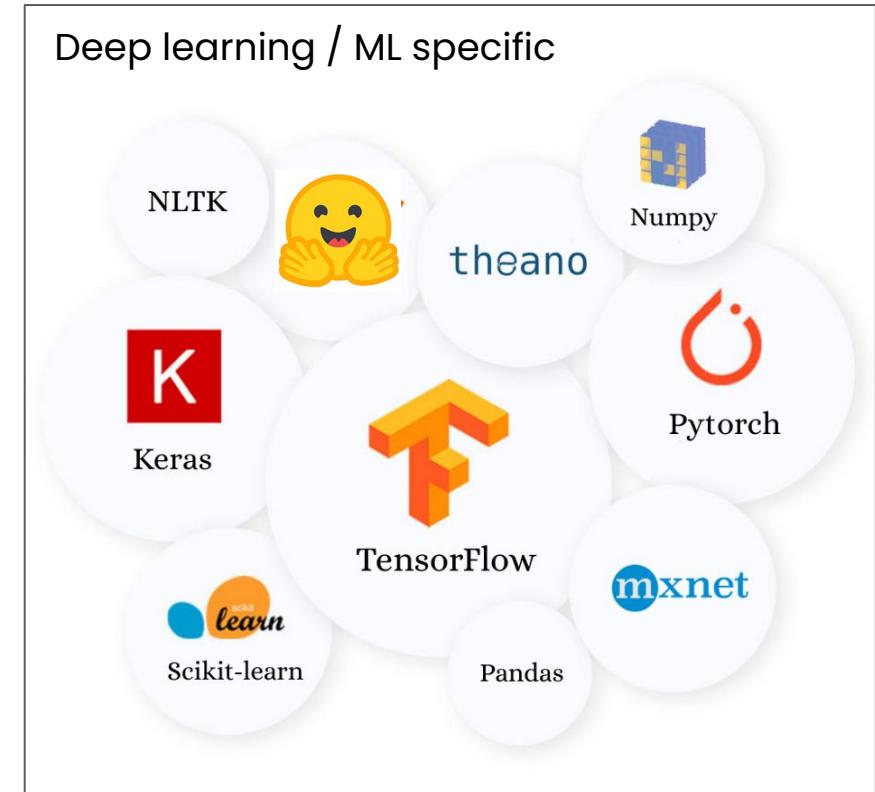




Python data science ecosystem dominating



Deep learning / ML specific



Origins of Ray

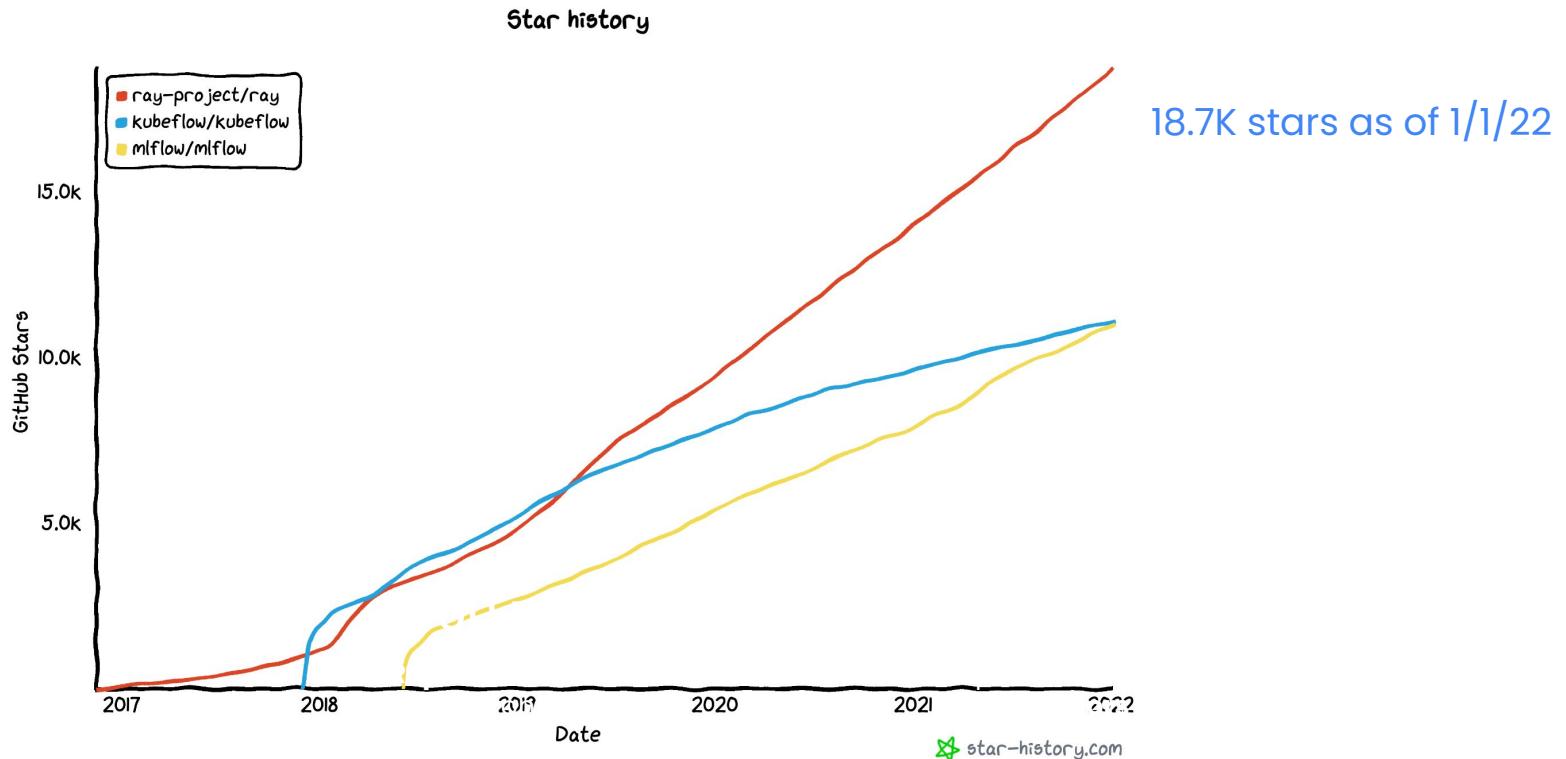


Caffe



RAY

Growth of open-source Ray adoption

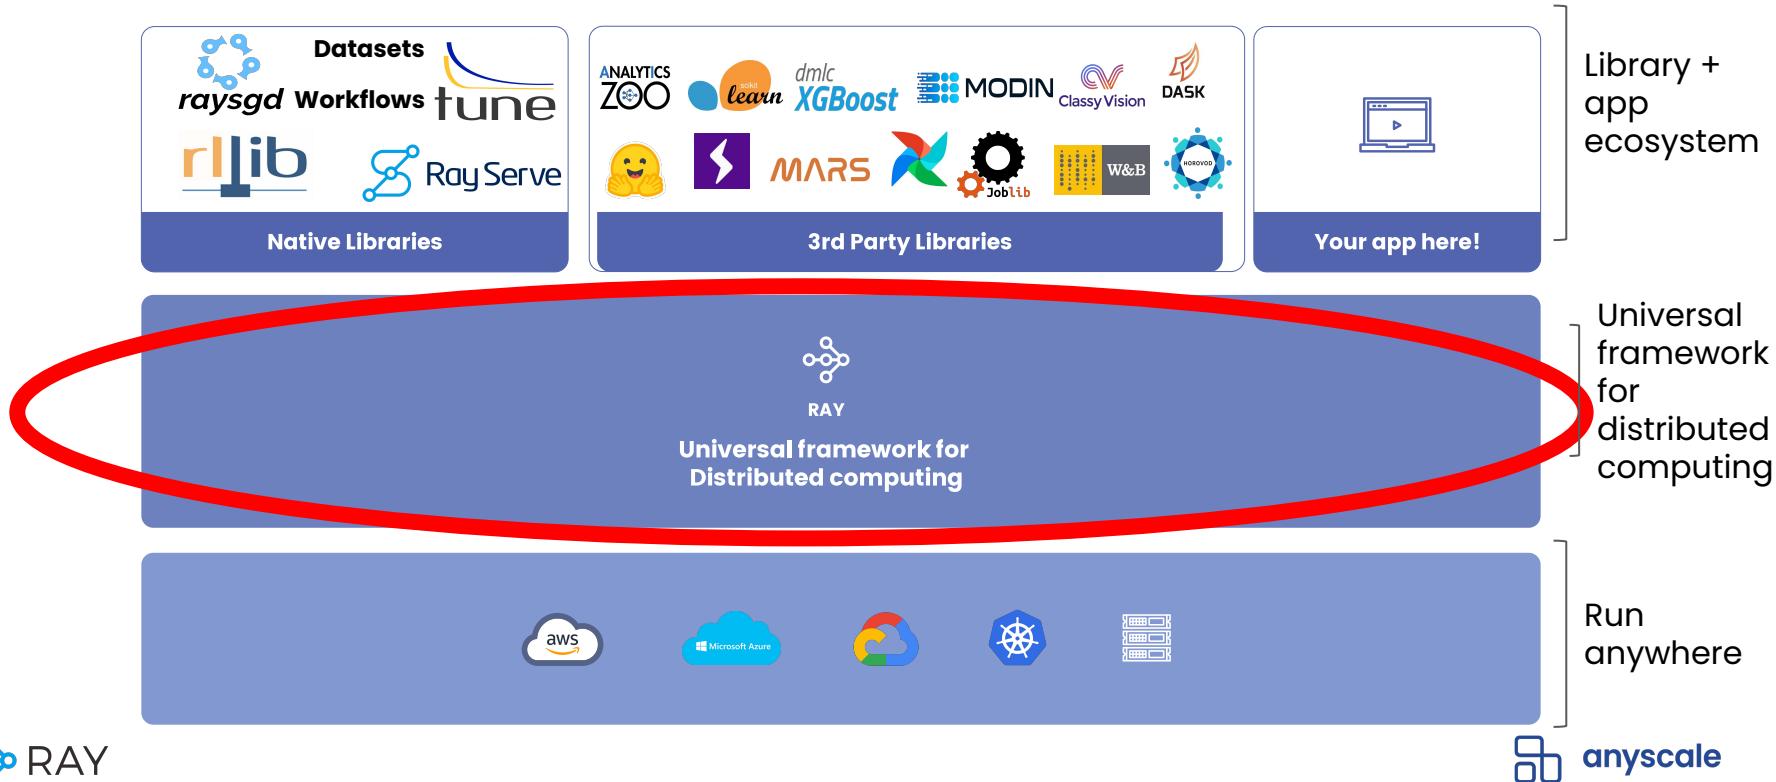


What is Ray?

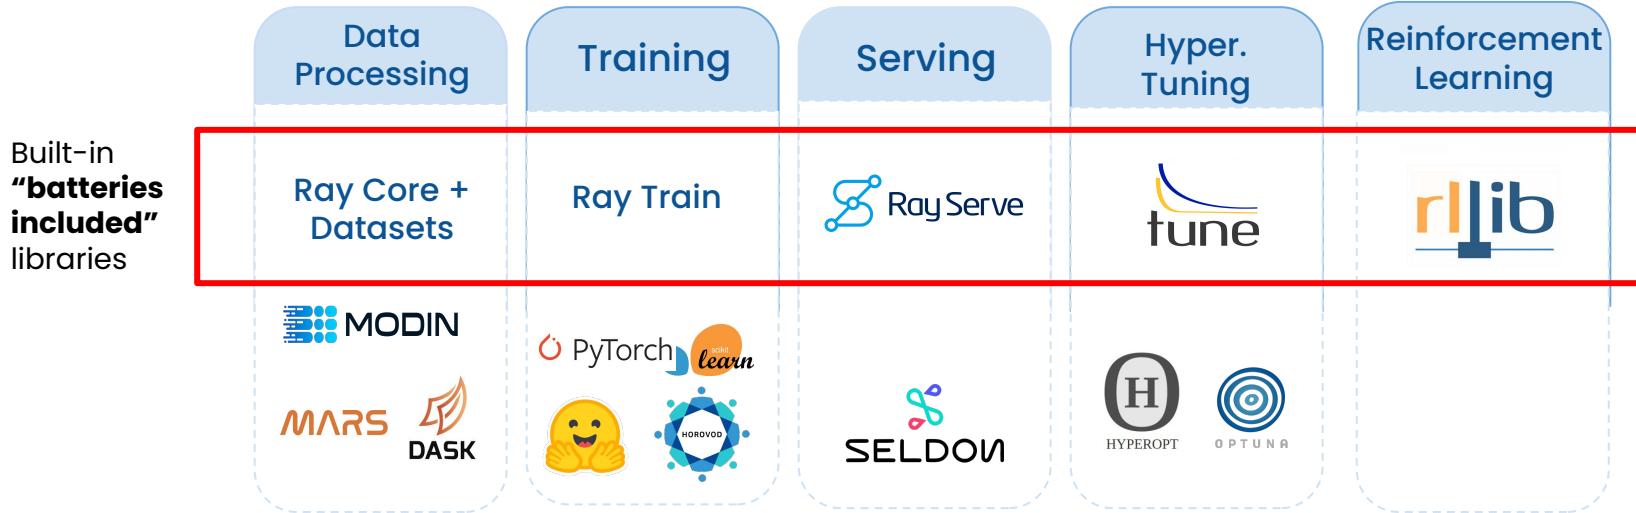
- A simple/general library for distributed computing
- An ecosystem of Python libraries (for scaling ML and more)
- Runs on laptop, public cloud, K8s, on-premise

A layered cake of functionality and capabilities for scaling ML workloads

The Ray Layered Cake and Ecosystem



Rich ecosystem for scaling ML workloads



Only use the libraries you need!

Companies scaling ML with Ray

amazon

intel

Microsoft

VISA

rbi
restaurant
brands
international

WILD
LIFE

shopify

IBM

dendra
SYSTEMS

Uber

McKinsey
& Company

ANT
GROUP

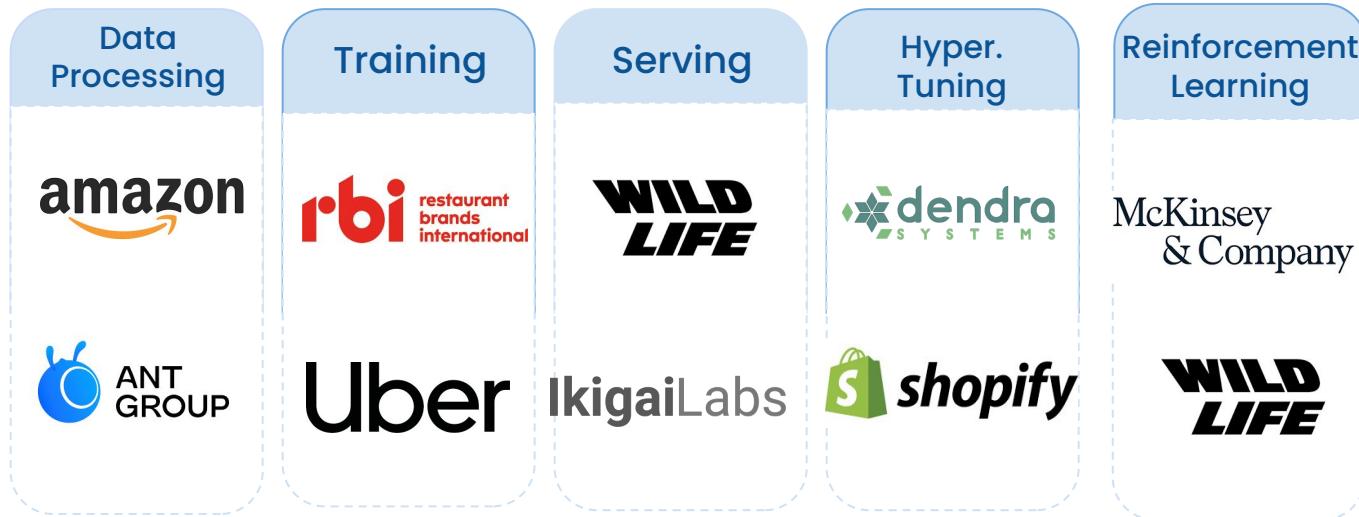
robinhood

NVIDIA.

2 σ TWO SIGMA

Alibaba.com

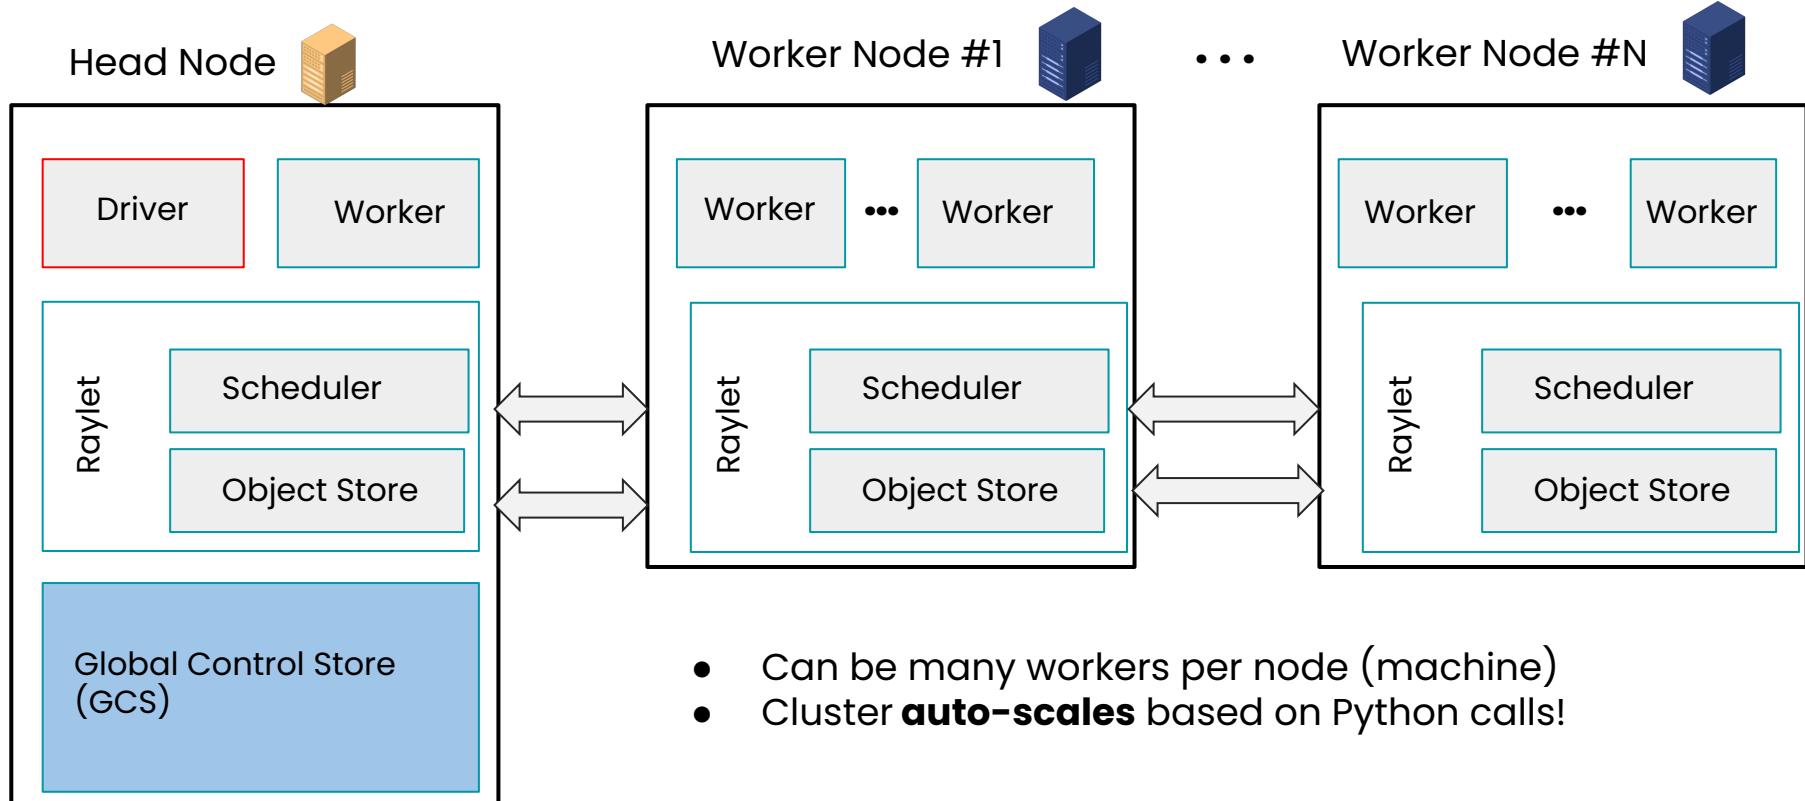
Companies scaling ML with Ray



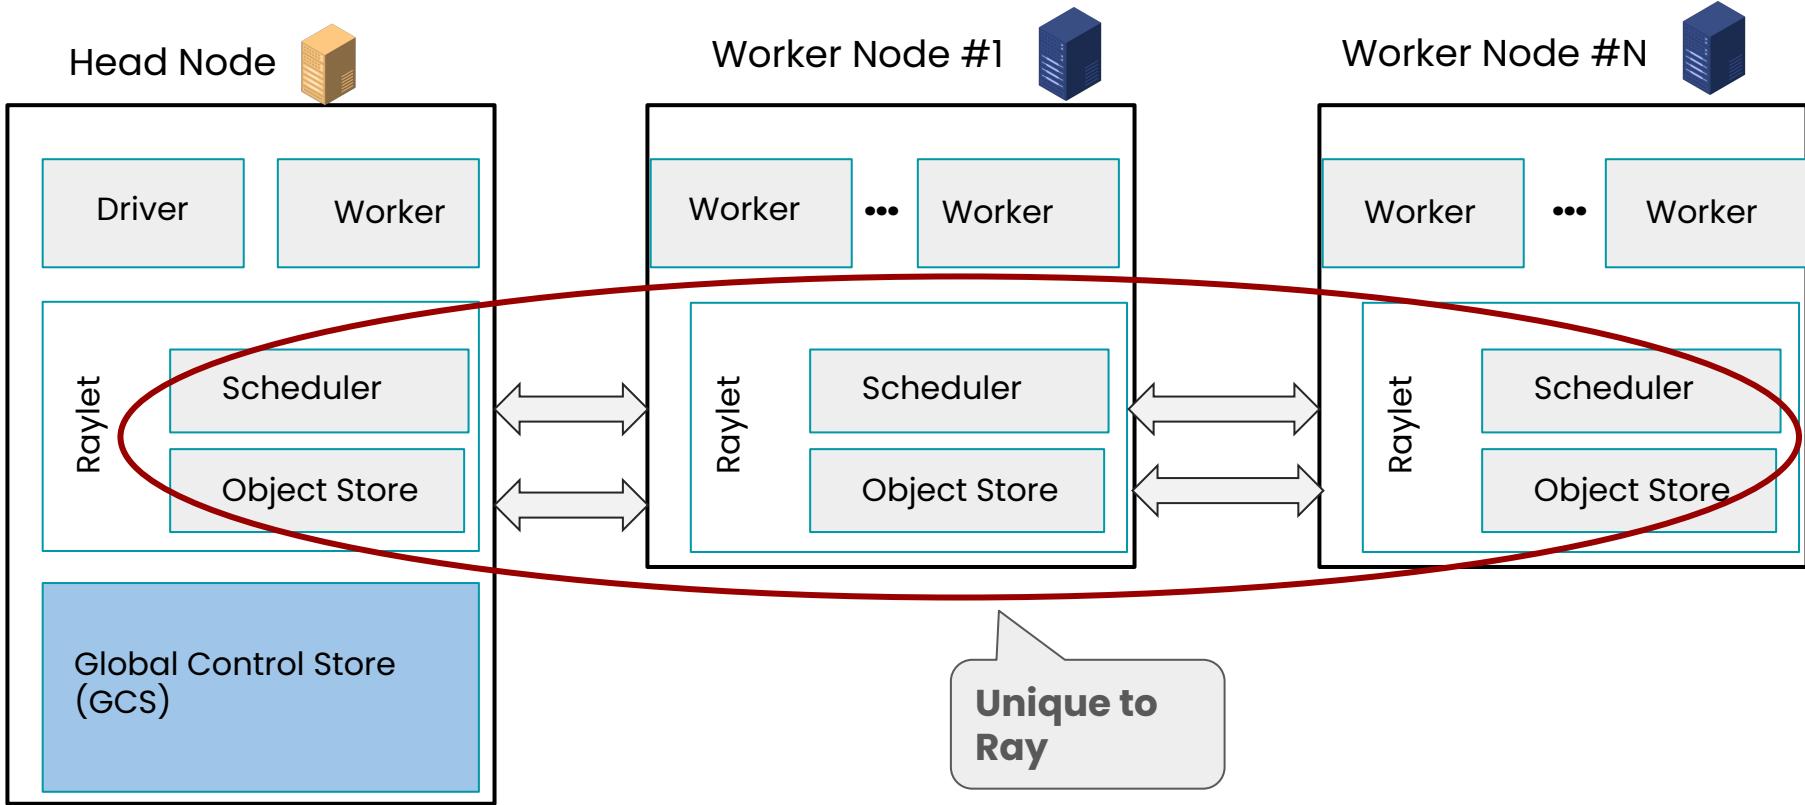
- <https://eng.uber.com/horovod-ray/>
- <https://www.anyscale.com/blog/wildlife-studios-serves-in-game-offers-3x-faster-at-1-10th-the-cost-with-ray>
- <https://www.ikigailabs.com/blog/how-ikigai-labs-serves-interactive-ai-workflows-at-scale-using-ray-serve>

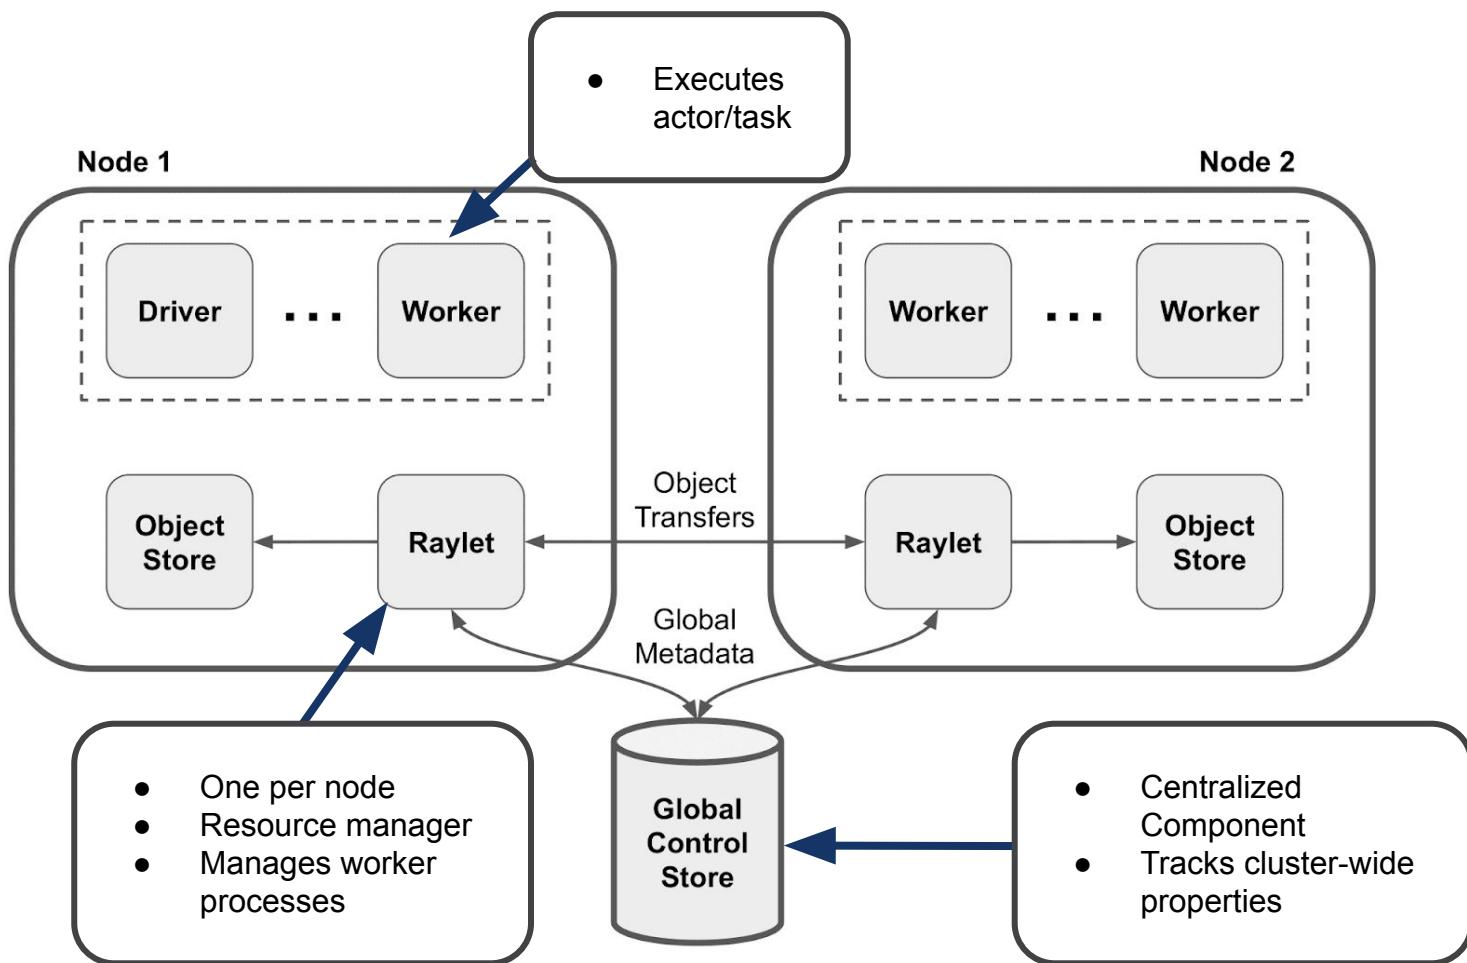
Ray Architecture & Components

Anatomy of a Ray cluster



Anatomy of a Ray cluster





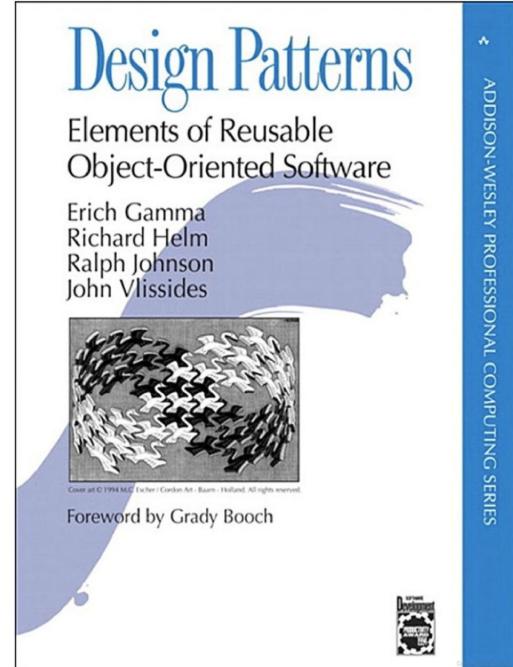
Ray Distributed Design Patterns & APIs



Ray Basic Design Patterns

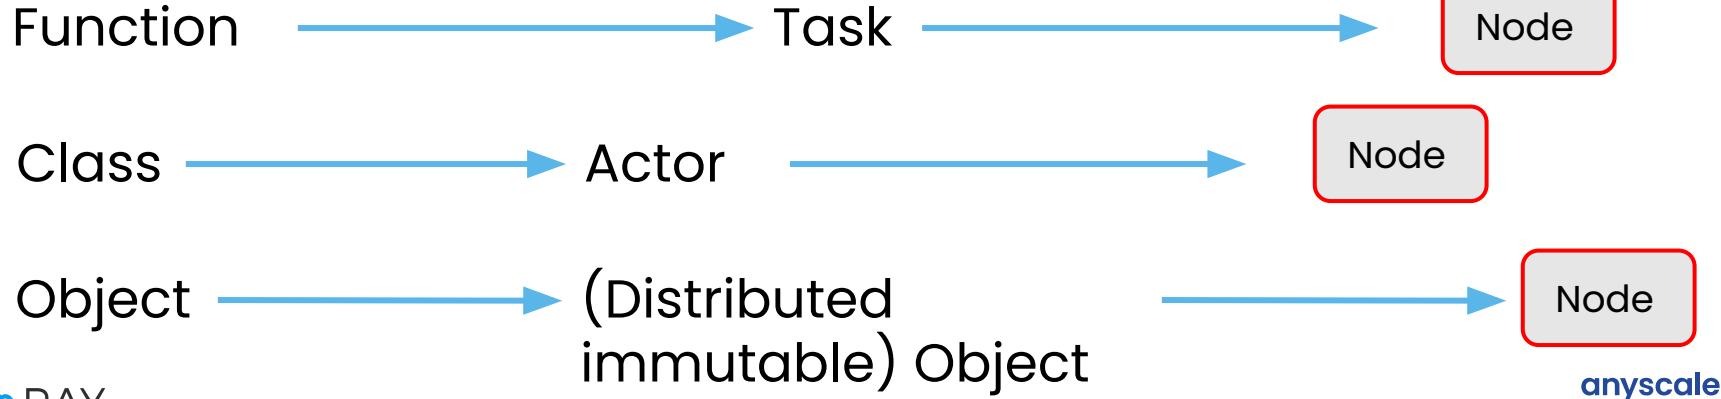
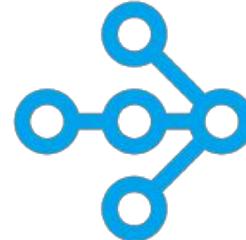
- **Ray Parallel Tasks**
 - Functions as stateless units of execution
 - Functions distributed across a clusters as tasks
- **Ray Objects or Futures**
 - Distributed (immutable) Object stored in cluster
 - Retrievable when available
 - Enable asynchronous execution of
- **Ray Actors**
 - Stateful service on a cluster
 - Message passing and maintains state

1. [Patterns for Parallel Programming](#)
2. [Ray Design Patterns](#)
3. [Ray Distributed Library Integration Patterns](#)





Python → Ray Basic Patterns



Function → Task

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```

Class → Actor

```
@ray.remote(num_gpus=1)
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
        return self.value

c = Counter.remote()
id4 = c.inc.remote()
id5 = c.inc.remote()
```

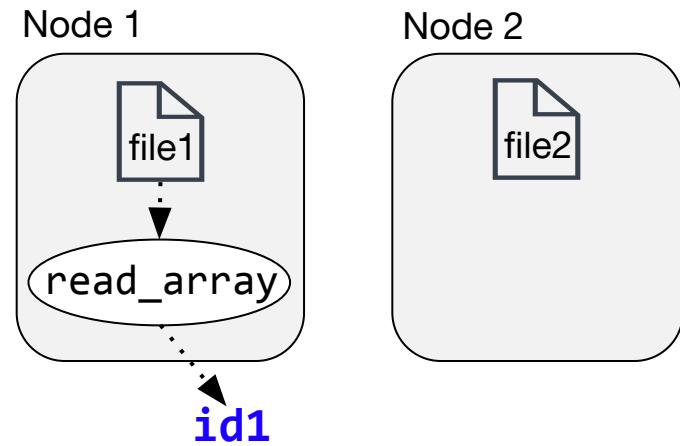
Task API

Blue variables are ObjectRef IDs
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```



Return **id1** (future) immediately,
before **read_array()** finishes

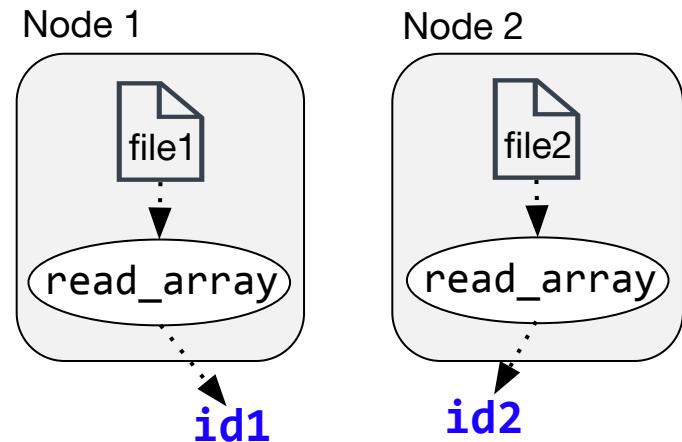
Task API

Blue variables are Object IDs
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```



Dynamic task graph:
build at runtime

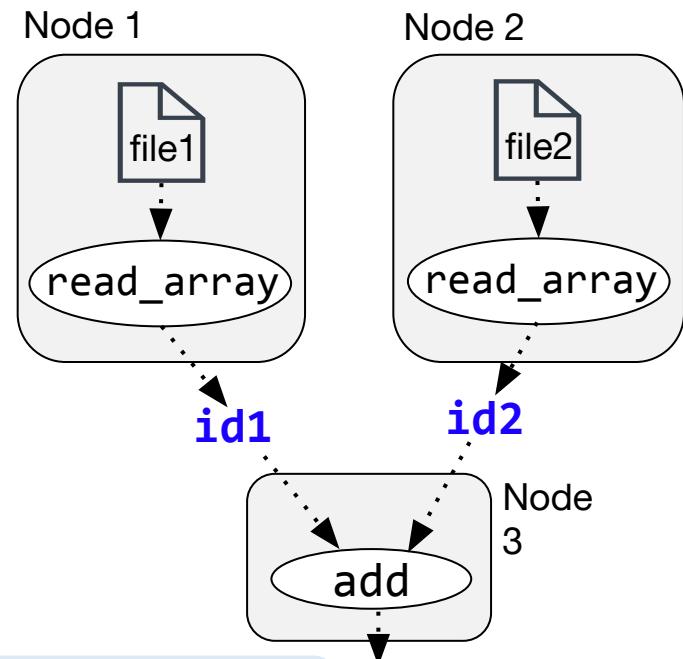
Task API

Blue variables are Object IDs
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```



Every task scheduled,
but not finished yet

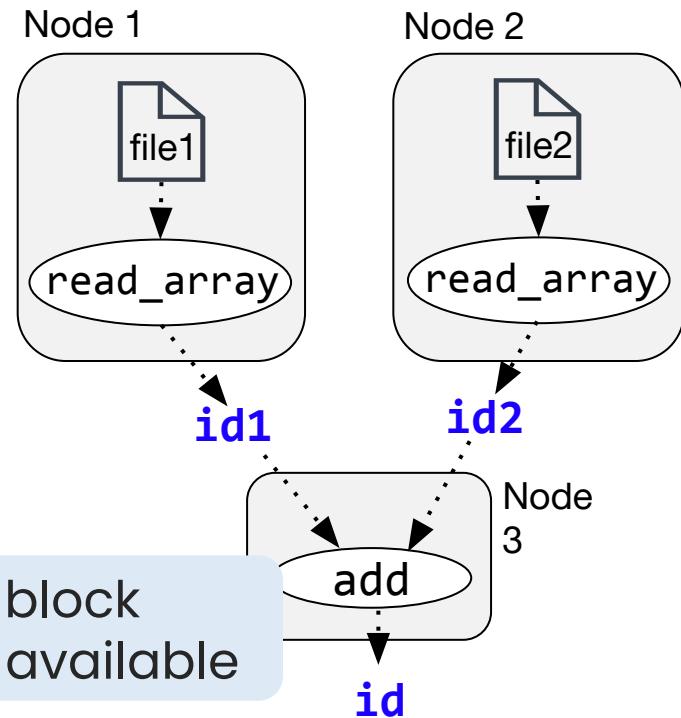
Task API

Blue variables are Object IDs
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```



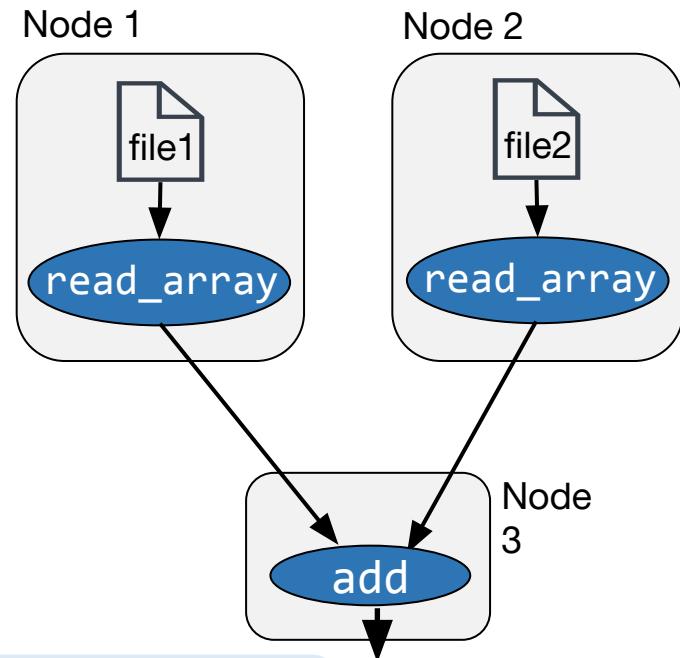
Task API

Blue variables are Object IDs
(similar to futures)

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

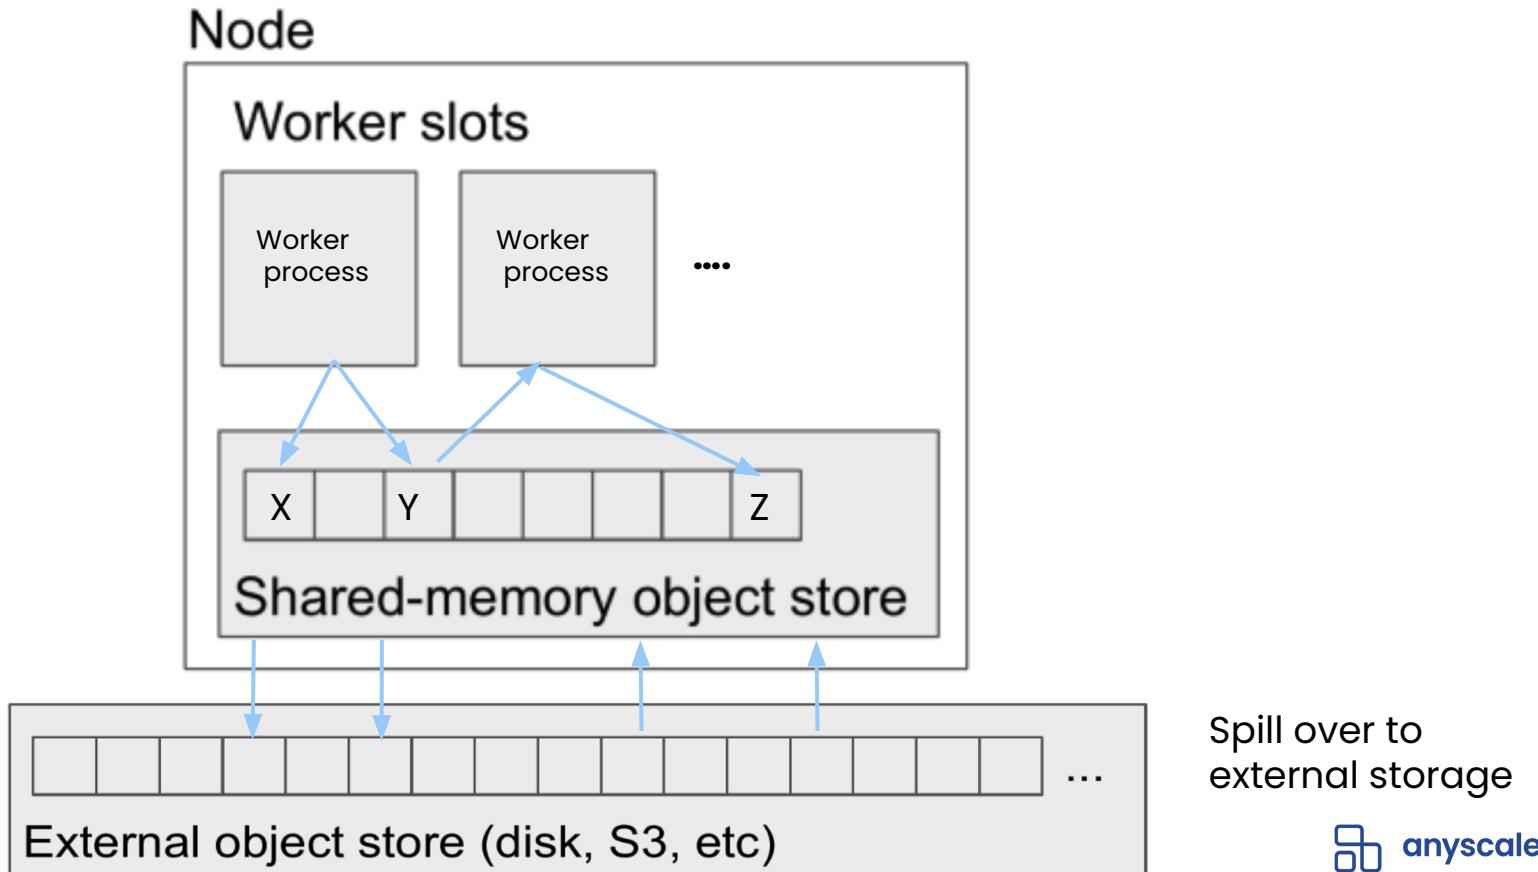
@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```

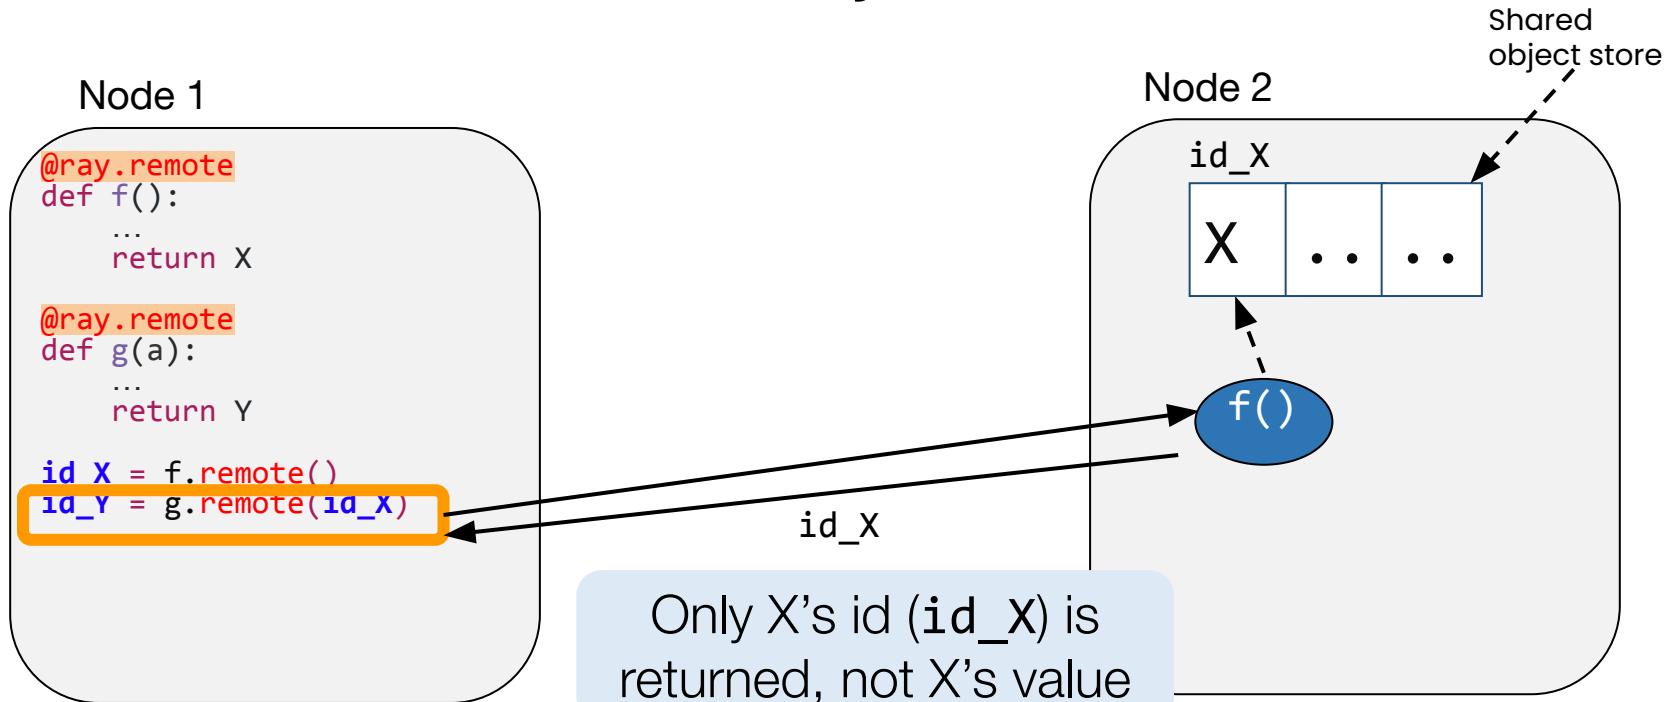


Task graph executed to sum
compute sum

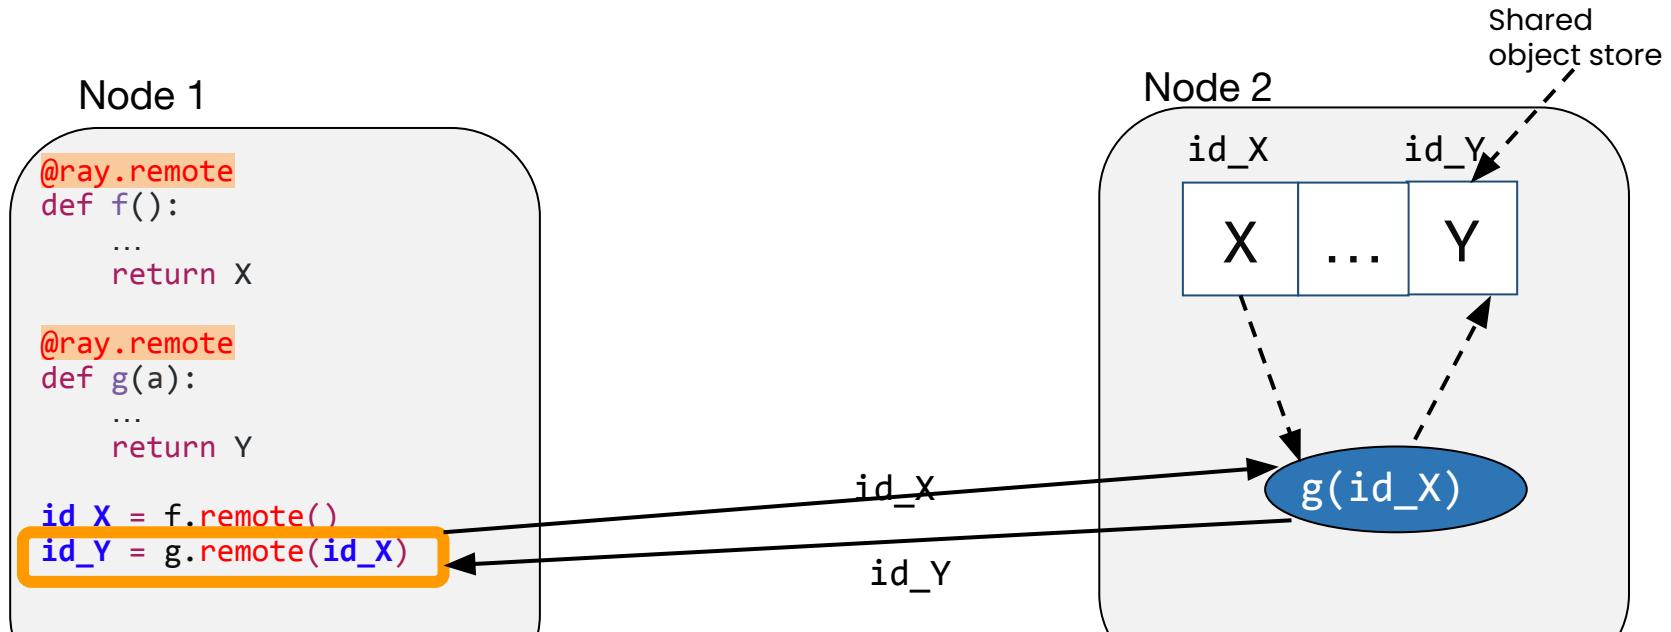
Distributed Immutable object store



Distributed Immutable object store



Distributed object store

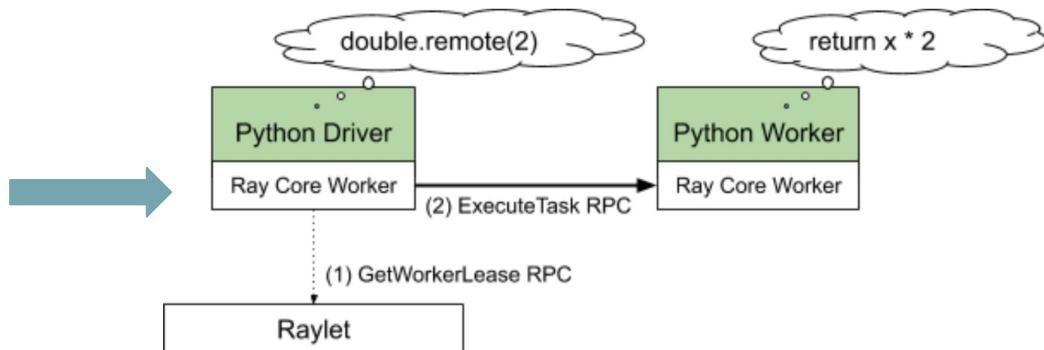


`g(id_X)` is scheduled on same node, so X is never transferred

How Raylet Schedules Tasks

Basic Ray Task Call

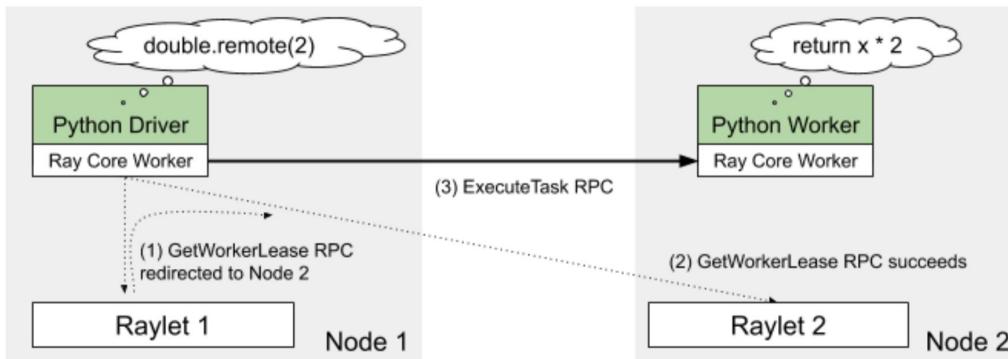
```
@ray.remote  
def double(x):  
    return x * 2  
  
fut1 = double.remote(2)  
assert ray.get(fut1) == 4
```



Components in green boxes represent Python code. Components in white boxes are part of the Ray common runtime written in C++. Joined boxes represent a process. Any Python driver or worker can call into the Ray C++ core worker library to execute further tasks. In this figure, all processes are running on the same machine. Ray uses gRPC as a unified communication layer for both local and remote procedure calls.

Scaling to Multiple Nodes

1. The driver asks Raylet 1 for a worker to execute `double`. It has no free workers, but Raylet 1 knows Raylet 2 has free resources, and redirects the request to Raylet 2.
2. The driver sends `ExecuteTask` to the remote Python worker leased from Raylet 2 over gRPC.

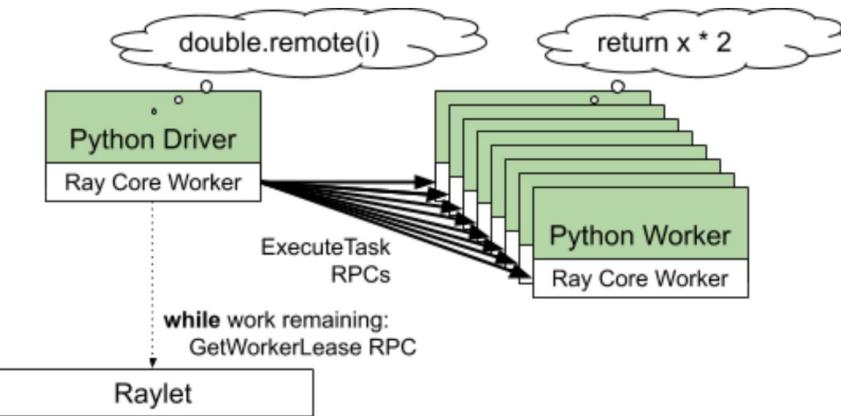


Tasks are sent to remote workers if there are no local resources available, transparently scaling Ray applications out to multiple nodes.

Caching Scheduling Decisions

```
futures = [double.remote(i)
for i in range(10000)]
ray.get(futures)

# [0, 2, 4, 6 ...]
```



Once a scheduling decision is made by the Raylet, the worker returned can be reused for other tasks with the same resource requirements and input dependencies. This amortizes scheduling RPC overhead when executing many similar tasks. To avoid unfair monopolization of workers when there are multiple processes trying to submit tasks, callers are only allowed to reuse workers within a few hundred milliseconds of initial grant.

Takeaways

- **Distributed computing is a necessity & norm**
- **Ray's vision: make distributed programming simple**
 - **Don't have to be distributed systems expert. Just use `@ray.remote` :)**
- **Scale your ML workloads with Ray Libraries**

Start learning Ray and contributing ...

Getting Started: pip install ray

Documentation (docs.ray.io)

Quick start example, reference guides, etc

Join Ray Meetup

Revived in Jan 2022. Next meetup March 2nd.

Meetup each month and publish recording to the members

<https://www.meetup.com/Bay-Area-Ray-Meetup/>

Forums (discuss.ray.io)

Learn / share with broader Ray community, including core team

Ray Slack

Connect with the Ray team and community

Social Media (@raydistrributed, @anyscalecompute)

Follow us on Twitter and linkedIn

GitHub

*Check out sources, file an issue, become a contributor, give us a **Star** :)*

<https://github.com/ray-project/ray>

Tutorials

<https://tinyurl.com/pycon-ray-tutorial>

Thank you!

Let's stay in touch:

jules@anyscale.com
<https://www.linkedin.com/in/dmatrix/>



@2twitme