

# Ray: A distributed framework for scaling AI & Python workloads



## UC San Diego

Jules S. Damji - [@2twitme](#)

May 18, 2023, Mandeville Auditorium, UCSD

# Few Important URLs

Keep these URLs open in your browser tabs

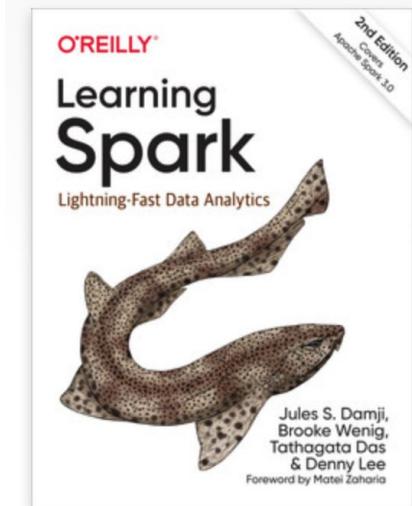
- GitHub: <https://github.com/dmatrix/ray-core-tutorial>
- Ray Documentation: <https://bit.ly/ray-core-docs>





# \$whoami

- Lead Developer Advocate, Anyscale & Ray Team
- Sr. Developer Advocate, Databricks, Apache Spark/MLflow Team
- Led Developer Advocacy, Hortonworks
- Held Software Engineering positions:
  - Sun Microsystems
  - Netscape
  - @Home
  - Loudcloud/Opsware
  - Verisign





# anyscale

**Who we are:** Original creators of Ray

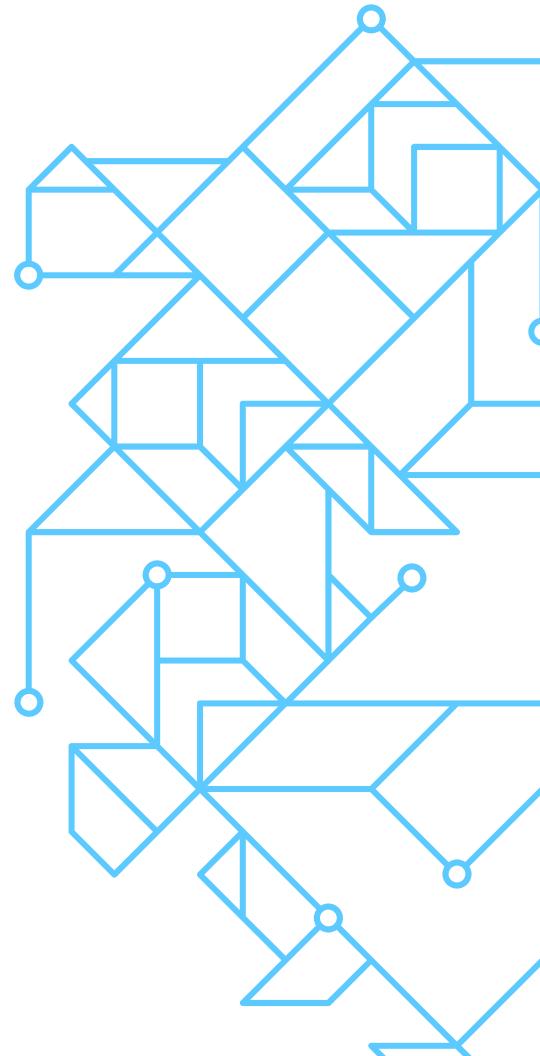
**What we do:** Unified compute platform to develop, deploy, and manage scalable AI & Python applications with Ray

**Why do it:** Scaling is a necessity, scaling is hard; make distributed computing easy and simple for everyone

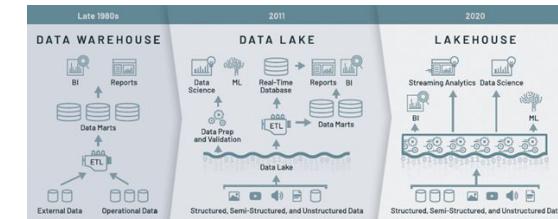
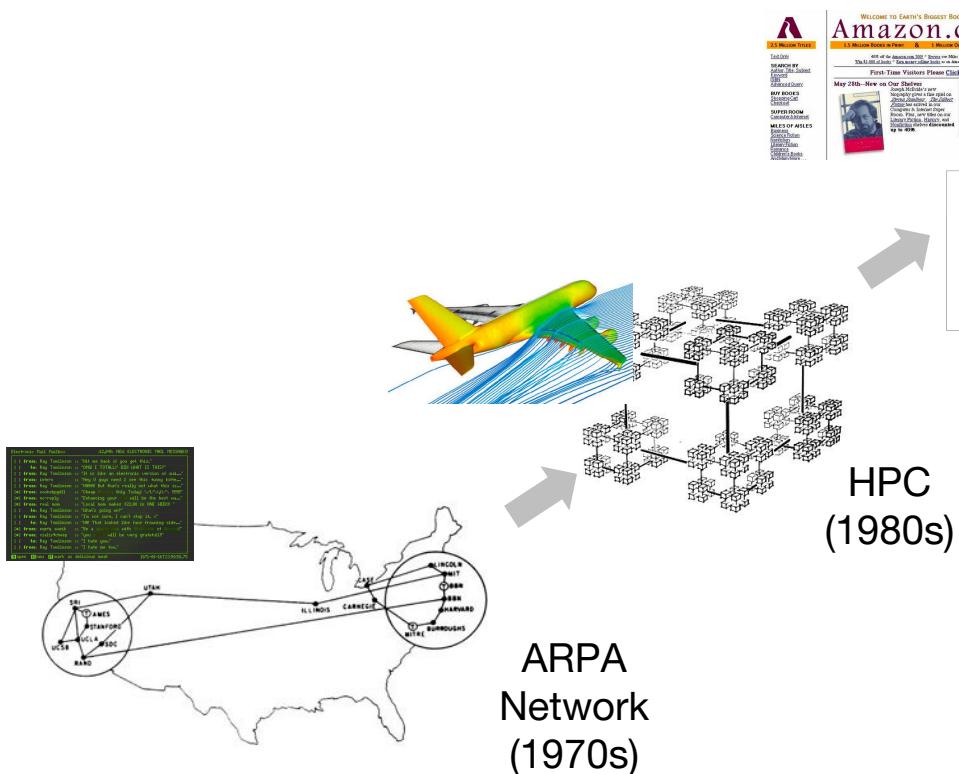
# Agenda

- Evolution of Distributed & Cloud Computing
- Distributed Computing: *Necessity not a norm*
- Why Ray & what is Ray
- Ray Ecosystem: Libraries & Integrations
- Ray & LLMs
- Ray core module - 1
- Q & A

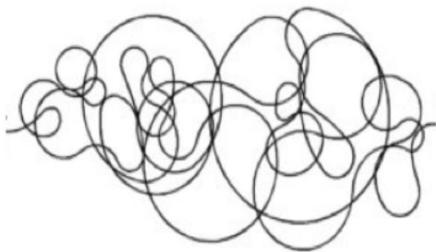
The future of computing is  
distributed



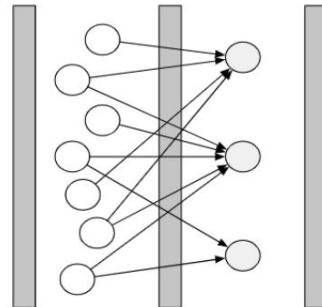
# Distributed systems not new...



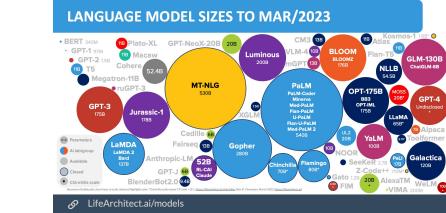
# Cluster usage, by generation

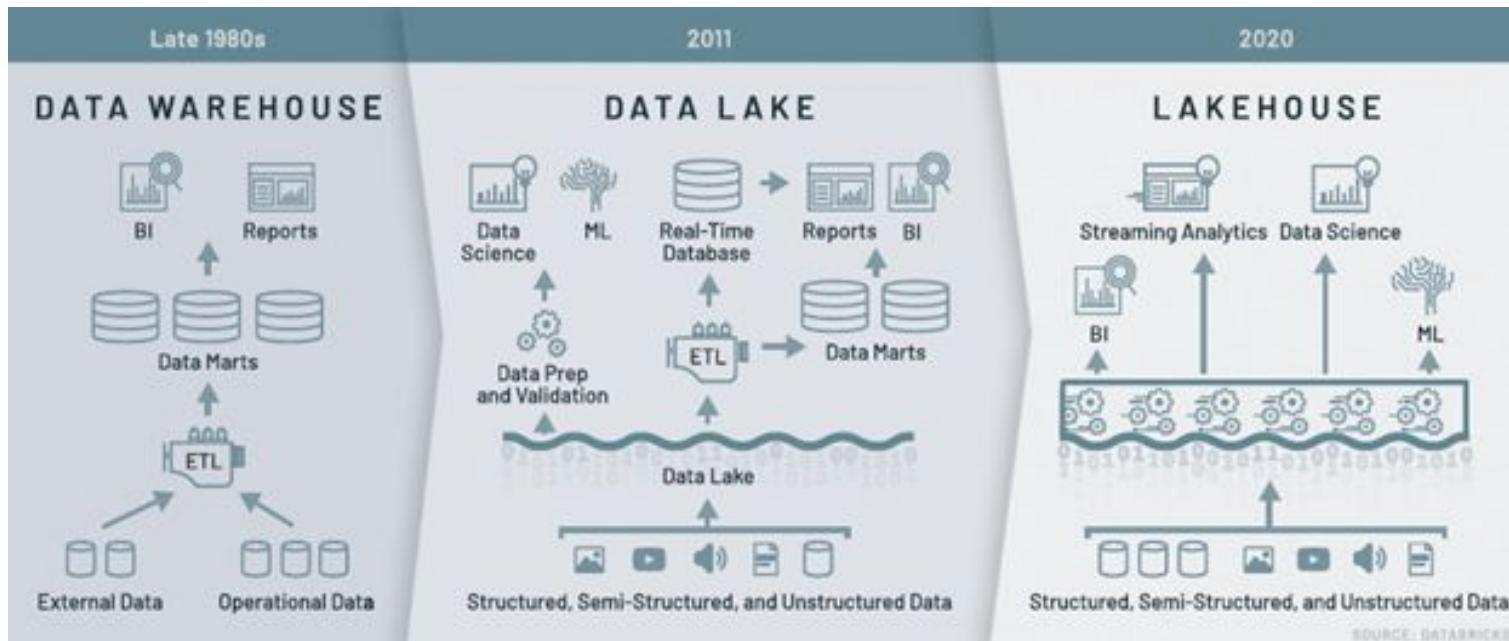


1990s



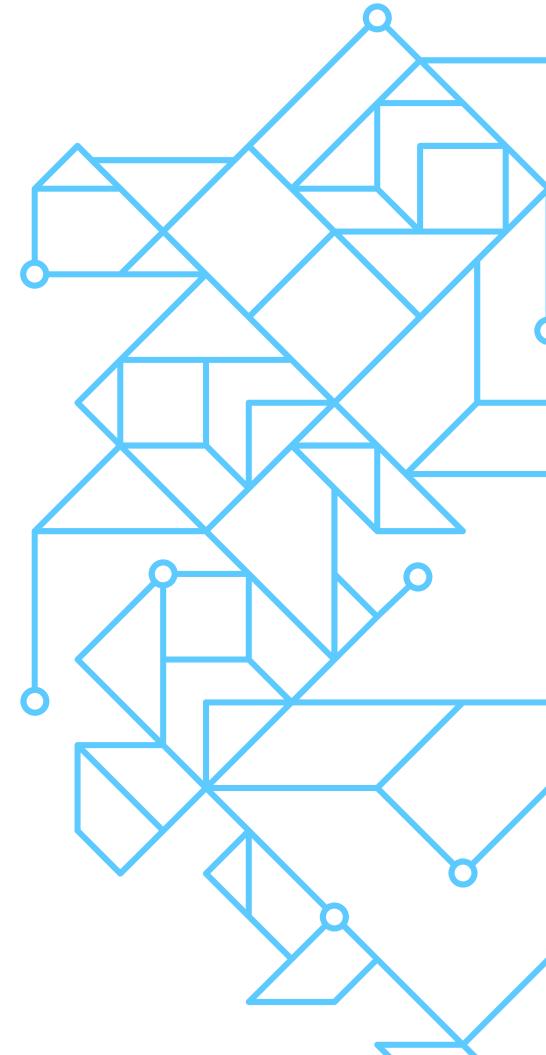
2000s





**How to scale AI applications using massive data in Lakehouse in the cloud?**

# The Evolution and trend in cloud computing

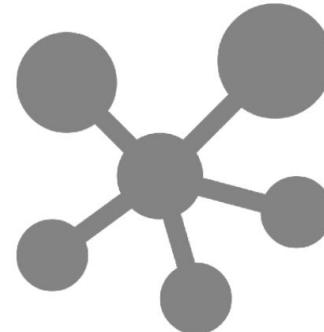


# Too Big To Fit, scale-up vs. scale-out

When an application becomes too big or too complex to run efficiently on a single server, there are some options:

- migrate to a larger server, and buy bigger licenses – that's called *vertical scale-up*
- distribute data+compute across multiple servers – that's called *horizontal scale-out*

The histories of **MPI**, **Hadoop**, **Spark**, **Dask**, etc., represent generations of scale-out, which imply **trade-offs** both for the risks (losing partitions, split-brain, etc.) as well as the inherent overhead costs



# Too Big To Fit, scale-up vs. scale-out

When an application becomes too big or too complex to run efficiently on a single server, there are some options:

- migrate to a bigger server  
that's called **scale-up**
- distribute data across multiple servers  
that's called **scale-out**

**Cloud Computing has arguably been an embodiment of distributed systems practices for the past 15 years, whether for scale-up or scale-out**

The histories of **HDFS**, **Hadoop**, **Spark**, **Dask**, etc., represent generations of scale-out, which imply **trade-offs** both for the risks (losing partitions, split-brain, etc.) as well as the inherent overhead costs



# Formal definition of Cloud Computing

Professors **Ion Stoica** and **David Patterson** led EECS grad students to define cloud computing **formally** in 2009

“More than 17,000 citations to this paper...”

2019 follow-up:

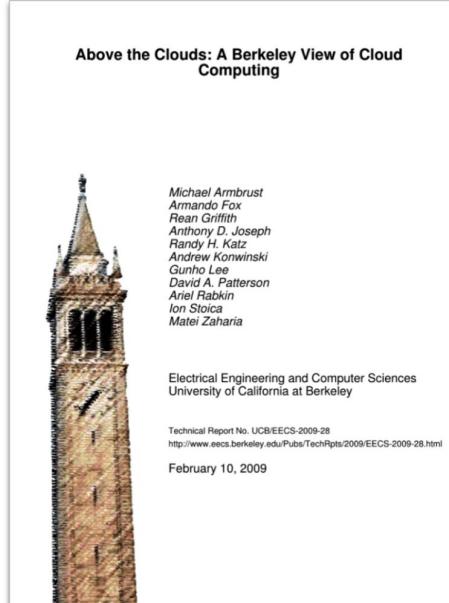
“We now predict ... that *Serverless Computing* will grow to dominate the future of cloud computing in the next decade”

The screenshot shows the Ariselab UC Berkeley website. At the top, there's a navigation bar with links for HOME, PEOPLE, PROJECTS, PUBLICATIONS, SPONSORS, DARE, ACADEMICS, NEWS, EVENTS, RISE CAMP, BLOGS, and JENKINS. The EVENTS link is highlighted with a red box. In the center, there's a large title for a publication: "Cloud Programming Simplified: A Berkeley View on Serverless Computing". Below the title, it says "BY ION STOICA / ON FEBRUARY 10, 2019 /". Underneath the title, it reads "David Patterson and Ion Stoica". A short description follows: "The publication of "Above the Clouds: A Berkeley View of Cloud Computing" on February 10, 2009 cleared up the considerable confusion about the new notion of "Cloud Computing." The paper defined what Cloud Computing was, where it came from, why some were excited by it, what were its technical advantages, and what were the obstacles and research opportunities for it to become even more popular. More than 17,000 citations to this paper and an abridged version in CACM—with more than 1000 in the past year—document that it continues to shape the discussions and the evolution of Cloud Computing."

## Evolution of cloud patterns

Initially, cloud services were simplified to make them more recognizable for IT staff accustomed to VMware

- Patterson, et al., developed industry research methodology and eventually also a **pattern language** to describe distributed systems
- **AMPLab** foresaw how cloud use cases would progress in industry over the next decade, which greatly informed **Apache Spark**, etc.



## Cloud patterns & architectures



# Evolution of cloud patterns

Initially, cloud services were simplified to make them more recognizable for IT staff accustomed to VMware

- Patterson, et al., research method also a [pattern language](#) for distributed systems
- [AMPLab](#) foresaw how cloud use cases would progress in industry over the next decade, which greatly informed [Apache Spark](#), etc.

**Apache Spark was an important outcome from this collective area of research**

Above the Clouds: A Berkeley View of Cloud Computing

*Michael Armbrust  
Armando Fox  
Rean Griffith  
Anthony D. Joseph  
Randy H. Katz  
Andrew Konwinski  
Gunho Lee  
David A. Patterson  
Ariel Rabkin  
Ion Stoica  
Matei Zaharia*

Electrical Engineering and Computer Sciences  
University of California at Berkeley

Technical Report No. UCB EECS 2009-28  
[Http://www.eecs.berkeley.edu/Pubs/TechRpts/2009/EECS-2009-28.html](http://www.eecs.berkeley.edu/Pubs/TechRpts/2009/EECS-2009-28.html)  
February 10, 2009

# Evolution of cloud patterns

**Eric Jonas** noted >50% of **RISElab** grad students had never used Spark; also, how cloud was evolving in its second decade:

- “Decoupling of computation and storage; they scale separately and are priced independently”
- “The abstraction of executing a piece of code instead of allocating resources on which to execute that code”
- “Paying for the code execution instead of paying for resources you have allocated toward executing the code”

Cloud Programming Simplified: A Berkeley View on  
Serverless Computing



Eric Jonas  
Johann Schleier-Smith  
Vikram Sreekanti  
Chia-Che Tsai  
Anurag Khandelwal  
Qifan Pu  
Vaishal Shankar  
Joao Menezes Carreira  
Karl Krauth  
Neeraja Yadwadkar  
Joseph Gonzalez  
Raluca Ada Popa  
Ion Stoica  
David A. Patterson

Electrical Engineering and Computer Sciences  
University of California at Berkeley

Technical Report No. UCB/EECS-2019-3  
<http://www2.eeecs.berkeley.edu/Pubs/TechRpts/2019/EECS-2019-3.html>

February 10, 2019

2019

# Evolution of cloud patterns

Eric Jonas noted >50% of RISElab grad students had never used Spark, plus how cloud was evolving in its second decade:

- “Decoupling of computation and storage so they scale separately and independently”
- “The abstraction of the cloud from the developer, so you can write code instead of allocating resources on which to execute that code”
- “Paying for the code execution instead of paying for resources you have allocated toward executing the code”

**Ray is an important outcome  
from this collective area of  
research**

Cloud Programming Simplified: A Berkeley View on  
Serverless Computing

Eric Jonas  
Johann Schleier-Smith  
Vikram Sreekanti  
Chia-Che Tsai  
Anurag Khandelwal  
Qifan Pu  
Vaishaal Shankar  
Joao Menezes Carreira  
Karl Krauth  
Neeraja Yadwadkar  
Joseph Gonzalez  
Raluca Ada Popa  
Ion Stoica  
David A. Patterson

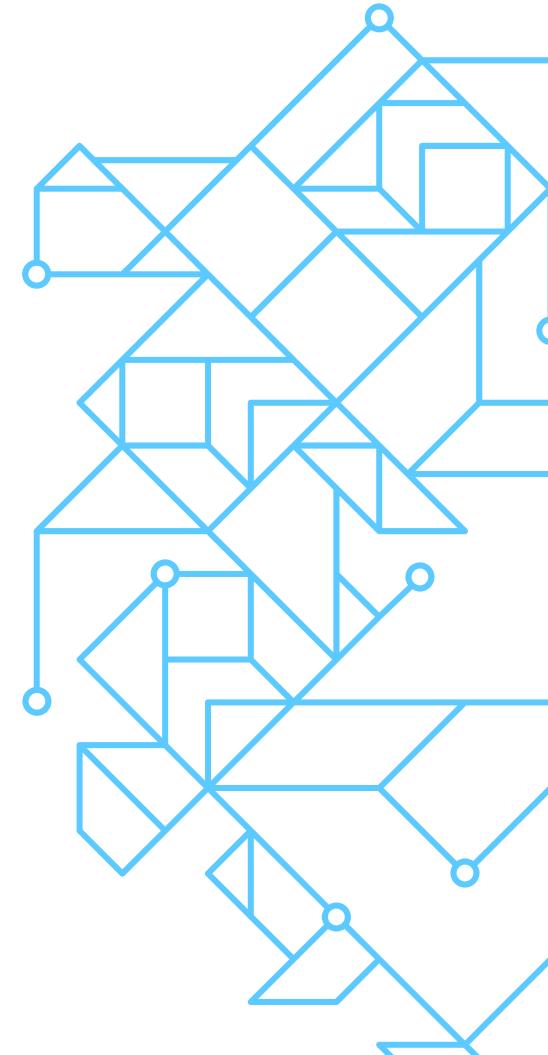
Electrical Engineering and Computer Sciences  
University of California at Berkeley

Technical Report No. UCB/EECS-2019-3  
<http://www2.eeecs.berkeley.edu/Pubs/TechRpts/2019/EECS-2019-3.html>

February 10, 2019

2019

# Why and what's Ray?



# Why Ray



Machine  
learning is  
pervasive

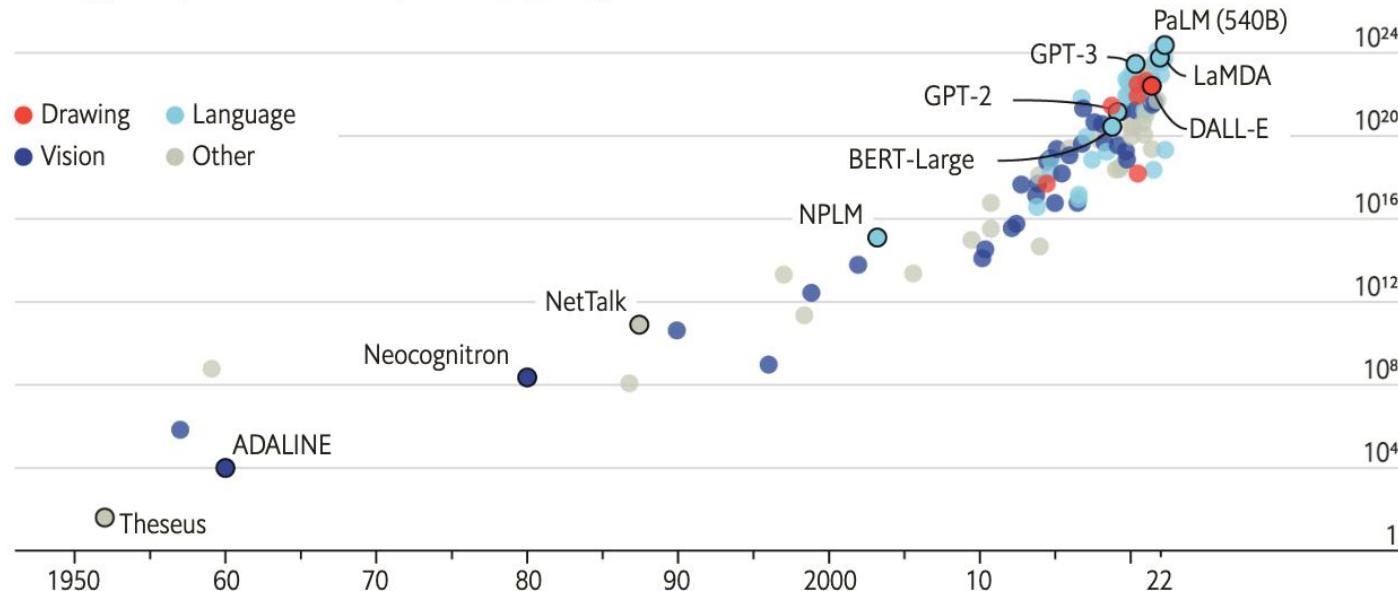
Distributed  
computing is a  
necessity

Python is the  
default  
language for  
DS/ML

# Blessings of scale ...

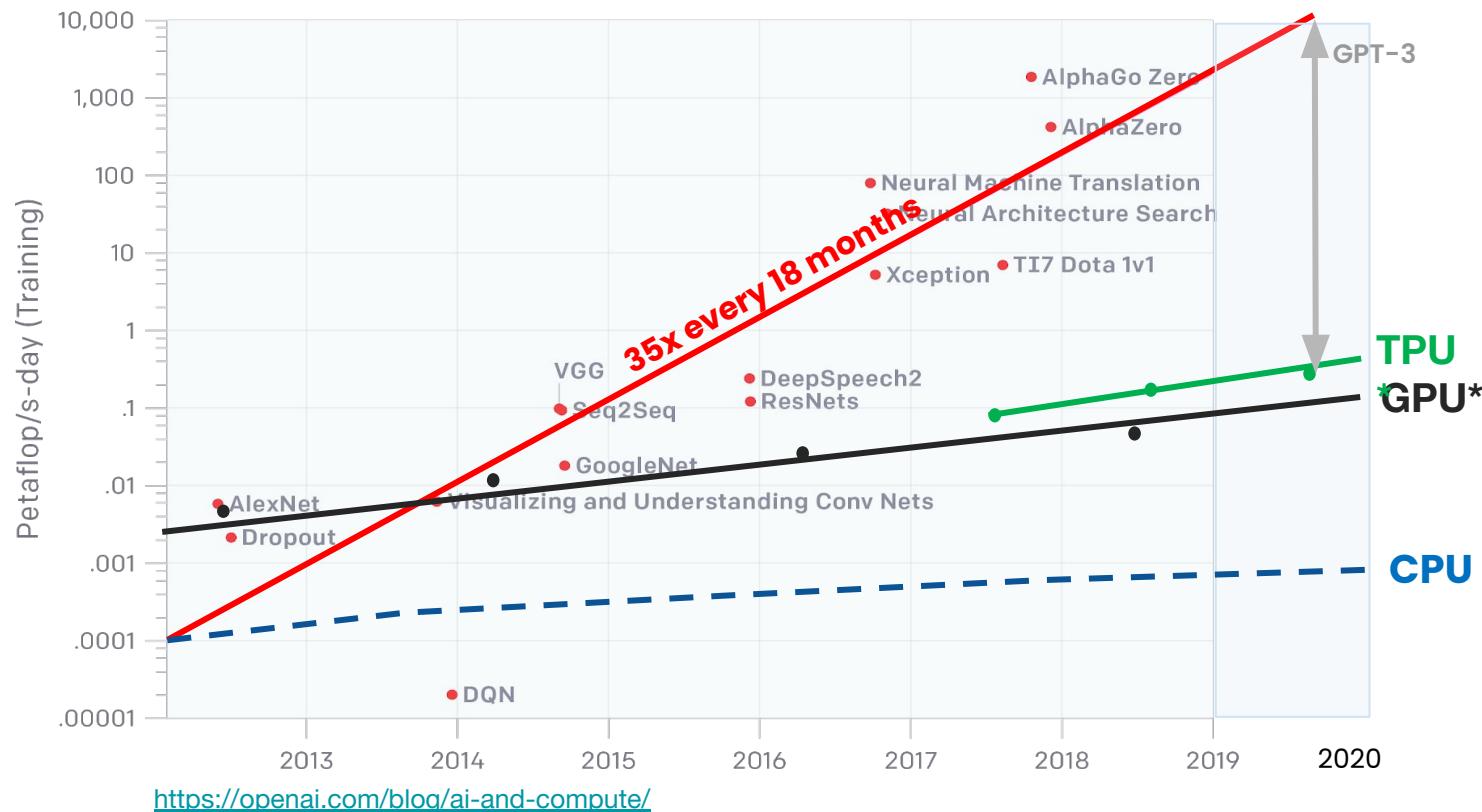
## The blessings of scale

AI training runs, estimated computing resources used  
Floating-point operations, selected systems, by type, log scale

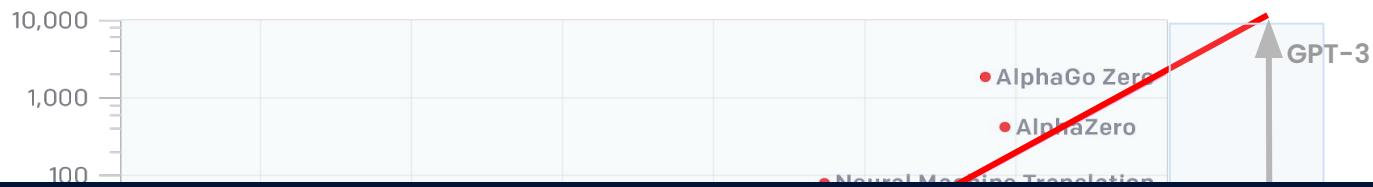


Sources: "Compute trends across three eras of machine learning", by J. Sevilla et al., arXiv, 2022; Our World in Data

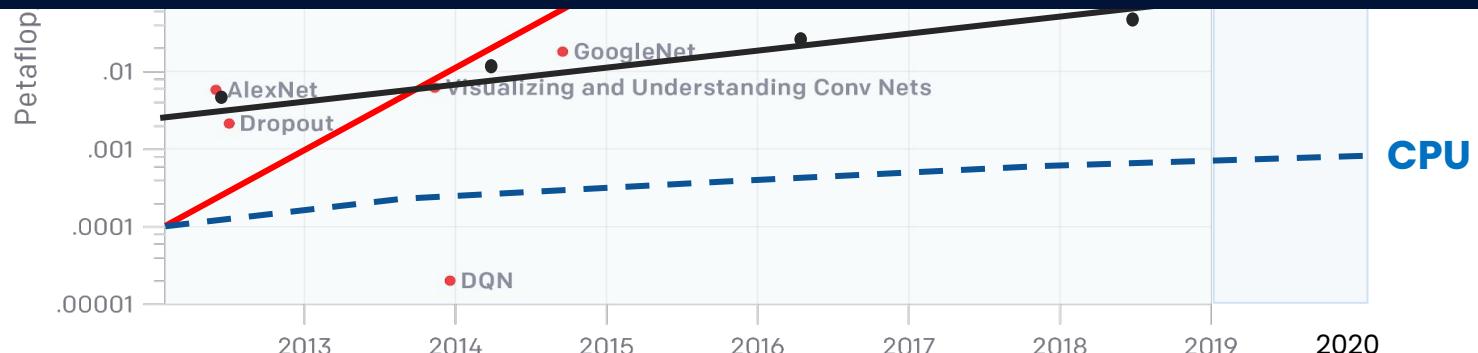
# Compute - supply demand problem



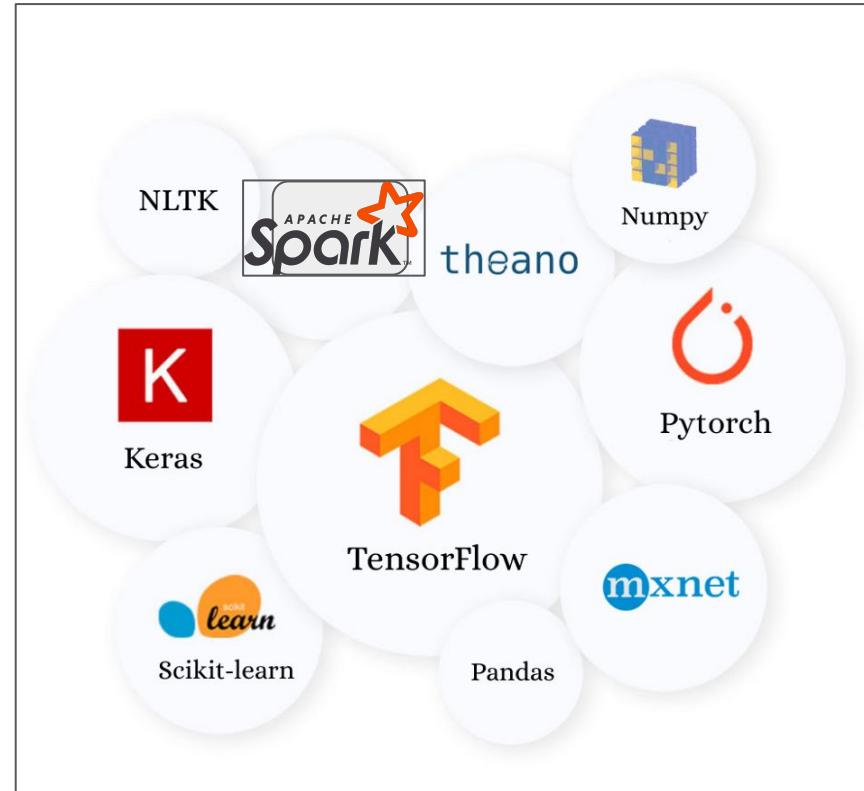
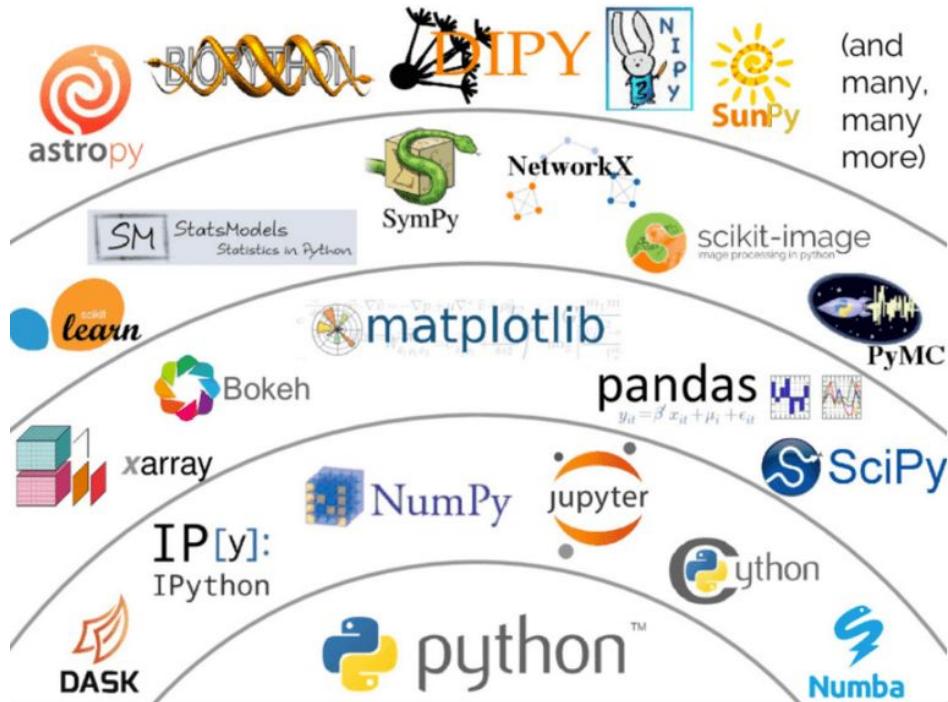
# Specialized hardware is not enough



## No way out but to distribute!



# Python data science ecosystem



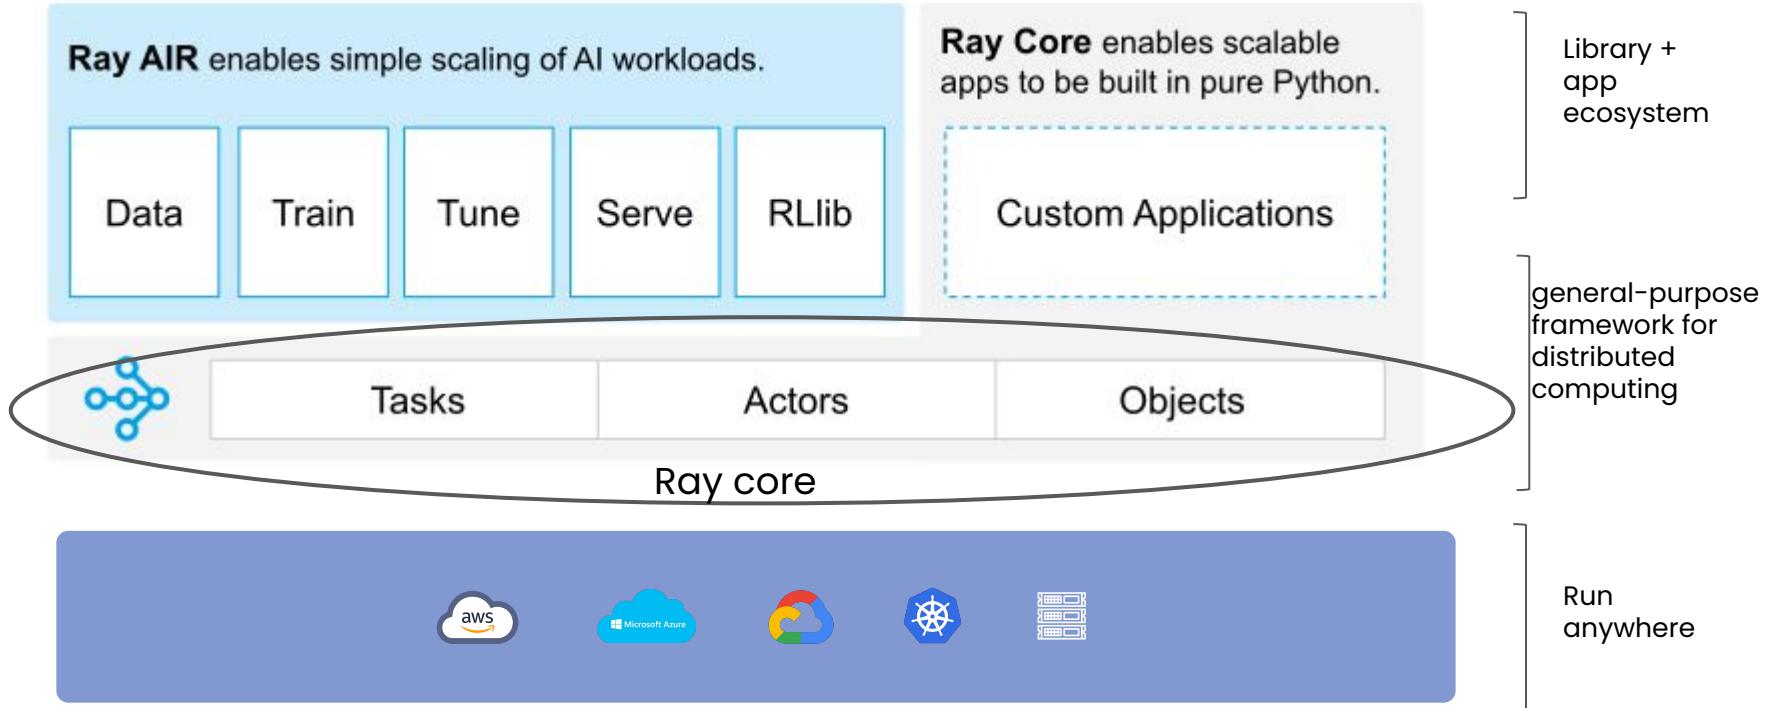
# What is Ray



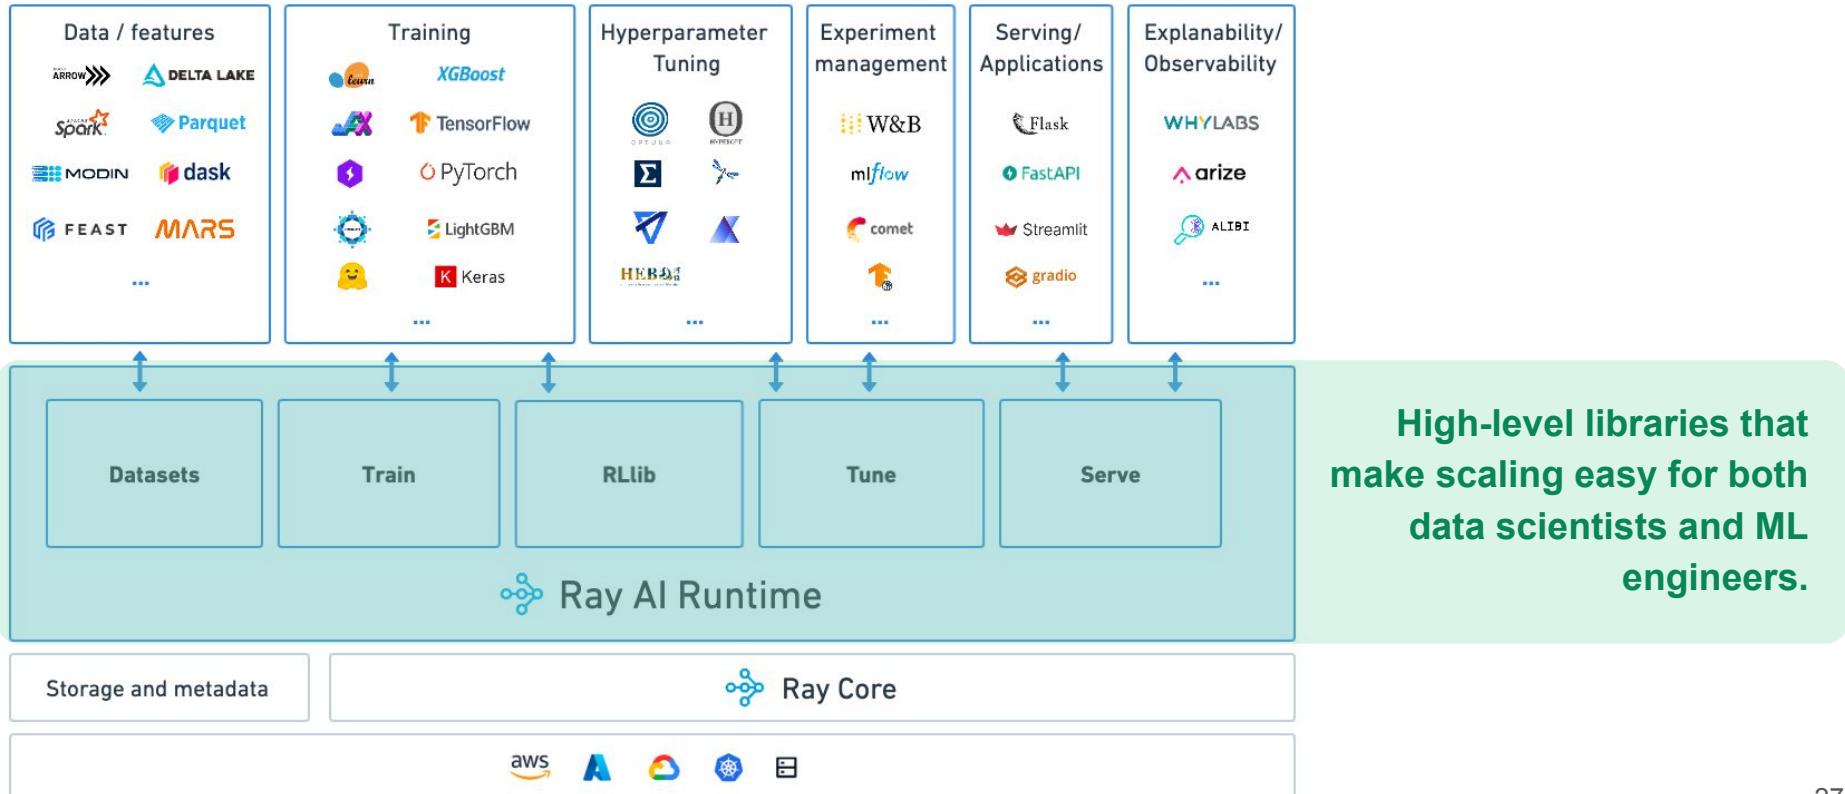
- A *simple/general-purpose* library for distributed computing
- An ecosystem of Python libraries (for scaling ML and more)
- Runs on laptop, public cloud, K8s, on-premise

A layered cake of functionality and capabilities for scaling ML workloads

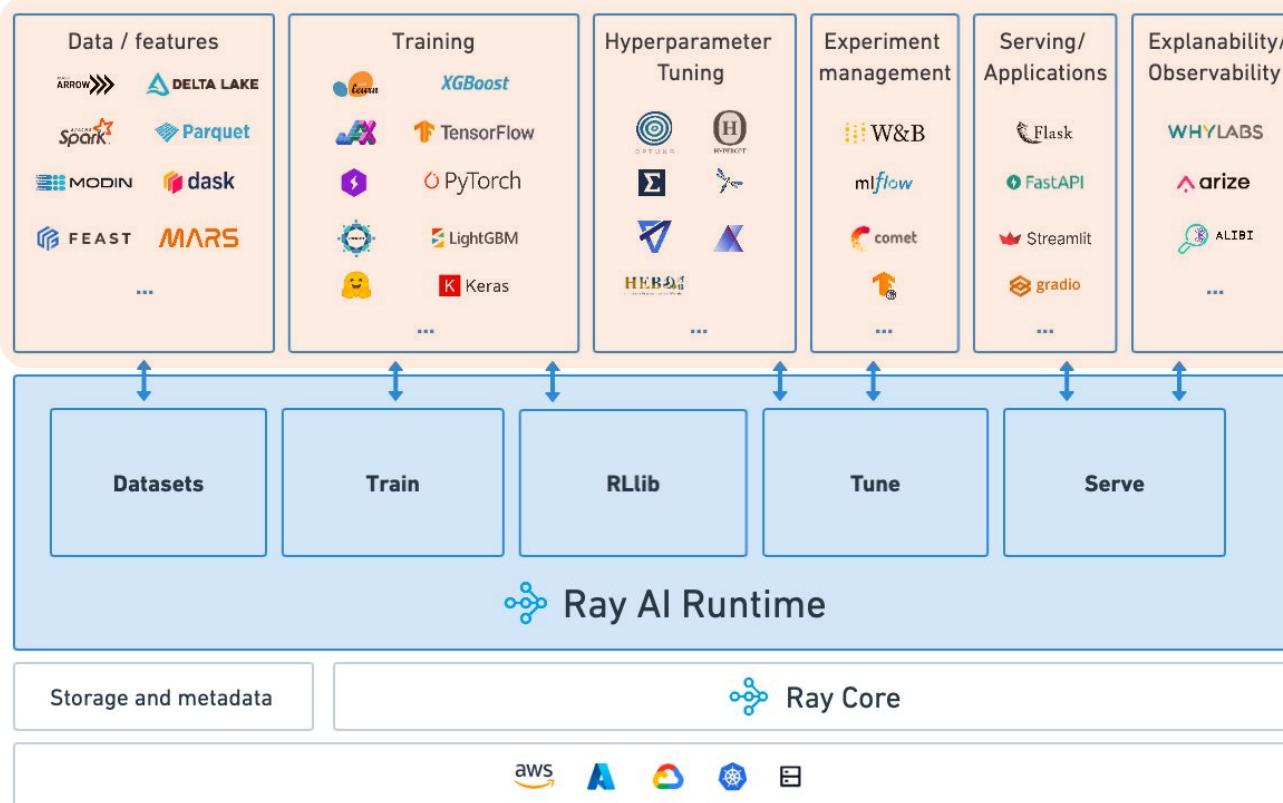
# A Layered Cake and Ecosystem



# Ray AI Runtime (AIR) is a scalable runtime for end-to-end ML applications

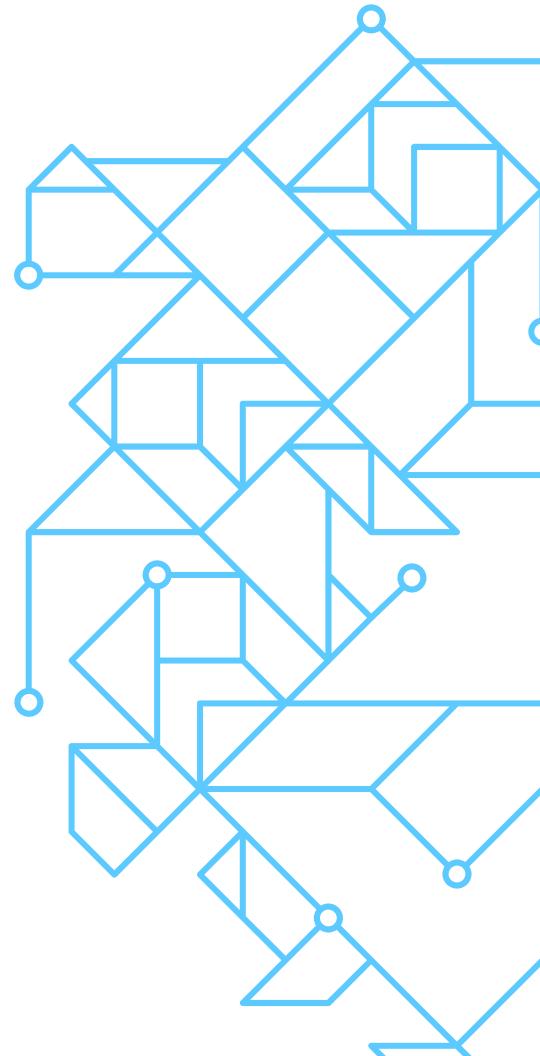


# Ray AI Runtime (AIR) is a scalable toolkit for end-to-end ML applications

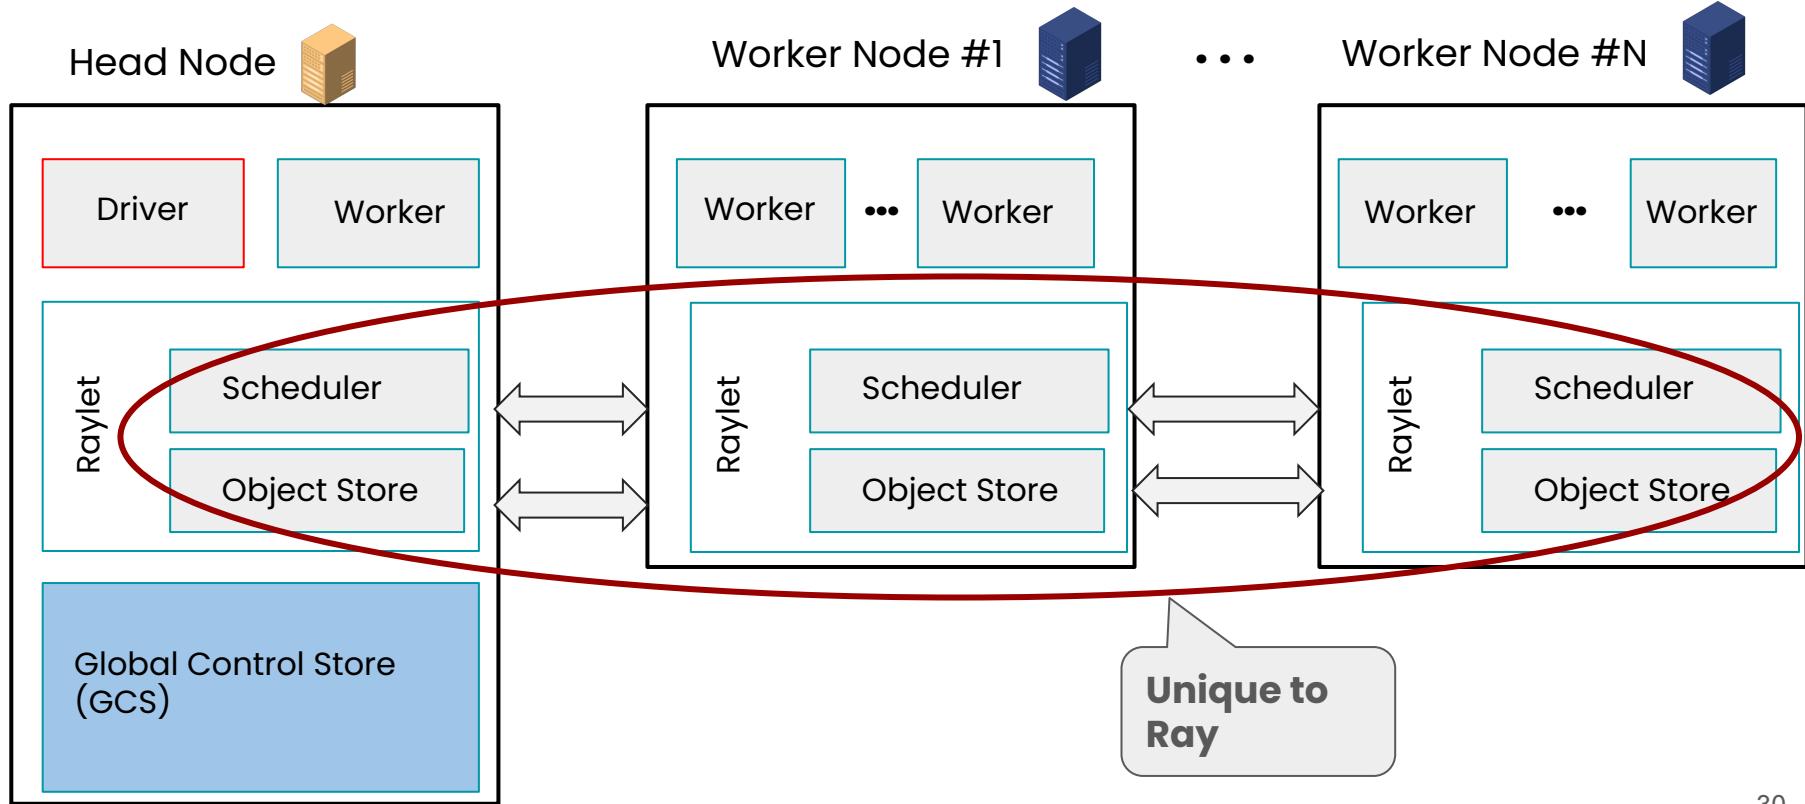


Scalable integrations with  
best-of-breed  
libraries/MLOps tools

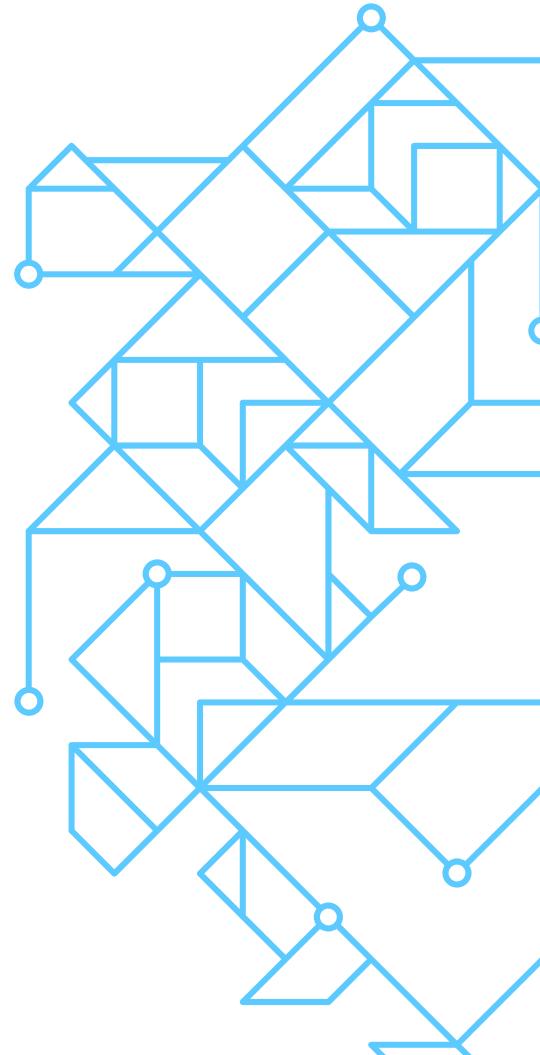
# Ray Architecture & Components



# An anatomy of a Ray cluster

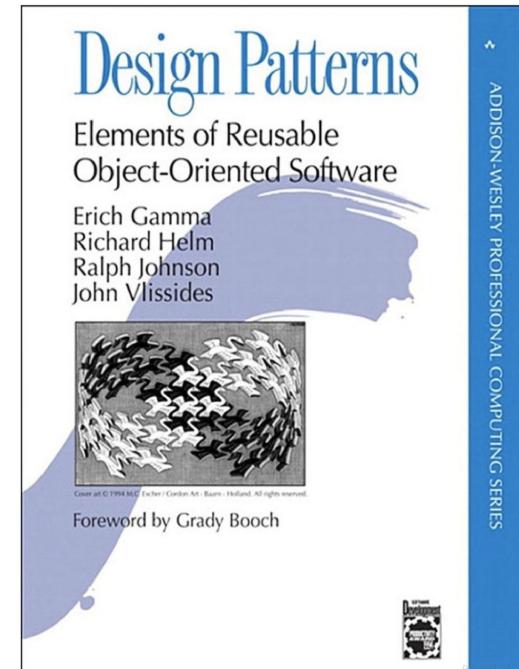


# Ray distributed design & scaling patterns & APIs



# Ray Basic Design Patterns

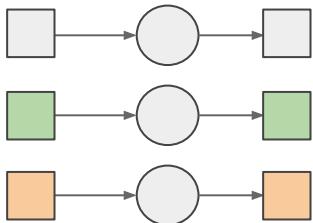
- Ray Parallel Tasks
  - Functions as stateless units of execution
  - Functions distributed across the cluster as tasks
- Ray Objects as Futures
  - Distributed (immutable objects) store in the cluster
  - Fetched when materialized
  - Enable massive asynchronous parallelism
- Ray Actors
  - Stateful service on a cluster
  - Enable Message passing



1. [Patterns for Parallel Programming](#)
2. [Ray Design Patterns](#)
3. [Ray Distributed Library Integration Patterns](#)

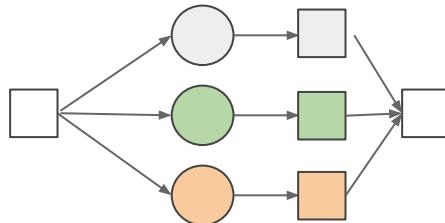
# Scaling Design Patterns

**Batch Training / Inference**



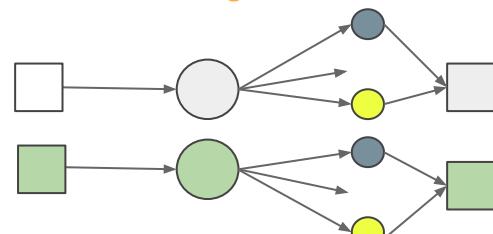
Different data / Same function

**AutoML**



Same data / Different function

**Batch Tuning**

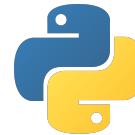


Different data / Same function /

Compute

Data

# Python → Ray APIs



```
def f(x):
    # do something with
    x:
        y = ...
    return y
```

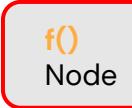
Task

```
@ray.remote
def f(x):
    # do something with
    x:
        Y = ...
    return y
f.remote()
@ray.remote
class Cls():
    def __init__(self, x):
        def f(self, a):
            ...
        def g(self, a):
            ...
cls = Cls.remote()
cls.f.remote(a)
```

Distributed



...



```
class Cls():
    def
    __init__(self, x):
        def f(self, a):
            ...
        def g(self, a):
            ...
...
```

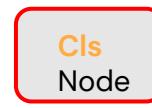
Actor

```
import numpy as np
a= np.arange(1, 10e6)
b = a * 2
```

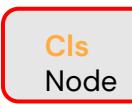
Distributed  
immutable  
object

```
import numpy as np
a = np.arange(1, 10e6)
obj_a = ray.put(a)
b = ray.get(obj_a) * 2
```

Distributed



...

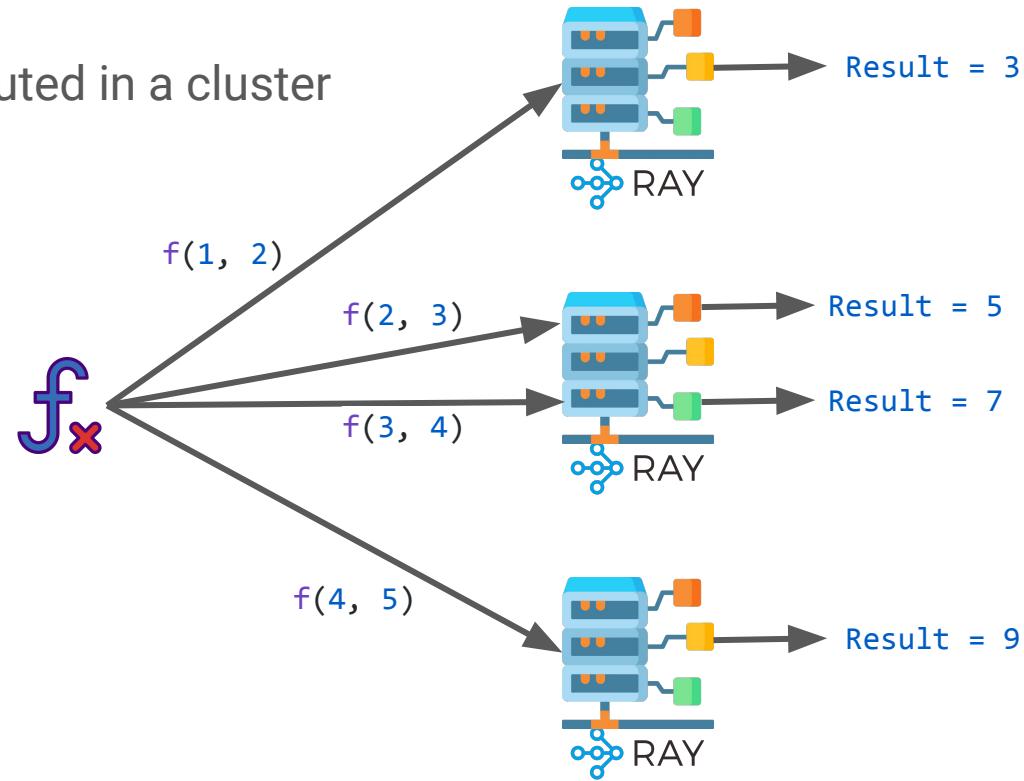


# Ray Task

A function  $f_x$  remotely executed in a cluster

```
@ray.remote(num_cpus=2)
def f(a, b):
    return a + b

f.remote(1, 2) # returns 3
```



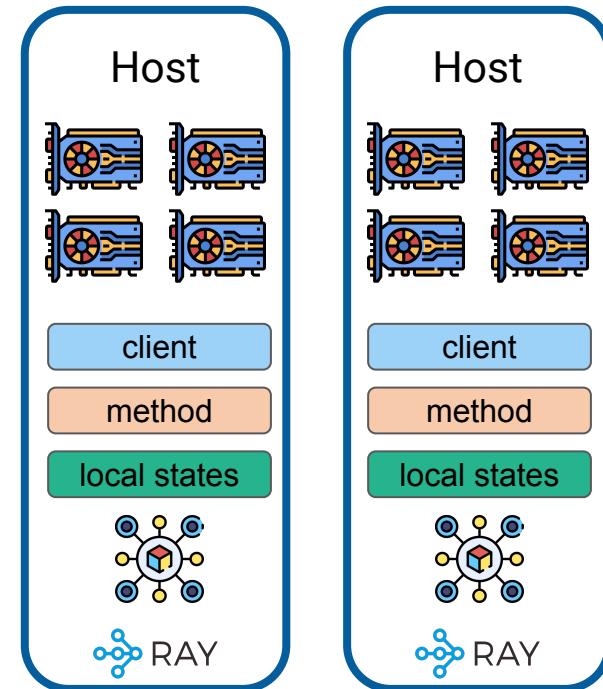
# Ray Actor

A **class**  remotely executed in a cluster

```
@ray.remote(num_gpus=4)
class HostActor:
    def __init__(self):
        self.num_devices = os.environ["CUDA_VISIBLE_DEVICES"]

    def f(self, output):
        return f"{output} {self.num_devices}"

actor = HostActor.remote() # Create an actor
actor.f.remote("hi") # returns "hi 0,1,2,3"
```



# Function → Task

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```

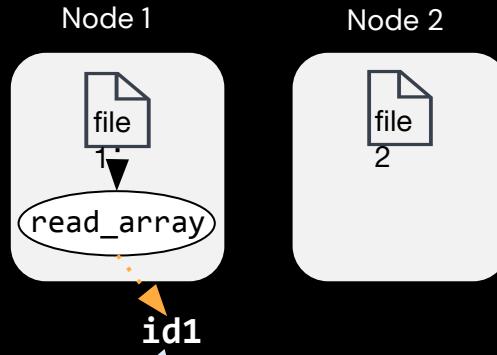
# Class → Actor

```
@ray.remote(num_gpus=1)
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
        return self.value
```

```
c = Counter.remote()
id4 = c.inc.remote()
id5 = c.inc.remote()
```

# Task API

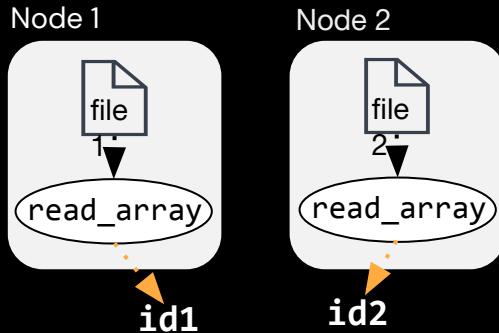
```
@ray.remote  
def read_array(file):  
    # read ndarray "a"  
    # from "file"  
    return a  
  
@ray.remote  
def add(a, b):  
    return np.add(a, b)  
  
id1 = read_array.remote(file1)  
id2 = read_array.remote(file2)  
id = add.remote(id1, id2)  
sum = ray.get(id)
```



Return `id1` (future) immediately,  
before `read_array()` finishes

# Task API

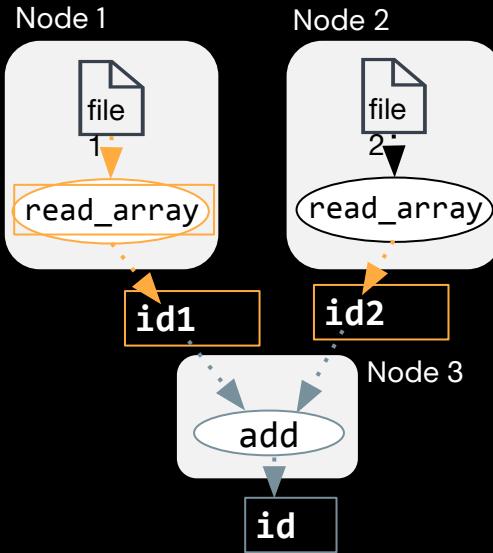
```
@ray.remote  
def read_array(file):  
    # read ndarray "a"  
    # from "file"  
    return a  
  
@ray.remote  
def add(a, b):  
    return np.add(a, b)  
  
id1 = read_array.remote(file1)  
id2 = read_array.remote(file2)  
id = add.remote(id1, id2)  
sum = ray.get(id)
```



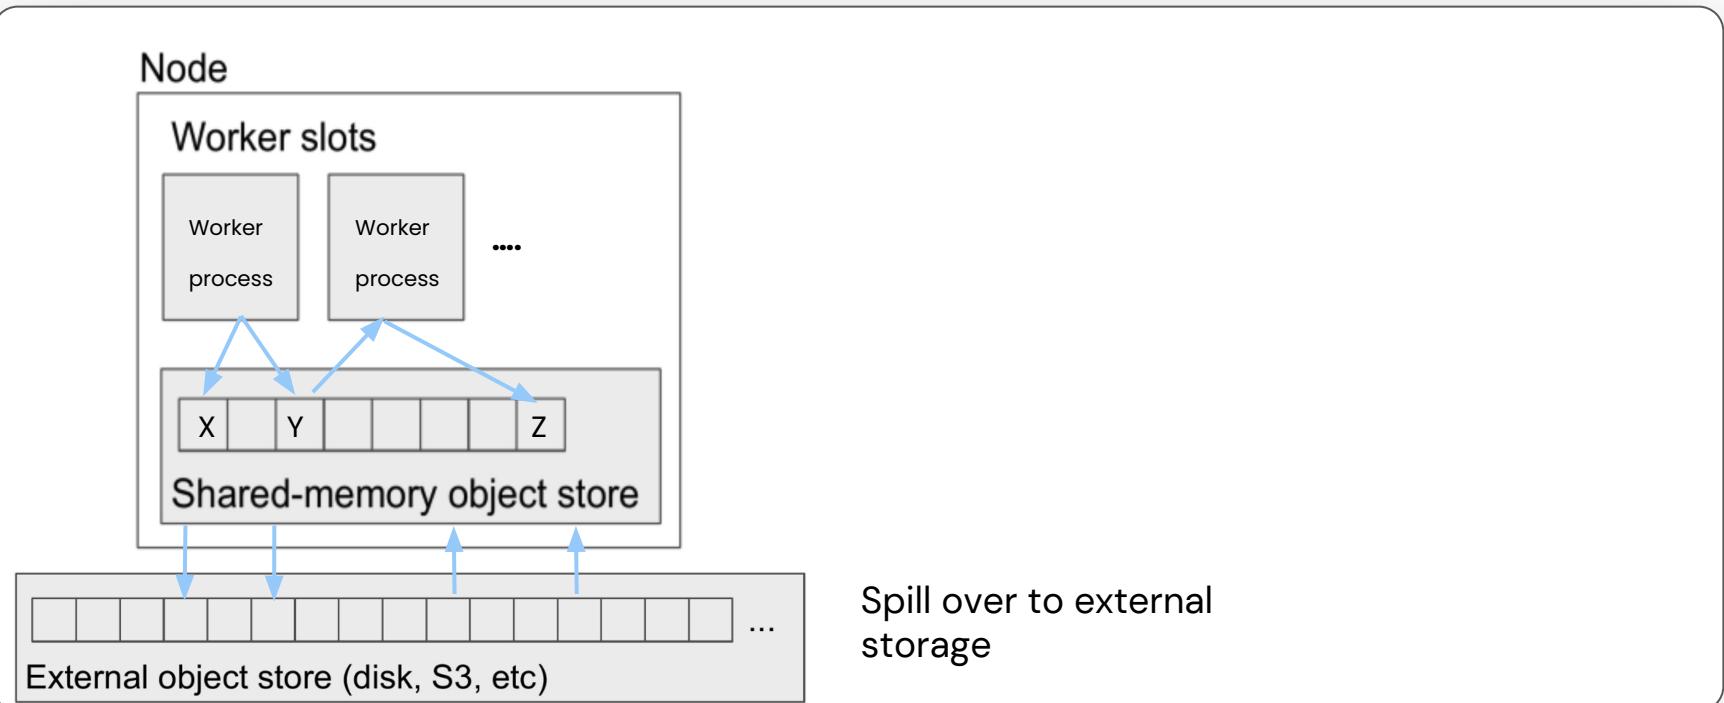
Dynamic task graph:  
build at runtime

# Task API

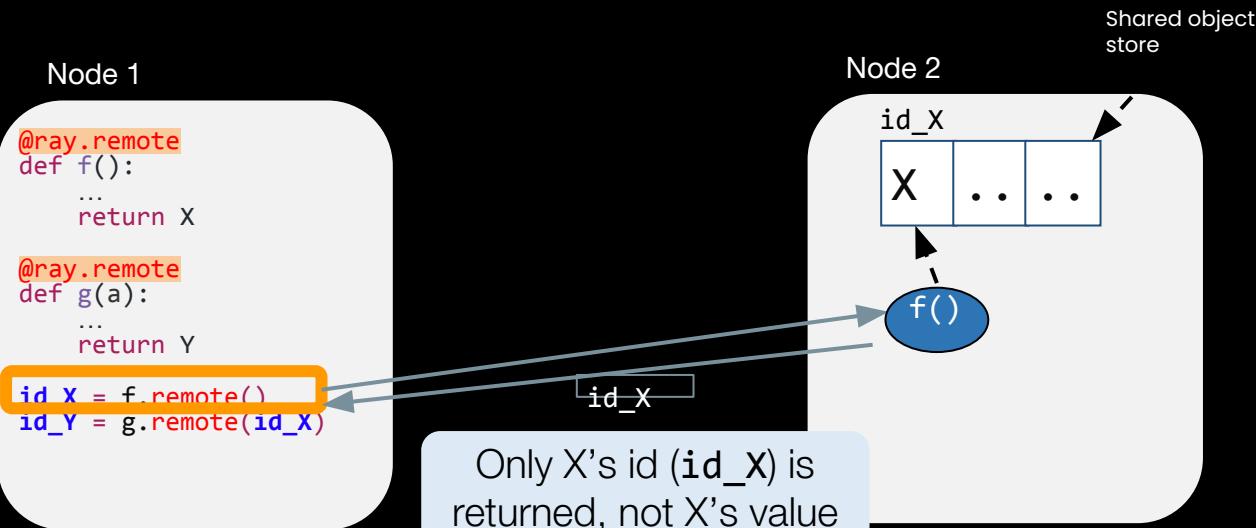
```
@ray.remote  
def read_array(file):  
    # read ndarray "a"  
    # from "file"  
    return a  
  
@ray.remote  
def add(a, b):  
    return np.add(a, b)  
  
id1 = read_array.remote(file1)  
id2 = read_array.remote(file2)  
id = add.remote(id1, id2)  
sum = ray.get(id) ray.get() block until  
result available
```



# Distributed Immutable object store

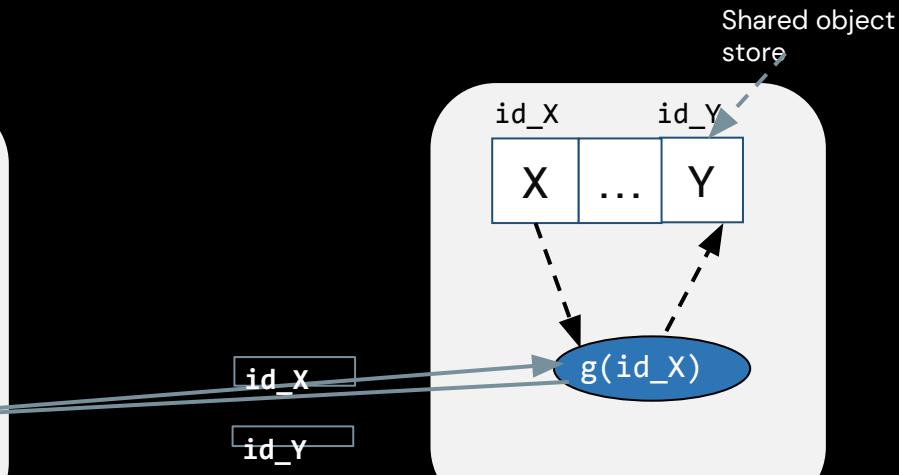


# Distributed object store



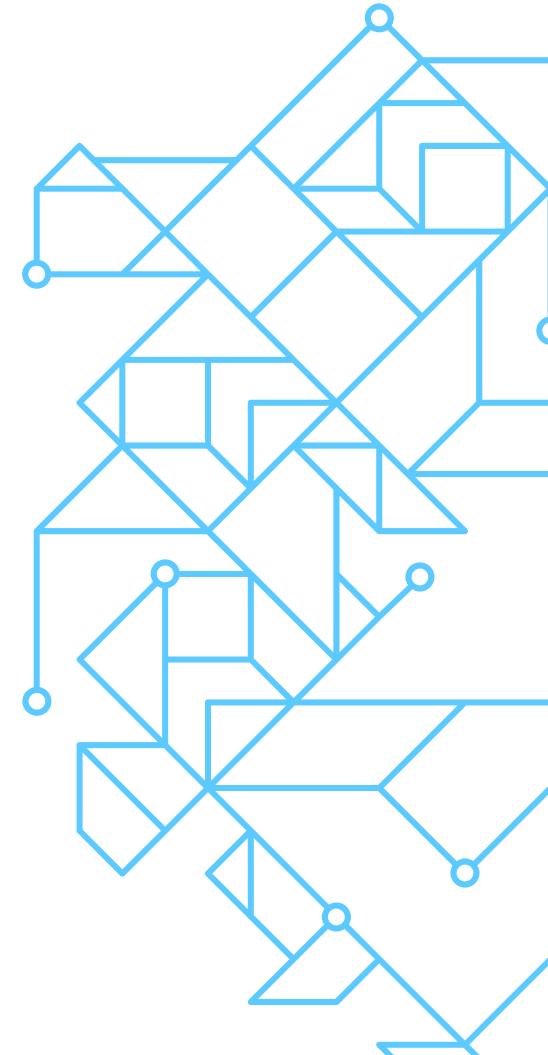
# Distributed object store

```
@ray.remote  
def f():  
    ...  
    return X  
  
@ray.remote  
def g(a):  
    ...  
    return Y  
  
id_X = f.remote()  
id_Y = g.remote(id_X)
```



`g(id_X)` is scheduled on same node, so `X` is never transferred

# Examples of Distributed Applications with Ray



# Distributed Applications with Ray

## ML Libraries

- Ray AI Runtime
- Distributed scikit-learn/Joblib
- Distributed XGBoost on Ray
- Ray Multiprocess Pool

All using Ray core APIs & patterns

## Experimenting & Monitoring Services

- WhyLabs
- Arize AI
- W & B
- MLflow

All using Ray core APIs & patterns

## ML Platforms & Integrations

- DoorDash ML platform
- AirFlow, PrefectePredibase AI
- Uber, Lyft
- Instacart
- Spotify

All using Ray core APIs & patterns

# Ray: Fastest Growing Scalable Compute Framework



McKinsey  
& Company



ERICSSON



cruise

Morgan Stanley



J.P.Morgan



Uber



RICARDO



verizon<sup>✓</sup>



**25,000+**

GitHub  
stars

**820+**

Community  
Contributors

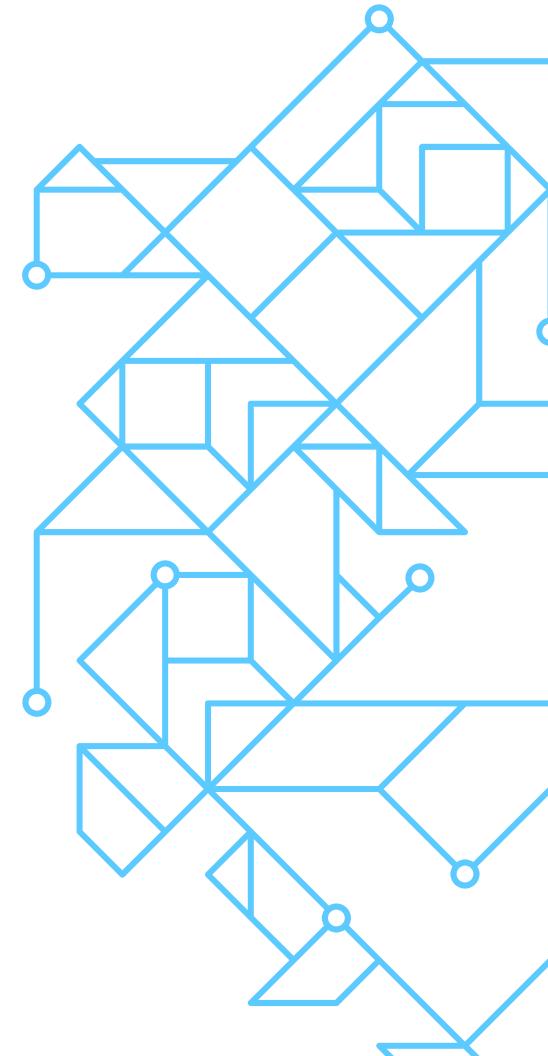
**5,000+**

Repositories  
Depend on Ray

**1,000+**

Organizations  
Using Ray

# Generative AI, LLMs & Ray



# Current state of the world..

LLM companies provide commercial APIs  
(note: co:here and OpenAI both use Ray internally)

co:here



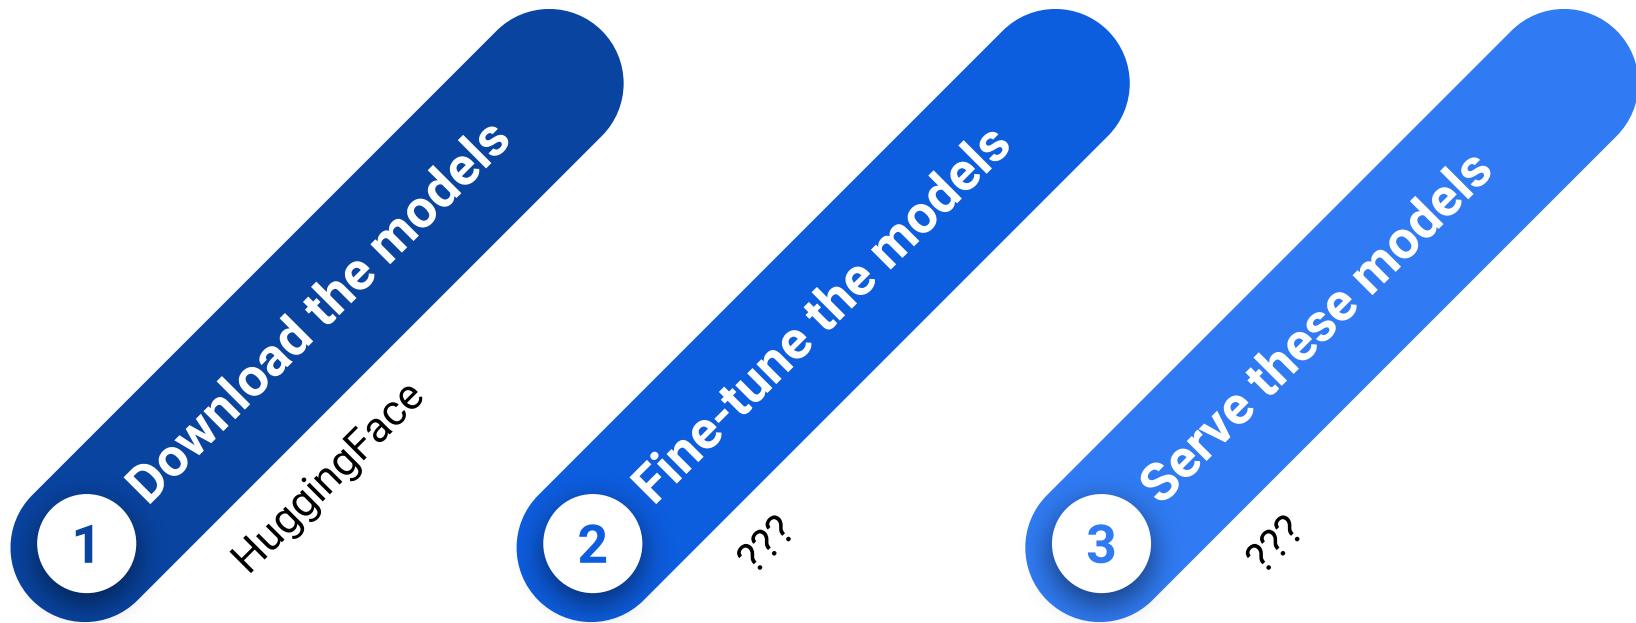
ANTHROPIC



Hugging Face



# Options for poor mortals ...



# What are pain points in training/serving LLMs?

01

**Scaling is costly and hard to manage**

- Need spot instance support
- Hard to run distributed workloads
- Hard to optimize CPUs/GPUs

02

**Existing serving / inference solutions don't scale**

- Individual replicas can't be distributed
- Need to be able to integrate business logic

03

**Distributed Training hard to get working right**

- Hyperparameters need to be tuned
- Need a platform to iterate very quickly at scale

# Ray provides generic platform for LLMs

01

## Simplify orchestration and scaling

- Spot instance support for data parallel training
- Easily spin up and run distributed workloads on any cloud
- Optimize CPUs/GPUs by pipelining w/ Ray Data

02

## Inference and serving

- Ability to support complex pipelines integrating business logic
- Ability to support multiple node serving

03

## Training

- Integrates distributed training with distributed hyperparameter tuning w/ ML frameworks

# What's the LLM stack for Generative AI?



## Technical stacks for Large language models (LLMs)



Model definition



Automatic model parallelization



GPU management and runtime orchestration



Compilation and runtime



GPU accelerator

HuggingFace for models

DeepSpeed for optimized Training

PyTorch for Framework

Ray for Orchestration

GPU or other hardware

Figure 5: Technical integration layered stack for LLM

# What about Generative AI?

## How Ray solves common production challenges for generative AI infrastructure

By Antoni Baum, Eric Liang, Jun Gong, Kai Fricke and Richard Liaw | March 20, 2023

This is part 1 of our generative AI blog series. In this post, we talk about how to use Ray to productionize common generative model workloads. An upcoming blog will deep dive into why projects like Alpa are using Ray to scale large models.

## Training 175B Parameter Language Models at 1000 GPU scale with Alpa and Ray

By Jiao Dong, Hao Zhang, Lianmin Zheng, Jun Gong, Jules S. Damji and Phi Nguyen | March 22, 2023

This is part 2 of our generative AI blog series. Here we cover how Ray empowers large language models (LLM) frameworks such as Alpa. To learn how to use Ray to productionize generative model workloads, see part 1.

<https://anyscale.com/blog>

## Faster stable diffusion fine-tuning with Ray AIR

By Kai Fricke | March 28, 2023

This is part 3 of our generative AI blog series that dives into a concrete example of how you can use Ray to scale the training of generative AI models. To learn more using Ray to productionize generative model workloads, see part 1. To learn about how Ray empowers LLM frameworks such as Alpa, see part 2.

## How to fine tune and serve LLMs simply, quickly and cost effectively using Ray + DeepSpeed + HuggingFace

By Waleed Kadous, Jun Gong, Antoni Baum and Richard Liaw | April 10, 2023

This is part 4 of our blog series on Generative AI. In the previous blog posts we explained why Ray is a sound platform for Generative AI, we showed how it can push the performance limits, and how you can use Ray for stable diffusion.

# Ray and Anyscale emerging as a standard for ML infrastructure

How Ray, a Distributed AI Framework, helps Power ChatGPT

**Ray breaks the \$1/TB barrier as the world's most cost-efficient sorting system**

By Frank Sifei Luan, UC Berkeley

**The Magic of Merlin: Shopify's New Machine Learning Platform**

by Isaac Vidas · Data Science & Engineering · Aug 6, 2020 · 11 minute read

Unleashing ML Innovation at Spotify with Ray



Contact Us Support My Account  
re:Invent Products Solutions Pricing Documentation Learn More

Blog Home

Category ▾

Edition ▾

AWS Big Data Blog

**Introducing AWS Glue for Ray: Scaling your data integration workloads using Python**

by Zach Mitchell, Ishan Gaur, Kinsukh Bahare, and Derek Liu | Oct 28, 2023 · 11 minute read



Han Li

Mar 17 · 9 min read · Listen



**Large Scale Deep Learning Training and Tuning with Ray at Uber**

Xu Ning, Di Yu, Michael Mui

## Distributed Machine Learning at Instacart

How Instacart uses distributed Machine Learning to efficiently thousands of models in production

Author



Netflix Technology Blog

Feb 13 · 11 min read · Listen



## Scaling Media Machine Learning at Netflix

# Key Takeaways

- Distributed computing is a necessity & norm
- Ray's vision: make distributed computing simple
  - Don't have to be distributed programming expert
- Build your own disruptive apps & libraries with Ray
- Scale your ML workloads with Ray libraries (Ray AIR)
- *Ray offers the compute substrate for Generative AI workloads*

# Ray Summit 2023



SPEAKERS

TRAINING

SPONSORS

WHO ATTENDS

REGISTER NOW

## THE PLACE FOR EVERYTHING RAY

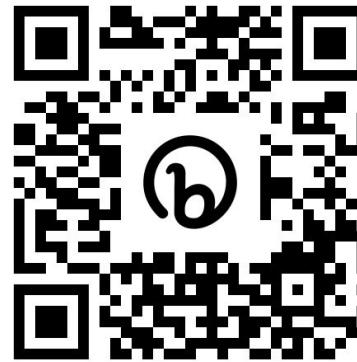
San Francisco Marriott Marquis | September 18-20

AI is moving fast. Get in front of what's next at Ray Summit 2023. Join the global Ray community in San Francisco for keynotes, Ray deep dives, lightning talks and more exploring the future of machine learning and scalable AI.

REGISTRATION IS OPEN



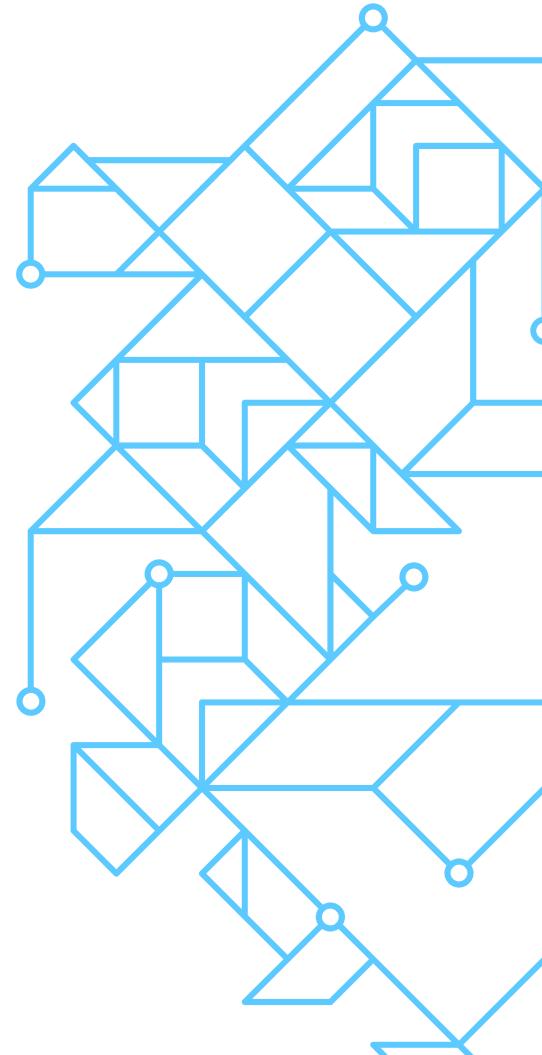
<https://bit.ly/raysummit2023>



Let's go with



<https://bit.ly/pydata-seattle-tutorial-2023>



# Thank you!

## Questions?

**email:** [jules@anyscale.com](mailto:jules@anyscale.com)

**twitter:** [@2twitme](https://twitter.com/2twitme)

